VDM-10 Language Manual

by

Peter Gorm Larsen
Kenneth Lausdahl
Nick Battle
John Fitzgerald
Sune Wolff
Shin Sahara
Marcel Verhoef
Peter Jørgensen
Tomohiro Oda

Overture – Open-source Tools for Formal Modelling

Document history

Month	Year	Version	Version of Overture.exe	Comment
April	2010		0.2	
May	2010	1	0.2	
February	2011	2	1.0.0	
July	2012	3	1.2.2	
April	2013	4	2.0.0	
March	2014	5	2.0.4	Includes RMs #16, #17, #18, #20
November	2014	6	2.1.2	Includes RMs #25, #26, #29

Contents

Chapter 1

Introduction

The Vienna Development Method (VDM) [Bjørner&78a, Jones90, Fitzgerald&08a] was originally developed at the IBM laboratories in Vienna in the 1970's and as such it is one of the longest established formal method. This document is a common language manual for the three dialects for VDM-SL, VDM++ and VDM-RT in the VDM-10 commonly agreed language revision. These dialects are supported by both VDMTools (minus VDM-RT) [Fitzgerald&08b] (in the appropriate version) as well as in the Overture open source tool [Larsen&10] built on top of the Eclipse platform. Whenever a construct is common to the three different dialects the term "VDM languages" will be used. Whenever a construct is specific to a subset of the VDM languages the specific dialect term mentioned above will be mentioned explicitly.

1.1 The VDM Specification Language (VDM-SL)

The syntax and semantics of the VDM-SL language is essentially the standard ISO/VDM-SL [ISOVDM96] with a modular extension[1]. Notice that all syntactically correct VDM-SL specifications are also correct in VDM-SL. Even though we have tried to present the language in a clear and understandable way the document is not a complete VDM-SL reference manual. For a more thorough presentation of the language we refer to the existing literature[2]. Wherever the VDM-SL notation differs from the VDM-SL standard notation the semantics will of course be carefully explained.

1.2 The VDM++ Language

VDM++ is a formal specification language intended to specify object oriented systems with parallel and real-time behaviour, typically in technical environments [Fitzgerald&05]. The language is based on VDM-SL [ISOVDM96], and has been extended with class and object concepts, which are

---

[1]A few other extensions are also included.

[2]A more tutorial like presentation is given in [Fitzgerald&98] whereas proofs in VDM-SL are treated best in [Jones90] and [Bicarregui&94].

also present in languages like Smalltalk-80 and Java. This combination facilitates the development of object oriented formal specifications.

1.3   The VDM Real Time Language (VDM-RT)

The VDM-RT language (formerly called VICE as an acronym for "VDM++ In Constrained Environments") is used to appropriately model and analyse Real-Time embedded and distributed systems [Mukherjee&00, Verhoef&06, Verhoef&07, Verhoef08, Larsen&09]. Thus VDM-RT is a pure extension of the VDM++ language.

1.4   Purpose of The Document

This document is the language reference manual for all the VDM-10 dialects. The syntax of VDM language constructs is defined using grammar rules. The meaning of each language construct is explained in an informal manner and some small examples are given. The description is supposed to be suited for 'looking up' information rather than for 'sequential reading'; it is a manual rather than a tutorial. The reader is expected be familiar with the concepts of object oriented programming/design.

We will use the ASCII (also called the interchange) concrete syntax but we will display all reserved words in a special keyword font. Note that general Unicode identifiers are allowed so it is for example possible to write Japanese characters directly.

1.5   Structure of the Document

Chapter 2 presents the BNF notation used for the description of syntactic constructs. The VDM notations are described in Chapter 3 to Chapter 16. The complete syntax of the language is described in Appendix A, the lexical specification in Appendix B and the operator precedence in Appendix C. Appendix D presents a list of the differences between symbols in the mathematical syntax and the ASCII concrete syntax.

Chapter 2

Concrete Syntax Notation

Wherever the syntax for parts of the language is presented in the document it will be described in a BNF dialect. The BNF notation used employs the following special symbols:

,	the concatenate symbol
=	the define symbol
|	the definition separator symbol (alternatives)
[ ]	enclose optional syntactic items
{ }	enclose syntactic items which may occur zero or more times
' '	single quotes are used to enclose terminal symbols
meta identifier	non-terminal symbols are written in lower-case letters (possibly including spaces)
;	terminator symbol to denote the end of a rule
( )	used for grouping, e.g. "a, (b | c)" is equivalent to "a, b | a, c".
–	denotes subtraction from a set of terminal symbols (e.g. "character – ('**"**')" denotes all characters excepting the double quote character.)

Chapter 3

Data Type Definitions

As in traditional programming languages it is possible to define data types in the VDM languages and give them appropriate names. Such an equation might look like:

```
types

   Amount = nat
```

Here we have defined a data type with the name "`Amount`" and stated that the values which belong to this type are natural numbers (**nat** is one of the basic types described below). One general point about the type system of the VDM languages which is worth mentioning at this point is that equality and inequality can be used between any value. In programming languages it is often required that the operands have the same type. Because of a construct called a union type (described below) this is not the case for the VDM languages.

In this chapter we will present the syntax of data type definitions. In addition, we will show how values belonging to a type can be constructed and manipulated (by means of built-in operators). We will present the basic data types first and then we will proceed with the compound types.

3.1 Basic Data Types

In the following a number of basic types will be presented. Each of them will contain:

- Name of the construct.

- Symbol for the construct.

- Special values belonging to the data type.

- Built-in operators for values belonging to the type.

- Semantics of the built-in operators.

- Examples illustrating how the built-in operators can be used.[1]

For each of the built-in operators the name, the symbol used and the type of the operator will be given together with a description of its semantics (except that the semantics of Equality and Inequality is not described, since it follows the usual semantics). In the semantics description identifiers refer to those used in the corresponding definition of operator type, e.g. a, b, x, y etc.

The basic types are the types defined by the language with distinct values that cannot be analysed into simpler values. There are five fundamental basic types: booleans, numeric types, characters, tokens and quote types. The basic types will be explained one by one in the following.

3.1.1 The Boolean Type

In general the VDM languages allow one to specify systems in which computations may fail to terminate or to deliver a result. To deal with such potential undefinedness, the VDM languages employs a three valued logic: values may be true, false or bottom (undefined). The semantics of the VDM interpreters differs from the ISO/VDM-SL standard in that it does not have an LPF (Logic of Partial Functions) three valued logic where the order of the operands is unimportant (see [Jones90]). The **and** operator, the **or** operator and the imply operator, though, have a conditional semantics meaning that if the first operand is sufficient to determine the final result, the second operand will not be evaluated. In a sense the semantics of the logic in the VDM interpreter can still be considered to be three-valued as for ISO/VDM-SL. However, bottom values may either result in infinite computation or a run-time error in the VDM interpreter.

Name: Boolean

Symbol: `bool`

Values: `true`, `false`

Operators: Assume that a and b in the following denote arbitrary boolean expressions:

Operator	Name	Type
not b	Negation	$\texttt{bool} \rightarrow \texttt{bool}$
a **and** b	Conjunction	$\texttt{bool} * \texttt{bool} \rightarrow \texttt{bool}$
a **or** b	Disjunction	$\texttt{bool} * \texttt{bool} \rightarrow \texttt{bool}$
a => b	Implication	$\texttt{bool} * \texttt{bool} \rightarrow \texttt{bool}$
a <=> b	Biimplication	$\texttt{bool} * \texttt{bool} \rightarrow \texttt{bool}$
a = b	Equality	$\texttt{bool} * \texttt{bool} \rightarrow \texttt{bool}$
a <> b	Inequality	$\texttt{bool} * \texttt{bool} \rightarrow \texttt{bool}$

Semantics of Operators: Semantically <=> and = are equivalent when we deal with boolean values. There is a conditional semantics for **and**, **or** and =>.

---

[1]In these examples the Meta symbol '≡' will be used to indicate what the given example is equivalent to.

We denote undefined terms (e.g. applying a map with a key outside its domain) by $\perp$. The truth tables for the boolean operators are then[2]:

Negation **not** b

b	**true**	**false**	$\perp$
not b	**false**	**true**	$\perp$

Conjunction a **and** b

$a \backslash b$	**true**	**false**	$\perp$
true	**true**	**false**	$\perp$
false	**false**	**false**	**false**
$\perp$	$\perp$	$\perp$	$\perp$

Disjunction a **or** b

$a \backslash b$	**true**	**false**	$\perp$
true	**true**	**true**	**true**
false	**true**	**false**	$\perp$
$\perp$	$\perp$	$\perp$	$\perp$

Implication a **=>** b

$a \backslash b$	**true**	**false**	$\perp$
true	**true**	**false**	$\perp$
false	**true**	**true**	**true**
$\perp$	$\perp$	$\perp$	$\perp$

Biimplication a **<=>** b

$a \backslash b$	**true**	**false**	$\perp$
true	**true**	**false**	$\perp$
false	**false**	**true**	$\perp$
$\perp$	$\perp$	$\perp$	$\perp$

Examples: Let a = **true** and b = **false** then:

not a	$\equiv$	**false**
a **and** b	$\equiv$	**false**
b **and** $\perp$	$\equiv$	**false**
a **or** b	$\equiv$	**true**
a **or** $\perp$	$\equiv$	**true**
a => b	$\equiv$	**false**
b => b	$\equiv$	**true**
b => $\perp$	$\equiv$	**true**
a <=> b	$\equiv$	**false**
a = b	$\equiv$	**false**
a <> b	$\equiv$	**true**
$\perp$ **or** **not** $\perp$	$\equiv$	$\perp$
(b **and** $\perp$) **or** ($\perp$ **and false**)	$\equiv$	$\perp$

---

[2]Notice that in standard VDM-SL all these truth tables (except =>) would be symmetric.

3.1.2 The Numeric Types

There are five basic numeric types: positive naturals, naturals, integers, rationals and reals. Except for three, all the numerical operators can have mixed operands of the three types. The exceptions are integer division, modulo and the remainder operation.

The five numeric types denote a subset hierarchy where **real** is the most general type followed by **rat**[3], **int**, **nat** and **nat1**. Note that no "casting" like it is done in many programming languages is needed in the VDM languages.

Type	Values
nat1	1, 2, 3, ...
nat	0, 1, 2, ...
int	..., −2, −1, 0, 1, ...
real	..., −12.78356, ..., 0, ..., 3, ..., 1726.34, ...

This means that any number of type **int** is also automatically of type **real** but not necessarily of type **nat**. Another way to illustrate this is to say that the positive natural numbers are a subset of the natural numbers which again are a subset of the integers which again are a subset of the rational numbers which finally are a subset of the real numbers. The following table shows some numbers and their associated type:

Number	Type
3	**real, rat, int, nat, nat1**
3.0	**real, rat, int, nat, nat1**
0	**real, rat, int, nat**
−1	**real, rat, int**
3.1415	**real, rat**

Note that all numbers are necessarily of type **real** (and **rat**).

Names: real, rational, integer, natural and positive natural numbers.

Symbols: real, **rat**, **int**, **nat**, **nat1**

Values: ..., −3.89, ..., −2, ..., 0, ..., 4, ..., 1074.345, ...

Operators: Assume in the following that x and y denote numeric expressions. No assumptions are made regarding their type.

---

[3]From the VDM interpreter's point of view there is no difference between **real** and **rat** because only rational numbers can be represented in a computer.

Operator	Name	Type
$-x$	Unary minus	**real** $\rightarrow$ **real**
abs x	Absolute value	**real** $\rightarrow$ **real**
floor x	Floor	**real** $\rightarrow$ **int**
x + y	Sum	**real** * **real** $\rightarrow$ **real**
x - y	Difference	**real** * **real** $\rightarrow$ **real**
x * y	Product	**real** * **real** $\rightarrow$ **real**
x / y	Division	**real** * **real** $\rightarrow$ **real**
x **div** y	Integer division	**int** * **int** $\rightarrow$ **int**
x **rem** y	Remainder	**int** * **int** $\rightarrow$ **int**
x **mod** y	Modulus	**int** * **int** $\rightarrow$ **int**
x**y	Power	**real** * **real** $\rightarrow$ **real**
x < y	Less than	**real** * **real** $\rightarrow$ **bool**
x > y	Greater than	**real** * **real** $\rightarrow$ **bool**
x <= y	Less or equal	**real** * **real** $\rightarrow$ **bool**
x >= y	Greater or equal	**real** * **real** $\rightarrow$ **bool**
x = y	Equal	**real** * **real** $\rightarrow$ **bool**
x <> y	Not equal	**real** * **real** $\rightarrow$ **bool**

The types stated for operands are the most general types allowed. This means for instance that unary minus works for operands of all five types (**nat1**, **nat**, **int**, **rat** and **real**).

Semantics of Operators: The operators Unary minus, Sum, Difference, Product, Division, Less than, Greater than, Less or equal, Greater or equal, Equal and Not equal have the usual semantics of such operators.

Operator Name	Semantics Description
Floor	yields the greatest integer which is equal to or smaller than x.
Absolute value	yields the absolute value of x, i.e. x itself if x >= 0 and $-x$ if x < 0.
Power	yields x raised to the y'th power.

There is often confusion on how integer division, remainder and modulus work on negative numbers. In fact, there are two valid answers to -14 **div** 3: either (the intuitive) -4 as in the VDM interpreters, or -5 as in e.g. Standard ML [Paulson91]. It is therefore appropriate to explain these operations in some detail.

Integer division is defined using **floor** and real number division:

```
    x/y <  0:    x div y = -floor(abs(-x/y))
    x/y >= 0:    x div y =  floor(abs(x/y))
```

Note that the order of **floor** and **abs** on the right-hand side makes a difference, the above example would yield -5 if we changed the order. This is because **floor** always yields a smaller (or equal) integer, e.g. **floor** $(14/3)$ is 4 while **floor** $(-14/3)$ is -5.

Remainder x **rem** y and modulus x **mod** y are the same if the signs of x and y are the same, otherwise they differ and **rem** takes the sign of x and **mod** takes the sign of y. The formulas for remainder and modulus are:

```
x rem y = x - y * (x div y)
x mod y = x - y * floor(x/y)
```

Hence, -14 **rem** 3 equals -2 and -14 **mod** 3 equals 1. One can view these results by walking the real axis, starting at -14 and making jumps of 3. The remainder will be the last negative number one visits, because the first argument corresponding to x is negative, while the modulus will be the first positive number one visit, because the second argument corresponding to y is positive.

Examples: Let a = 7, b = 3.5, c = 3.1415, d = -3, e = 2 then:

```
        - a                   ≡  -7
        abs a                 ≡  7
        abs d                 ≡  3
        floor a <= a          ≡  true
        a + d                 ≡  4
        a * b                 ≡  24.5
        a / b                 ≡  2
        a div e               ≡  3
        a div d               ≡  -2
        a mod e               ≡  1
        a mod d               ≡  -2
        -a mod d              ≡  -1
        a rem e               ≡  1
        a rem d               ≡  1
        -a rem d              ≡  -1
        3**2 + 4**2 = 5**2    ≡  true
        b < c                 ≡  false
        b > c                 ≡  true
        a <= d                ≡  false
        b >= e                ≡  true
        a = e                 ≡  false
        a = 7.0               ≡  true
        c <> d                ≡  true
        abs c < 0             ≡  false
```

$$(a \; \textbf{div} \; e) \; \ast \; e \qquad \equiv \quad 6$$

3.1.3 The Character Type

The character type contains all the single character elements of the VDM character set (see Table B.1 on page 204).

Name: Char

Symbol: `char`

Values: `'a'`, `'b'`, ..., `'1'`, `'2'`, ... `'+'`, `'-'` ...

Operators: Assume that `c1` and `c2` in the following denote arbitrary characters:

Operator	Name	Type
`c1 = c2`	Equal	**char** * **char** $\rightarrow$ **bool**
`c1 <> c2`	Not equal	**char** * **char** $\rightarrow$ **bool**

Examples:

```
'a' = 'b'   ≡   false
'1' = 'c'   ≡   false
'd' <> '7'  ≡   true
'e' = 'e'   ≡   true
```

3.1.4 The Quote Type

The quote type corresponds to enumerated types in a programming language like Pascal. However, instead of writing the different quote literals between curly brackets in VDM it is done by letting a quote type consist of a single quote literal and then let them be a part of a union type.

Name: Quote

Symbol: e.g. `<QuoteLit>`

Values: `<RED>`, `<CAR>`, `<QuoteLit>`, ...

Operators: Assume that `q` and `r` in the following denote arbitrary quote values belonging to an enumerated type `T`:

Operator	Name	Type
`q = r`	Equal	T * T $\rightarrow$ **bool**
`q <> r`	Not equal	T * T $\rightarrow$ **bool**

Examples: Let `T` be the type defined as:

```
T = <France> | <Denmark> | <SouthAfrica> | <SaudiArabia>
```

If for example `a = <France>` then:

`<France> = <Denmark>`	$\equiv$	**false**
`<SaudiArabia> <> <SouthAfrica>`	$\equiv$	**true**
`a <> <France>`	$\equiv$	**false**

3.1.5   The Token Type

The token type consists of a countably infinite set of distinct values, called tokens. The only operations that can be carried out on tokens are equality and inequality. In VDM, tokens cannot be individually represented whereas they can be written with a **mk_token** around an arbitrary expression. This is a way of enabling testing of specifications which contain token types. However, in order to resemble the ISO/VDM-SL standard these token values cannot be decomposed by means of any pattern matching and they cannot be used for anything other than equality and inequality comparisons.

Name: Token

Symbol: token

Values: mk_token(5), **mk_token**({9, 3}), **mk_token**([**true**, {}]),...

Operators: Assume that `s` and `t` in the following denote arbitrary token values:

Operator	Name	Type
`s = t`	Equal	**token** * **token** $\rightarrow$ **bool**
`s <> t`	Not equal	**token** * **token** $\rightarrow$ **bool**

Examples: Let for example `s = ` **mk_token**(6) and let `t = ` **mk_token**(1) in:

`s = t`	$\equiv$	**false**
`s <> t`	$\equiv$	**true**
`s = ` **mk_token**(6)	$\equiv$	**true**

3.2   Compound Types

In the following compound types will be presented. Each of them will contain:

- The syntax for the compound type definition.

- An equation illustrating how to use the construct.

- Examples of how to construct values belonging to the type. In most cases there will also be given a forward reference to the section where the syntax of the basic constructor expressions is given.

- Built-in operators for values belonging to the type[4].

- Semantics of the built-in operators.

- Examples illustrating how the built-in operators can be used.

For each of the built-in operators the name, the symbol used and the type of the operator will be given together with a description of its semantics (except that the semantics of Equality and Inequality is not described, since it follows the usual semantics). In the semantics description identifiers refer to those used in the corresponding definition of operator type, e.g. `m`, `m1`, `s`, `s1` etc.

3.2.1 Set Types

A set is an unordered collection of values, all of the same type[5], which is treated as a whole. All sets in VDM languages are finite, i.e. they contain only a finite number of elements. The elements of a set type can be arbitrarily complex, they could for example be sets themselves.

In the following this convention will be used: `A` is an arbitrary type, `S` is a set type, `s`, `s1`, `s2` are set values, `ss` is a set of set values, `e`, `e1`, `e2` and `en` are elements from the sets, `bd1`, `bd2`, ..., `bdm` are bindings of identifiers to sets or types, and `P` is a logical predicate.

Syntax: type = set type
                    | ... ;

          set type = '**set of**', type ;

Equation: `S` = **set of** `A`

Constructors:

Set enumeration: `{e1, e2, ..., en}` constructs a set of the enumerated elements. The empty set is denoted by `{}`.

Set comprehension: `{e | bd1, bd2, ..., bdm & P}` constructs a set by evaluating the expression `e` on all the bindings for which the predicate `P` evaluates to **true**. A binding is either a set binding or a type binding[6]. A set bind `bdn` has the form `pat1, ..., patp` **in set** `s`, where `pati` is a pattern (normally simply an identifier), and `s` is a set constructed by an expression. A type binding is similar, in the sense that **in set** is replaced by a colon and `s` is replaced with a type expression.

---

[4]These operators are used in either unary or binary expressions which are given with all the operators in section 6.3.
[5]Note however that it is always possible to find a common type for two values by the use of a union type (see section 3.2.6.)
[6]Notice that type bindings cannot be executed by the VDM interpreters because in general they are not executable (see section 8 for further information about this).

The syntax and semantics for all set expressions are given in section 6.7.

Operators:

Operator	Name	Type
e **in set** s1	Membership	A * **set of** A → **bool**
e **not in set** s1	Not membership	A * **set of** A → **bool**
s1 **union** s2	Union	**set of** A * **set of** A → **set of** A
s1 **inter** s2	Intersection	**set of** A * **set of** A → **set of** A
s1 \ s2	Difference	**set of** A * **set of** A → **set of** A
s1 **subset** s2	Subset	**set of** A * **set of** A → **bool**
s1 **psubset** s2	Proper subset	**set of** A * **set of** A → **bool**
s1 = s2	Equality	**set of** A * **set of** A → **bool**
s1 <> s2	Inequality	**set of** A * **set of** A → **bool**
card s1	Cardinality	**set of** A → **nat**
dunion ss	Distributed union	**set of set of** A → **set of** A
dinter ss	Distributed intersection	**set of set of** A → **set of** A
power s1	Finite power set	**set of** A → **set of set of** A

Note that the types A, **set of** A and **set of set of** A are only meant to illustrate the structure of the type. For instance it is possible to make a union between two arbitrary sets s1 and s2 and the type of the resultant set is the union type of the two set types. Examples of this will be given in section 3.2.6.

Semantics of Operators:

Operator Name	Semantics Description
Membership	tests if e is a member of the set s1
Not membership	tests if e is not a member of the set s1
Union	yields the union of the sets s1 and s2, i.e. the set containing all the elements of both s1 and s2.
Intersection	yields the intersection of sets s1 and s2, i.e. the set containing the elements that are in both s1 and s2.
Difference	yields the set containing all the elements from s1 that are not in s2. s2 need not be a subset of s1.
Subset	tests if s1 is a subset of s2, i.e. whether all elements from s1 are also in s2. Notice that any set is a subset of itself.
Proper subset	tests if s1 is a proper subset of s2, i.e. it is a subset and s2\s1 is non-empty.
Cardinality	yields the number of elements in s1.
Distributed union	the resulting set is the union of all the elements (these are sets themselves) of ss, i.e. it contains all the elements of all the elements/sets of ss.

Operator Name	Semantics Description
Distributes intersection	the resulting set is the intersection of all the elements (these are sets themselves) of, i.e. it contains the elements that are in all the elements/sets of `ss`. `ss` must be non-empty.
Finite power set	yields the power set of `s1`, i.e. the set of all subsets of `s1`.

Examples: Let `s1` = {`<France>`,`<Denmark>`,`<SouthAfrica>`,`<SaudiArabia>`}, `s2` = {2, 4, 6, 8, 11} and `s3` = {} then:

```
<England> in set s1                          ≡  false
10 not in set s2                             ≡  true
s2 union s3                                  ≡  {2, 4, 6, 8, 11}
s1 inter s3                                  ≡  {}
(s2 \ {2,4,8,10}) union {2,4,8,10} = s2      ≡  false
s1 subset s3                                 ≡  false
s3 subset s1                                 ≡  true
s2 psubset s2                                ≡  false
s2 <> s2 union {2, 4}                        ≡  false
card s2 union {2, 4}                         ≡  5
dunion {s2, {2,4}, {4,5,6}, {0,12}}          ≡  {0,2,4,5,6,8,11,12}
dinter {s2, {2,4}, {4,5,6}}                  ≡  {4}
dunion power {2,4}                           ≡  {2,4}
dinter power {2,4}                           ≡  {}
```

3.2.2 Sequence Types

A sequence value is an ordered collection of elements of some type indexed by 1, 2, ..., n; where n is the length of the sequence. A sequence type is the type of finite sequences of elements of a type, either including the empty sequence (seq0 type) or excluding it (seq1 type). The elements of a sequence type can be arbitrarily complex; they could e.g. be sequences themselves.

In the following this convention will be used: A is an arbitrary type, L is a sequence type, S is a set type, l, l1, l2 are sequence values, ll is a sequence of sequence values. e1, e2 and en are elements in these sequences, i will be a natural number, P is a predicate and e is an arbitrary expression.

Syntax: type = seq type
                | ... ;

           seq type = seq0 type
                    | seq1 type ;

seq0 type = '**seq of**', type ;

seq1 type = '**seq1 of**', type ;

Equation: L = **seq of** A or L = **seq1 of** A

Constructors:

Sequence enumeration: [e1, e2,..., en] constructs a sequence of the enumerated elements. The empty sequence will be written as []. A text literal is a shorthand for enumerating a sequence of characters (e.g. "ifad" = ['i','f','a','d']).

Sequence comprehension: [e | id **in set** S & P] constructs a sequence by evaluating the expression e on all the bindings for which the predicate P evaluates to **true**. The expression e will use the identifier id. S is a set of numbers and id will be matched to the numbers in the normal order (the smallest number first).

The syntax and semantics of all sequence expressions are given in section 6.8.

Operators:

Operator	Name	Type
hd l	Head	**seq1 of** A $\rightarrow$ A
tl l	Tail	**seq1 of** A $\rightarrow$ **seq of** A
len l	Length	**seq of** A $\rightarrow$ **nat**
elems l	Elements	**seq of** A $\rightarrow$ **set of** A
inds l	Indexes	**seq of** A $\rightarrow$ **set of nat1**
reverse l	Reverse	**seq of** A $\rightarrow$ **seq of** A
l1 ^ l2	Concatenation	(**seq of** A) * (**seq of** A) $\rightarrow$ **seq of** A
conc ll	Distributed concatenation	**seq of seq of** A $\rightarrow$ **seq of** A
l ++ m	Sequence modification	**seq of** A * **map nat1 to** A $\rightarrow$ **seq of** A
l(i)	Sequence application	**seq of** A * **nat1** $\rightarrow$ A
l1 = l2	Equality	(**seq of** A) * (**seq of** A) $\rightarrow$ **bool**
l1 <> l2	Inequality	(**seq of** A) * (**seq of** A) $\rightarrow$ **bool**

The type A is an arbitrary type and the operands for the concatenation and distributed concatenation operators do not have to be of the same (A) type. The type of the resultant sequence will be the union type of the types of the operands. Examples will be given in section 3.2.6.

Semantics of Operators:

Operator Name	Semantics Description
Head	yields the first element of l. l must be a non-empty sequence.

Operator Name	Semantics Description
Tail	yields the subsequence of `l` where the first element is removed. `l` must be a non-empty sequence.
Length	yields the length of `l`.
Elements	yields the set containing all the elements of `l`.
Indexes	yields the set of indexes of `l`, i.e. the set $\{1,...,\text{len } l\}$.
Reverse	yields a new sequence where the order of the elements have been reversed.
Concatenation	yields the concatenation of `l1` and `l2`, i.e. the sequence consisting of the elements of `l1` followed by those of `l2`, in order.
Distributed concatenation	yields the sequence where the elements (these are sequences themselves) of `ll` are concatenated: the first and the second, and then the third, etc.
Sequence modification	the elements of `l` whose indexes are in the domain of `m` are modified to the range value that the index maps into. `dom m` must be a subset of `inds l`
Sequence application	yields the element of index from `l`. `i` must be in the indexes of `l`.

Examples: Let `l1 = [3,1,4,1,5,9,2]`, `l2 = [2,7,1,8]`,
   `l3 = [<England>, <Rumania>, <Colombia>, <Tunisia>]` then:

```
len l1                                   ≡  7
hd (l1^l2)                               ≡  3
tl (l1^l2)                               ≡  [1,4,1,5,9,2,2,7,1,8]
l3(len l3)                               ≡  <Tunisia>
"England"(2)                             ≡  'n'
reverse l1                               ≡  [2,9,5,1,4,1,3]
conc [l1,l2] = l1^l2                     ≡  true
conc [l1,l1,l2] = l1^l2                  ≡  false
elems l3                                 ≡  { <England>,
                                              <Rumania>,
                                              <Colombia>,<Tunisia>}
(elems l1) inter (elems l2)              ≡  {1,2}
inds l1                                  ≡  {1,2,3,4,5,6,7}
(inds l1) inter (inds l2)                ≡  {1,2,3,4}
l3 ++ {2 |-> <Germany>,4 |-> <Nigeria>}  ≡  [ <England>,
                                              <Germany>,
                                              <Colombia>,
                                              <Nigeria>]
```

17

3.2.3 Map Types

A map type from a type A to a type B is a type that associates with each element of A (or a subset of A) an element of B. A map value can be thought of as an unordered collection of pairs. The first element in each pair is called a key, because it can be used as a key to get the second element (called the information part) in that pair. All key elements in a map must therefore be unique. The set of all key elements is called the domain of the map, while the set of all information values is called the range of the map. All maps in VDM languages are finite. The domain and range elements of a map type can be arbitrarily complex, they could e.g. be maps themselves.

A special kind of map is the injective map. An injective map is one for which no element of the range is associated with more than one element of the domain. For an injective map it is possible to invert the map.

In the following this convention will be used: m, m1 and m2 are maps from an arbitrary type A to another arbitrary type B, ms is a set of map values, a, a1, a2 and an are elements from A while b, b1, b2 and bn are elements from B and P is a logic predicate. e1 and e2 are arbitrary expressions and s is an arbitrary set.

Syntax: type = map type
                | … ;

map type = general map type
                | injective map type ;

general map type = '**map**', type, '**to**', type ;

injective map type = '**inmap**', type, '**to**', type ;

Equation: M = **map** A **to** B or M = **inmap** A **to** B

Constructors:

Map enumeration: {a1 |-> b1, a2 |-> b2, ..., an |-> bn} constructs a mapping of the enumerated maplets. The empty map will be written as {|->}.

Map comprehension: {ed |-> er | bd1, ..., bdn & P} constructs a mapping by evaluating the expressions ed and er on all the possible bindings for which the predicate P evaluates to **true**. bd1, ..., bdn are bindings of free identifiers from the expressions ed and er to sets or types.

The syntax and semantics of all map expressions are given in section 6.9.

Operators:

Operator	Name	Type
dom m	Domain	(**map** A **to** B) → **set of** A
rng m	Range	(**map** A **to** B) → **set of** B
m1 **munion** m2	Merge	(**map** A **to** B) * (**map** A **to** B) → **map** A **to** B
m1 ++ m2	Override	(**map** A **to** B) * (**map** A **to** B) → **map** A **to** B
merge ms	Distributed merge	**set of** (**map** A **to** B) → **map** A **to** B
s <: m	Domain restrict to	(**set of** A) * (**map** A **to** B) → **map** A **to** B
s <-: m	Domain restrict by	(**set of** A) * (**map** A **to** B) → **map** A **to** B
m :> s	Range restrict to	(**map** A **to** B) * (**set of** B) → **map** A **to** B
m :-> s	Range restrict by	(**map** A **to** B) * (**set of** B) → **map** A **to** B
m(d)	Map apply	(**map** A **to** B) * A → B
m1 **comp** m2	Map composition	(**map** B **to** C) * (**map** A **to** B) → **map** A **to** C
m ** n	Map iteration	(**map** A **to** A) * **nat** → **map** A **to** A
m1 = m2	Equality	(**map** A **to** B) * (**map** A **to** B) → **bool**
m1 <> m2	Inequality	(**map** A **to** B) * (**map** A **to** B) → **bool**
inverse m	Map inverse	**inmap** A **to** B → **inmap** B **to** A

Semantics of Operators: Two maps m1 and m2 are compatible if any common element of **dom**m1 and **dom**m2 is mapped to the same value by both maps.

Operator Name	Semantics Description
Domain	yields the domain (the set of keys) of m.
Range	yields the range (the set of information values) of m.
Merge	yields a map combined by m1 and m2 such that the resulting map maps the elements of **dom**m1 as does m1, and the elements of domm2 as does m2. The two maps must be compatible.
Override	overrides and merges m1 with m2, i.e. it is like a merge except that m1 and m2 need not be compatible; any common elements are mapped as by m2 (so m2 overrides m1).
Distributed merge	yields the map that is constructed by merging all the maps in ms. The maps in ms must be compatible.
Domain restricted to	creates the map consisting of the elements in m whose key is in s. s need not be a subset of **dom**m.
Domain restricted by	creates the map consisting of the elements in m whose key is not in s. s need not be a subset of **dom**m.
Range restricted to	creates the map consisting of the elements in m whose information value is in s. s need not be a subset of **rng**m.

Operator Name	Semantics Description
Range restricted by	creates the map consisting of the elements in m whose information value is not in s. s need not be a subset of **rng**m.
Map apply	yields the information value whose key is d. d must be in the domain of m.
Map composition	yields the the map that is created by composing m2 elements with m1 elements. The resulting map is a map with the same domain as m2. The information value corresponding to a key is the one found by first applying m2 to the key and then applying m1 to the result. **rng**m2 must be a subset of **dom**m1.
Map iteration	yields the map where m is composed with itself n times. n=0 yields the identity map where each element of **dom**m is map into itself; n=1 yields m itself. For n>1, the range of m must be a subset of **dom**m.
Map inverse	yields the inverse map of m. m must be a 1-to-1 mapping.

Examples: Let

```
m1 = { <France> |-> 9, <Denmark> |-> 4,
       <SouthAfrica> |-> 2, <SaudiArabia> |-> 1},
m2 = { 1 |-> 2, 2 |-> 3, 3 |-> 4, 4 |-> 1 },
Europe = { <France>, <England>, <Denmark>, <Spain> }
```

then:

dom m1 ≡ {<France>, <Denmark>, <SouthAfrica>, <SaudiArabia>}

rng m1 ≡ {1,2,4,9}

m1 **munion** {<England> |-> 3} ≡ {<France> |-> 9, <Denmark> |-> 4, <England> |-> 3, <SaudiArabia> |-> 1, <SouthAfrica> |-> 2}

```
m1 ++ {<France>  |-> 8,          ≡  {<France>  |-> 8,
        <England> |-> 4}             <Denmark> |-> 4,
                                     <SouthAfrica> |-> 2,
                                     <SaudiArabia> |-> 1,
                                     <England> |-> 4}


merge{ {<France>  |-> 9,         ≡  {<France>  |-> 9,
        <Spain>   |-> 4}            <England> |-> 3,
       {<France>  |-> 9,            <Spain>  |-> 4,
        <England> |-> 3,            <UnitedStates> |-> 1}
        <UnitedStates> |-> 1}}


Europe <:  m1                    ≡  {<France>  |-> 9,
                                    <Denmark> |-> 4}


Europe <-:  m1                   ≡  {<SouthAfrica> |-> 2,
                                    <SaudiArabia> |-> 1}


m1 :> {2,...,10}                 ≡  {<France>  |-> 9,
                                    <Denmark> |-> 4,
                                    <SouthAfrica> |-> 2}


m1 :-> {2,...,10}                ≡  {<SaudiArabia> |-> 1}


m1 comp ({"France" |-> <France>}) ≡  {"France" |-> 9}


m2 ** 3                          ≡  {1 |-> 4, 2 |-> 1,
                                     3 |-> 2, 4 |-> 3 }


inverse m2                       ≡  {2 |-> 1, 3 |-> 2,
                                     4 |-> 3, 1 |-> 4 }


m2 comp (inverse m2)             ≡  {1 |-> 1, 2 |-> 2,
                                     3 |-> 3, 4 |-> 4 }
```

3.2.4 Product Types

The values of a product type are called tuples. A tuple is a fixed length list where the i'th element of the tuple must belong to the i'th element of the product type.

Syntax: type = product type
         | ... ;

product type = type, '*', type, { '*', type } ;

A product type consists of at least two subtypes.

Equation: `T = A1 * A2 * ... * An`

Constructors: The tuple constructor: **mk_**`(a1, a2, ..., an)`

The syntax and semantics for the tuple constructor are given in section 6.10.

Operators:

Operator	Name	Type
`t.#n`	Select	T * **nat** $\rightarrow$ Ti
`t1 = t2`	Equality	T * T $\rightarrow$ **bool**
`t1 <> t2`	Inequality	T * T $\rightarrow$ **bool**

The only operators working on tuples are component select, equality and inequality. Tuple components may be accessed using the select operator or by matching against a tuple pattern. Details of the semantics of the tuple select operator and an example of its use are given in section 6.12.

Examples: Let a = **mk_**`(1, 4, 8)`, b = **mk_**`(2, 4, 8)` then:

```
a = b          ≡  false
a <> b         ≡  true
a = mk_(2,4)   ≡  false
```

3.2.5 Composite Types

Composite types correspond to record types in programming languages. Thus, elements of this type are somewhat similar to the tuples described in the section about product types above. The difference between the record type and the product type is that the different components of a record can be directly selected by means of corresponding selector functions. In addition records are tagged with an identifier which must be used when manipulating the record. The only way to tag a type is by defining it as a record. It is therefore common usage to define records with only one field in order to give it a tag. This is another difference to tuples as a tuple must have at least two entries whereas records can be empty.

In VDM languages, **is_** is a reserved prefix for names and it is used in an *is expression*. This is a built-in operator which is used to determine which record type a record value belongs to. It is often used to discriminate between the subtypes of a union type and will therefore be explained further in section 3.2.6. In addition to record types the **is_** operator can also determine if a value is of one of the basic types.

In the following this convention will be used: A is a record type, `A1, ..., Am` are arbitrary types, `r`, `r1`, and `r2` are record values, `i1, ..., im` are selectors from the `r` record value, `e1, ..., em` are arbitrary expressions.

Syntax: type = composite type
              | … ;

        composite type = '**compose**', identifier, '**of**', field list, '**end**' ;

        field list = { field } ;

        field = [identifier, ':'], type
              | [identifier, ':-'], type ;

or the shorthand notation

        composite type = identifier, '::', field list ;

where identifier denotes both the type name and the tag name.

Equation:

```
     A :: selfirst : A1
          selsec   : A2
```

or

```
     A :: selfirst :  A1
          selsec   :- A2
```

or

```
     A :: A1 A2
```

In the second notation, an *equality abstraction* field is used for the second field `selsec`. The minus indicates that such a field is ignored when comparing records using the equality operator. In the last notation the fields of `A` can only be accessed by pattern matching (like it is done for tuples) as the fields have not been named.

In the last notation the fields of `A` can only be accessed by pattern matching (as is done for tuples) since the fields have not been named.

The shorthand notation `::` used in the two previous examples where the tag name equals the type name, is the notation most used. The more general **compose** notation is typically used if a composite type has to be specified directly as a component of a more complex type:

```
     T = map S to compose A of A1 A2 end
```

It should be noted however that composite types can only be used in type definitions, and not e.g. in signatures to functions or operations.

Typically composite types are used as alternatives in a union type definition (see section 3.2.6) such as:

```
MasterA = A | B | ...
```

where `A` and `B` are defined as composite types themselves. In this situation the **is_** predicate can be used to distingush the alternatives.

Constructors: The record constructor: **mk_**A(a, b) where a belongs to the type A1 and b belongs to the type A2.

The syntax and semantics for all record expressions are given in section 6.11.

Operators:

Operator	Name	Type
`r.i`	Field select	A * Id $\rightarrow$ Ai
`r1 = r2`	Equality	A * A $\rightarrow$ **bool**
`r1 <> r2`	Inequality	A * A $\rightarrow$ **bool**
is_A(r1)	Is	Id * MasterA $\rightarrow$ **bool**

Semantics of Operators:

Operator Name	Semantics Description
Field select	yields the value of the field with fieldname `i` in the record value `r`. `r` must have a field with name `i`.

Examples: Let `Score` be defined as

```
Score :: team   : Team
         won    : nat
         drawn  : nat
         lost   : nat
         points : nat;
Team = <Brazil> | <France> | ...
```

and let

```
sc1 = mk_Score (<France>, 3, 0, 0, 9),
sc2 = mk_Score (<Denmark>, 1, 1, 1, 4),
sc3 = mk_Score (<SouthAfrica>, 0, 2, 1, 2) and
sc4 = mk_Score (<SaudiArabia>, 0, 1, 2, 1).
```

Then

```
sc1.team              ≡   <France>
sc4.points            ≡   1
sc2.points > sc3.points ≡  true
is_Score(sc4)         ≡   true
is_bool(sc3)          ≡   false
is_int(sc1.won)       ≡   true
sc4 = sc1             ≡   false
sc4 <> sc2            ≡   true
```

The equality abstraction field, written using ':-' instead of ':', may be useful, for example, when working with lower level models of an abstract syntax of a programming language. For example, one may wish to add a position information field to a type of identifiers without affecting the true identity of identifiers:

```
Id :: name :  seq of char
      pos  :- nat
```

The effect of this will be that the `pos` field is ignored in equality comparisons, e.g. the following would evaluate to true:

```
mk_Id("x",7) = mk_Id("x",9)
```

In particular this can be useful when looking up in an environment which is typically modelled as a map of the following form:

```
Env = map Id to Val
```

Such a map will contain at most one index for a specific identifier, and a map lookup will be independent of the `pos` field.

Moreover, the equality abstraction field will affect set expressions. For example,

```
{mk_Id("x",7),mk_Id("y",8),mk_Id("x",9)}
```

will be equal to

```
{mk_Id("x",?),mk_Id("y",8)}
```

where the question mark stands for 7 or 9.

Finally, note that for equality abstraction fields valid patterns are limited to don't care and identifier patterns. Since equality abstraction fields are ignored when comparing two values, it does not make sense to use more complicated patterns.

3.2.6 Union and Optional Types

The union type corresponds to a set-theoretic union, i.e. the type defined by means of a union type will contain all the elements from each of the components of the union type. It is possible to use types that are not disjoint in the union type, even though such usage would be bad practice. However, the union type is normally used when something belongs to one type from a set of possible types. The types which constitute the union type are often composite types. This makes it possible, using the **is_** operator, to decide which of these types a given value of the union type belongs to.

The optional type [T] is a kind of shorthand for a union type T | **nil**, where **nil** is used to denote the absence of a value. However, it is not possible to use the set {**nil**} as a type so the only types **nil** will belong to will be optional types.

Syntax: type = union type
              | optional type
              | … ;

        union type = type, '|', type, { '|', type } ;

        optional type = '[', type, ']' ;

Equation: B = A1 | A2 | ... | An

Constructors: None.

Operators:

Operator	Name	Type
t1 = t2	Equality	A * A → **bool**
t1 <> t2	Inequality	A * A → **bool**

Examples: In this example Expr is a union type whereas Const, Var, Infix and Cond are composite types defined using the shorthand :: notation.

```
Expr  = Const | Var | Infix | Cond;
Const :: nat | bool;
Var   :: id:Id
         tp: [<Bool> | <Nat>];
Infix :: Expr * Op * Expr;
Cond  :: test : Expr
```

```
            cons : Expr
            altn : Expr
```

and let `expr` = **mk_**Cond(**mk_**Var("b",<Bool>),**mk_**Const(3),
mk_ Var("v",**nil**)) then:

is_ Cond(expr)	≡	**true**
is_ Const(expr.cons)	≡	**true**
is_ Var(expr.altn)	≡	**true**
is_ Infix(expr.test)	≡	**false**

Using union types we can extend the use of previously defined operators. For instance, interpreting = as a test over **bool** | **nat** we have

1 = **false** ≡ **false**

Similarly we can take use union types for taking unions of sets and concatenating sequences:

{1,2} **union** {**false**,**true**}	≡	{1,2, **false**,**true**}
['a','b']^[<c>,<d>]	≡	['a','b', <c>,<d>]

In the set union, we take the union over sets of type **nat** | **bool**; for the sequence concatenation we are manipulating sequences of type **char** | <c> | <d>.

3.2.7 The Object Reference Type (VDM++ and VDM-RT)

The object reference type has been added as part of the standard VDM-SL types. Therefore there is no direct way of restricting the use of object reference types (and thus of objects) in a way that conforms to pure object oriented principles; no additional structuring mechanisms than classes are foreseen. From these principles it follows that the use of an object reference type in combination with a type constructor (record, map, set, etc.) should be treated with caution.

A value of the object reference type can be regarded as a *reference* to an object. If, for example, an instance variable (see section 10.1) is defined to be of this type, this makes the class in which that instance variable is defined, a 'client' of the class in the object reference type; a *clientship relation* is established between the two classes.

An object reference type is denoted by a class name. The class name in the object reference type must be the name of a class defined in the specification.

The only operators defined for values of this type is the test for equality ('=') and inequality ('<>'). Equality is based on references rather than values. That is, if `o1` and `o2` are two distinct objects which happen to have the same contents, `o1 = o2` will yield false.

Constructors Object references are constructed using the new expression (see section 6.13).

Operators

Operator	Name	Type
t1 = t2	Equality	A * A → **bool**
t1 <> t2	Inequality	A * A → **bool**

Examples An example of the use of object references is in the definition of the class of binary trees:

```
class Tree

types

  protected tree = <Empty> | node;

  public node :: lt   : Tree
                 nval : int
                 rt   : Tree

instance variables

  protected root: tree := <Empty>;
end Tree
```

Here we define the type of nodes, which consist of a node value, and references to left and right tree objects. Details of access specifiers may be found in section 13.3.3.

3.2.8 Function Types

In the VDM languages function types can also be used in type definitions. A function type from a type A (actually a list of types as a tuple type) to a type B is a type that associates with each element of A an element of B. A function value can be thought of as a function in a programming language which has no side-effects (i.e. it does not use any global variables).

Such usage can be considered advanced in the sense that functions are used as values (thus this section may be skipped during the first reading). Function values may be created by lambda expressions (see below), or by function definitions, which are described in section 5. Function values can be of higher order in the sense that they can take functions as arguments or return functions as results. In this way functions can be Curried such that a new function is returned when the first set of parameters are supplied (see the examples below).

Syntax: type = partial function type
                   | … ;

          function type = partial function type
                          | total function type ;

          partial function type = discretionary type, '->', type ;

> total function type = discretionary type, '+>', type ;
>
> discretionary type = type | '(',')' ;

Equation: `F = A +> B`[7] or `F = A -> B`

Constructors: In addition to the traditional function definitions the only way to construct functions is by the lambda expression: **lambda** `pat1 : T1, ..., patn : Tn &` `body` where the `patj` are patterns, the `Tj` are type expressions, and `body` is the body expression which may use the pattern identifiers from all the patterns.

The syntax and semantics for the lambda expression are given in section 6.16.

Operators:

Operator	Name	Type
`f(a1,...,an)`	Function apply	$A1 * \cdots * An \rightarrow B$
`f1` **comp** `f2`	Function composition	$(B \rightarrow C) * (A \rightarrow B) \rightarrow (A \rightarrow C)$
`f ** n`	Function iteration	$(A \rightarrow A) *$ **nat** $\rightarrow (A \rightarrow A)$
`t1 = t2`	Equality	$A * A \rightarrow$ **bool**
`t1 <> t2`	Inequality	$A * A \rightarrow$ **bool**

Note that equality and inequality between type values should be used with great care. In the VDM languages this corresponds to the mathematical equality (and inequality) which is not computable for infinite values like general functions. Thus, in the VDM interpreters the equality is on the abstract syntax of the function value (see `inc1` and `inc2` below).

Semantics of Operators:

Operator Name	Semantics Description
Function apply	yields the result of applying the function `f` to the values of a_j. See the definition of apply expressions in Section 6.12.
Function composition	it yields the function equivalent to applying first `f2` and then applying `f1` to the result. `f1`, but not `f2` may be Curried.
Function iteration	yields the funciton equivalent to applying `f` n times. `n=0` yields the identity function which just returns the value of its parameter; `n=1` yields the function itself. For `n>1`, the result of `f` must be contained in its parameter type.

---

[7]Note that the total function arrow can only be used in signatures of totally defined functions and thus not in a type definition.

Examples: Let the following function values be defined:

```
f1 = lambda x : nat & lambda y : nat & x + y
f2 = lambda x : nat & x + 2
inc1 = lambda x : nat & x + 1
inc2 = lambda y : nat & y + 1
```

then the following holds:

```
f1(5)        ≡ lambda y :nat & 5 + y
f2(4)        ≡ 6
f1 comp f2   ≡ lambda x :nat & lambda y :nat & (x + 2) + y
f2 ** 4      ≡ lambda x :nat & x + 8
inc1 = inc2  ≡ false
```

Notice that the equality test does not yield the expected result with respect to the semantics of the VDM languages. Thus, one should be **very** careful with the usage of equality for infinite values like functions.

3.3 Invariants

If the data types specified by means of equations as described above contain values which should not be allowed, then it is possible to restrict the values in a type by means of an invariant. The result is that the type is restricted to a subset of its original values. Thus, by means of a predicate the acceptable values of the defined type are limited to those where this expression is true.

The general scheme for using invariants looks like this:

```
Id = Type
inv pat == expr
```

where `pat` is a pattern matching the values belonging to the type `Id`, and `expr` is a truth-valued expression, involving some or all of the identifiers from the pattern `pat`.

If an invariant is defined, a new (total) function is implicitly created with the signature:

```
inv_Id : Type +> bool
```

This function can be used within other invariant, function or operation definitions.

For instance, recall the record type `Score` defined on page 24. We can ensure that the number of points awarded is consistent with the number of games won and drawn using an invariant:

```
Score :: team   : Team
         won    : nat
```

```
          drawn  : nat
          lost   : nat
          points : nat
 inv sc == sc.points = 3 * sc.won + sc.drawn;
```

The invariant function implicitly created for this type is:

```
  inv_Score : Score +> bool
  inv_Score (sc) ==
    sc.points = 3 * sc.won + sc.drawn;
```

Chapter 4

Algorithm Definitions

In the VDM languages algorithms can be defined by both functions and operations. However, they do not directly correspond to functions in traditional programming languages. What separates functions from operations in the VDM languages is the use of local and global variables. Operations can manipulate both the global variables and any local variables. Both local and global variables will be described later. Functions are pure in the sense that they cannot access global variables and they are not allowed to define local variables. Thus, functions are purely applicative while operations are imperative.

Functions and operations can be defined both explicitly (by means of an explicit algorithm definition) or implicitly (by means of a pre-condition and/or a post condition). An explicit algorithm definition for a function is called an expression while for an operation it is called a statement. A pre-condition is a truth-valued expression which specifies what must hold before the function/operation is evaluated. A pre-condition can only refer to parameter values and global variables (if it is an operation). A post-condition is also a truth valued expression which specifies what must hold after the function/operation is evaluated. A post-condition can refer to the result identifier, the parameter values, the current values of global variables and the old values of global variables. The old values of global variables are the values of the variables as they were before the operation was evaluated. Only operations can refer to the old values of global variables in a post-condition as functions are not allowed access to the global variables in any way.

However, in order to be able to execute both functions and operations by the VDM interpreters they must be defined explicitly[1]. In the VDM languages it is also possible for explicit function and operation definitions to specify an additional pre- and a post-condition. In the post-condition of explicit function and operation definitions the result value must be referred to by the reserved word **RESULT**.

---

[1]Implicitly specified functions and operations cannot in general be executed because their post-condition does not need to directly relate the output to the input. Often it is done by specifying the properties the output must satisfy.

Chapter 5

Function Definitions

In the VDM languages we can define first order and higher order functions. A higher order function is either a Curried function (a function that returns a function as result), or a function that takes functions as arguments. Furthermore, both first order and higher order functions can be polymorphic. In VDM++ and VDM-RT functions are automatically available in a static form (i.e. without having an instance of the defining class). Thus there is no need to use the **static** keyword that can be used for operations in VDM++ and VDM-RT. In general, the syntax for the definition of a function is:

> function definitions = '**functions**', [access function definition],
>   { ';' }, access function definition function definition, [';'] ;

> access function definition = [access], function definition ;

> access = '**public**'
> | '**private**'
> | '**protected**' ;

> function definition = explicit function definition
> | implicit function definition
> | extended explicit function definition ;

> explicit function definition = identifier,
>   [type variable list], ':', function type,
>   identifier, parameters list, '==',
>   function body,
>   ['**pre**', expression],
>   ['**post**', expression],
>   ['**measure**', name] ;

> implicit function definition = identifier, [type variable list],
>   parameter types, identifier type pair list,
>   ['**pre**', expression],
>   '**post**', expression ;

extended explicit function definition = identifier, [type variable list],
 parameter types,
 identifier type pair list,
 '==', function body,
 ['**pre**', expression],
 ['**post**', expression] ;

type variable list = '[', type variable identifier,
 { ',', type variable identifier }, ']' ;

identifier type pair list = identifier, ':', type,
 { ',', identifier, ':', type } ;

parameter types = '(', [pattern type pair list], ')' ;

pattern type pair list = pattern list, ':', type,
 { ',', pattern list, ':', type } ;

function type = partial function type
 | total function type ;

partial function type = discretionary type, '->', type ;

total function type = discretionary type, '+>', type ;

discretionary type = type | '(',')' ;

parameters = '(', [pattern list], ')' ;

pattern list = pattern, { ',', pattern } ;

function body = expression
 | '**is not yet specified**'
 | '**is subclass responsibility**' ;

Here **is not yet specified** may be used as the function body during development of a model; whereas the **is subclass responsibility** indicates that implementation of this body must be undertaken by any subclasses so that can only be used in VDM++ and VDM-RT.

A simple example of an explicit function definition is the function map_inter which takes two compatible maps over natural numbers and returns those maplets common to both

```
map_inter: (map nat to nat) * (map nat to nat) -> map nat to nat
map_inter (m1,m2) ==
  (dom m1 inter dom m2) <: m1
pre forall d in set dom m1 inter dom m2 & m1(d) = m2(d)
```

Note that we could also use the optional post condition to allow assertions about the result of the function:

```
map_inter: (map nat to nat) * (map nat to nat) -> map nat to nat
map_inter (m1,m2) ==
  (dom m1 inter dom m2) <: m1
pre forall d in set dom m1 inter dom m2 & m1(d) = m2(d)
post dom RESULT = dom m1 inter dom m2
```

The same function can also be defined implicitly:

```
map_inter2 (m1,m2: map nat to nat) m: map nat to nat
pre forall d in set dom m1 inter dom m2 & m1(d) = m2(d)
post dom m = dom m1 inter dom m2 and
     forall d in set dom m & m(d) = m1(d);
```

A simple example of an extended explicit function definition (non-standard) is the function map_disj which takes a pair of compatible maps over natural numbers and returns the map consisting of those maplets unique to one or other of the given maps:

```
map_disj (m1:map nat to nat,m2:map nat to nat)
          res : map nat to nat ==
  (dom m1 inter dom m2) <-: m1 munion
  (dom m1 inter dom m2) <-: m2
pre forall d in set dom m1 inter dom m2 & m1(d) = m2(d)
post dom res = (dom m1 union dom m2) ++ (dom m1 inter dom m2)
     and
     forall d in set dom res & res(d) = m1(d) or res(d) = m2(d)
```

(Note here that an attempt to interpret the post-condition could potentially result in a run-time error since m1(d) and m2(d) need not both be defined simultaneously.)

The functions map_inter and map_disj can be evaluated by the VDM interpreters, but the implicit function map_inter2 cannot be evaluated. However, in all three cases the pre- and post-conditions can be used in other functions; for instance from the definition of map_inter2 we get functions **pre_**map_inter2 and **post_**map_inter2 with the following signatures:

```
  pre_map_inter2 : (map nat to nat) * (map nat to nat) +> bool
  post_map_inter2 : (map nat to nat) * (map nat to nat) *
                    (map nat to nat) +> bool
```

These kinds of functions are automatically created by the VDM interpreters and they can be used in other definitions (this technique is called quoting). In general, for a function f with signature

```
f : T1 * ... * Tn -> Tr
```

defining a pre-condition for the function causes creation of a function **pre_**f with signature

```
pre_f : T1 * ... * Tn +> bool
```

and defining a post-condition for the function causes creation of a function **post_**f with signature

```
post_f : T1 * ... * Tn * Tr +> bool
```

Functions can also be defined using recursion (i.e. by calling themselves). When this is done one is recommended to add a '**measure**' function that can be used in the proof obligations generated from the model such that termination proofs can be carried out. The **measure** function shall take the same type of parameters as the recursive function itself and it yield a natural number. A simple example here could be the traditional factorial function defined as:

```
functions

fac: nat +> nat
fac(n) ==
  if n = 0
  then 1
  else n * fac(n - 1)
measure id
```

where id would be defined as:

```
id: nat +> nat
id(n) == n
```

Here the proof obligation will become:

```
forall n:nat &
  (not (n = 0) =>
    id(n) > id((n - 1)))
```

This proof obligation will ensure that the recursive function will terminate and thus sooner or later reach the base case.

5.1 Polymorphic Functions

Functions can also be polymorphic in VDM. This means that we can create generic functions that can be used on values of several different types. For this purpose type parameters (or type variables which are written like normal identifiers prefixed with a @ sign) are used. Consider the polymorphic function to create an empty bag:[1]

```
empty_bag[@elem] : () +> (map @elem to nat1)
empty_bag() ==
  { |-> }
```

Before we can use the above function, we have to instantiate the function empty_bag with a type, for example integers (see also section 6.12):

```
emptyInt = empty_bag[int]
```

Now we can use the function emptyInt to create a new bag to store integers. More examples of polymorphic functions are:

```
num_bag[@elem] : @elem * (map @elem to nat1) +> nat
num_bag(e, m) ==
  if e in set dom m
  then m(e)
  else 0;

plus_bag[@elem] : @elem * (map @elem to nat1) +>
                  (map @elem to nat1)
plus_bag(e, m) ==
  m ++ { e |-> num_bag[@elem](e, m) + 1 }
```

If pre- and or post-conditions are defined for polymorphic functions, the corresponding predicate functions are also polymorphic. For instance if num_bag was defined as

```
  num_bag[@elem] : @elem * (map @elem to nat1) +> nat
```

---

[1]The examples for polymorphic functions are taken from [Dawes91]. Bags are modelled as maps from the elements to their multiplicity in the bag. The multiplicity is at least 1, i.e. a non-element is not part of the map, rather than being mapped to 0.

```
num_bag(e, m) ==
  m(e)
pre e in set dom m
```

then the pre-condition function would be

```
pre_num_bag[@elem] :@elem * (map @elem to nat1) +> bool
```

In case functions are defined polymorphic a **measure** should also be used.

Finally, VDM makes no assumptions about the type bound to a type parameter. This binding must be made explicitly, either by making an assertion as part of the pre-condition of the function, or by making a test before the part of the function body that restricts the type parameter. For an example, consider the function given below:

```
public f[@T]: seq of @T -> seq of @T
f(s) ==
  if is_(s, seq of real) then
        [s(i) + 1 | i in set inds s]
  else
        reverse s;
```

Here the type parameter `@T` denotes the element type of the sequence `s` passed to the function. Note how the body of the function assumes the type parameter to be defined for the + operator. This only works since the **if**-clause explicitly states that the parameter `s` is of type **seq of real**. Without the **if**-clause use of the + operator would fail due to `@T` not being of a numeric type. Alternatively, the assumption could be stated using a pre-condition. This puts a restriction on the type parameter that applies to the entire body of the function, though with simple examples like this the effect is the same as explicitly declaring the type of the parameter in the function signature.

5.2   Higher Order Functions

Functions are allowed to receive other functions as arguments. A simple example of this is the function `nat_filter` which takes a sequence of natural numbers, and a predicate, and returns the subsequence that satisfies this predicate:

```
nat_filter : (nat -> bool) * seq of nat -> seq of nat
nat_filter (p,ns) ==
  [ns(i) | i in set inds ns & p(ns(i))];
```

Then `nat_filter` (**lambda** x:**nat** & x **mod** 2 = 0, [1,2,3,4,5]) $\equiv$ [2,4]. In fact, this algorithm is not specific to natural numbers, so we may define a polymorphic version of this function:

```
filter[@elem]: (@elem -> bool) * seq of @elem -> seq of @elem
filter (p,l) ==
  [l(i) | i in set inds l & p(l(i))];
```

so `filter`[**real**](**lambda** x:**real** & **floor** x = x, [2.3,0.7,-2.1,3]) $\equiv$ [3].

Functions may also return functions as results. An example of this is the function `fmap`:

```
fmap[@elem]: (@elem -> @elem) -> seq of @elem -> seq of @elem
fmap (f)(l) ==
  if l = []
  then []
  else [f(hd l)] ^ (fmap[@elem] (f)(tl l));
```

So `fmap`[**nat**](**lambda** x:**nat** & x * x)([1,2,3,4,5]) $\equiv$ [1,4,9,16,25].

Since the `fmap` function is recursive, it ought to have a **measure** function defined. In the case of curried functions, the measure function's parameters are the same as a de-curried version of the recursive function's parameters. For the `fmap` example, this would be:

```
m[@elem]: (@elem -> @elem) * seq of @elem -> nat
m(-, l) == len l;
```

Note that the measure function is also polymorphic, and must have the same type parameters as the function it measures. The proof obligation will also be polymorphic:

```
(forall f:(@elem -> @elem), l:seq of (@elem) &
  (not (l = []) =>
    m[@elem](f, l) > m[@elem](f, (tl l))))
```

Chapter 6

Expressions

In this chapter we will describe the different kinds of expressions one by one. Each of them will be described by means of:

- A syntax description in BNF.

- An informal semantics description.

- An example illustrating its usage.

6.1   Let Expressions

Syntax: expression = let expression
                         | let be expression
                         | … ;

let expression = '**let**', local definition { ',', local definition },
                      '**in**', expression ;

let be expression = '**let**', multiple bind, ['**be**', '**st**', expression], '**in**',
                         expression ;

local definition = value definition
                         | function definition ;

value definition = pattern, [':', type], '=', expression ;

where the "function definition" component is described in section 5.

Semantics: A simple *let expression* has the form:

```
    let p1 = e1, ..., pn = en in e
```

43

where p1, ..., pn are patterns, e1, ..., en are expressions which match the corresponding pattern pi, and e is an expression, of any type, involving the pattern identifiers of p1, ..., pn. It denotes the value of the expression e in the context in which the patterns p1, ..., pn are matched against the corresponding expressions e1, ..., en.

More advanced let expressions can also be made by using local function definitions. The semantics of doing so is simply that the scope of such locally defined functions is restricted to the body of the let expression.

In standard VDM-SL the collection of definitions may be mutually recursive. However, in the VDM languages this is not supported by the VDM interpreters. Furthermore, the definitions must be ordered such that all constructs are defined before they are used.

A *let-be-such-that expression* has the form:

```
let mb be st e1 in e2
```

where mb is a multi-binding of one or more patterns (mostly just one pattern) to a set value (or a type), e1 is a boolean expression, and e2 is an expression, of any type, involving the pattern identifiers of the patterns from mb. The **be st** e1 part is optional. The expression denotes the value of the expression e2 in the context in which all the patterns from mb has been matched against either an element in the set from mb or against a value from the type in mb[1]. If the **st** e1 expression is present, only such bindings where e1 evaluates to true in the matching context are used.

Examples: *Let expressions* are useful for improving readability especially by contracting complicated expressions used more than once. For instance, we can improve the function map_disj from page 37:

```
map_disj : (map nat to nat) * (map nat to nat) ->
             map nat to nat
map_disj (m1,m2) ==
  let inter_dom = dom m1 inter dom m2
  in
     inter_dom <-: m1 munion inter_dom <-: m2
pre forall d in set dom m1 inter dom m2 & m1(d) = m2(d)
```

They are also convenient for decomposing complex structures into their components. For instance, using the previously defined record type Score (see page 24) we can test whether one score is greater than another:

```
let mk_Score(-,w1,-,-,p1) = sc1,
```

---

[1]Remember that only the set bindings can be executed by means of the VDM interpreters.

```
        mk_Score(-,w2,-,-,p2) = sc2
  in (p1 > p2) or (p1 = p2 and w1 > w2)
```

In this particular example we extract the second and fifth components of the two scores. Note that don't care patterns (see page 77) are used to indicate that the remaining components are irrelevant for the processing done in the body of this expression.

Let-be-such-that expressions are useful for abstracting away the non-essential choice of an element from a set, in particular in formulating recursive definitions over sets. An example of this is a version of the sequence filter function (see page 41) over sets:

```
set_filter[@elem] : (@elem -> bool) -> (set of @elem) ->
                    (set of @elem)
set_filter(p)(s) ==
  if s = {}
  then {}
  else let x in set s
       in (if p(x) then {x} else {}) union
          set_filter[@elem](p)(s \ {x});
```

We could alternatively have defined this function using a set comprehension (described in section 6.7):

```
   set_filter[@elem] : (@elem -> bool) -> (set of @elem) ->
                       (set of @elem)
   set_filter(p)(s) ==
     { x | x in set s & p(x)};
```

The last example shows how the optional "be such that" part (**be st**) can be used. This part is especially useful when it is known that an element with some property exists but an explicit expression for such an element is not known or difficult to write. For instance we can exploit this expression to write a selection sort algorithm:

```
   remove : nat * seq of nat -> seq of nat
   remove (x,l) ==
     let i in set inds l be st l(i) = x
     in
       l(1,...,i-1) ^ l(i+1,...,len l)
   pre x in set elems l;

   selection_sort : seq of nat -> seq of nat
```

```
selection_sort (l) ==
  if l = []
  then []
  else let m in set elems l be st
          forall x in set elems l & m <= x
        in
          [m] ^ (selection_sort (remove(m,l)))
```

Here the first function removes a given element from the given list; the second function repeatedly removes the least element in the unsorted portion of the list, and places it at the head of the sorted portion of the list.

6.2 The Define Expression

This expression can only be used inside operations which will be described in section 11. In order to deal with global variables inside the expression part an extra expression construct is available inside operations.

Syntax: expression = …
                   | def expression
                   | … ;

def expression =   '**def**', pattern bind, '=', expression,
                   { ';', pattern bind, '=', expression }, [';'],
                   '**in**', expression ;

Semantics: A *define expression* has the form:

```
def pb1 = e1;
    ...
    pbn = en
in
  e
```

The *define expression* corresponds to a let expression except that the right hand side expressions may depend on the value of the local and/or global variable and that it may not be mutually recursive. It denotes the value of the expression e in the context in which the patterns (or binds) pb1, ..., pbn are matched against the corresponding expressions e1, ..., en[2].

---

[2]If binds are used, it simply means that the values which can match the pattern are further constrained by the type or set expression as explained in Chapter 7.

Examples: The *define expression* is used in a pragmatic way, in order to make the reader aware of the fact that the value of the expression depends upon the global variable.

This can be illustrated by a small example:

```
def user = lib(copy)
in
  if user = <OUT>
  then true
  else false
```

where `copy` is defined in the context, `lib` is global variable (thus `lib(copy)` can be considered as looking up the contents of a part of the variable).

The operation `GroupRunnerUp_expl` in section 12.1 also gives an example of a define expression.

6.3 Unary and Binary Expressions

Syntax: expression = …
                  | unary expression
                  | binary expression
                  | … ;

   unary expression = prefix expression
                  | map inverse ;

   prefix expression = unary operator, expression ;

   unary operator = '+' | '−' | '**abs**' | '**floor**' | '**not**' | '**reverse**'
               | '**card**' | '**power**' | '**dunion**' | '**dinter**'
               | '**hd**' | '**tl**' | '**len**' | '**elems**' | '**inds**' | '**conc**'
               | '**dom**' | '**rng**' | '**merge**' ;

   map inverse = '**inverse**', expression ;

   binary expression = expression, binary operator, expression ;

   binary operator = '+' | '−' | '*' | '/'
               | '**rem**' | '**div**' | '**mod**' | '**'
               | '**union**' | '**inter**' | '\' | '**subset**'
               | '**psubset**' | '**in set**' | '**not in set**'
               | '^'

47

$$
\begin{array}{ll}
| & \text{`++'} \ | \ \textbf{`munion'} \ | \ \text{`<:'} \ | \ \text{`<-:'} \ | \ \text{`:>'} \ | \ \text{`:->'} \\
| & \textbf{`and'} \ | \ \textbf{`or'} \\
| & \text{`=>'} \ | \ \text{`<=>'} \ | \ \text{`='} \ | \ \text{`<>'} \\
| & \text{`<'} \ | \ \text{`<='} \ | \ \text{`>'} \ | \ \text{`>='} \\
| & \textbf{`comp'} \ ;
\end{array}
$$

Semantics: Unary and binary expressions are a combination of operands and operators denoting a value of a specific type. The signature of all these operators is already given in Chapter 3, so no further explanation will be provided here. The map inverse unary operator is treated separately because it is written with postfix notation in the mathematical syntax.

Examples: Examples using these operators were given in Chapter 3, so none will be provided here.

6.4   Conditional Expressions

Syntax: expression = …
| if expression
| cases expression
| … ;

if expression = '**if**', expression, '**then**', expression,
      { elseif expression }, '**else**', expression ;

elseif expression = '**elseif**', expression, '**then**', expression ;

cases expression = '**cases**', expression, ':',
      cases expression alternatives,
      [',', others expression], '**end**' ;

cases expression alternatives = cases expression alternative,
      { ',', cases expression alternative } ;

cases expression alternative = pattern list, '->', expression ;

others expression = '**others**', '->', expression ;

Semantics: *If expressions* and *cases expressions* allow the choice of one from a number of expressions on the basis of the value of a particular expression.

The *if expression* has the form:

```
     if e1
     then e2
     else e3
```

where e1 is a boolean expression, while e2 and e3 are expressions of any type. The if expression denotes the value of e2 evaluated in the given context if e1 evaluates to true in the given context. Otherwise the if expression denotes the value of e3 evaluated in the given context. The use of an **elseif** expression is simply a shorthand for a nested if then else expression in the **else** part of the expression.

The *cases expression* has the form

```
     cases e :
       p11, p12, ..., p1n -> e1,
       ...                 -> ...,
       pm1, pm2, ..., pmk -> em,
       others              -> emplus1
     end
```

where e is an expression of any type, all pij's are patterns which are matched one by one against the expression e. The ei's are expressions of any type, and the keyword **others** and the corresponding expression emplus1 are optional. The cases expression denotes the value of the ei expression evaluated in the context in which one of the pij patterns has been matched against e. The chosen ei is the first entry where it has been possible to match the expression e against one of the patterns. If none of the patterns match e an **others** clause must be present, and then the cases expression denotes the value of emplus1 evaluated in the given context.

Examples: The if expression in the VDM languages corresponds to what is used in most programming languages, while the cases expression in the VDM languages is more general than most programming languages. This is shown by the fact that real pattern matching is taking place, but also because the patterns do not have to be constants as in most programming languages.

An example of the use of conditional expressions is provided by the specification of the mergesort algorithm:

```
  lmerge : seq of nat * seq of nat -> seq of nat
  lmerge (s1,s2) ==
    if s1 = []
    then s2
    elseif s2 = []
    then s1
```

```
    elseif (hd s1) < (hd s2)
    then [hd s1] ^ (lmerge (tl s1, s2))
    else [hd s2] ^ (lmerge (s1, tl s2));

mergesort : seq of nat -> seq of nat
mergesort (l) ==
  cases l:
    []      -> [],
    [x]     -> [x],
    l1 ^ l2 -> lmerge (mergesort(l1), mergesort(l2))
  end
```

The pattern matching provided by cases expressions is useful for manipulating members of type unions. For instance, using the type definition `Expr` from page 27 we have:

```
print_Expr : Expr -> seq1 of char
print_Expr (e) ==
  cases e:
    mk_Const(-) -> "Const of" ^ (print_Const(e)),
    mk_Var(id,-) -> "Var of" ^ id,
    mk_Infix(mk_(e1,op,e2)) -> "Infix of" ^ print_Expr(e1)^","
                                  ^ print_Op(op) ^ ","
                                  ^ print_Expr(e2),
    mk_Cond(t,c,a) -> "Cond of" ^ print_Expr(t) ^ ","
                          ^ print_Expr(c) ^ ","
                          ^ print_Expr(a)
  end;

print_Const : Const -> seq1 of char
print_Const(mk_Const(c)) ==
  if is_nat(c)
  then "nat"
  else -- must be bool
      "bool";
```

The function `print_Op` would be defined similarly.

6.5 Quantified Expressions

Syntax: expression = …

$$| \quad \text{quantified expression}$$
$$| \quad \dots \;;$$

quantified expression $=$ all expression
$$| \quad \text{exists expression}$$
$$| \quad \text{exists unique expression} \;;$$

all expression $=$ '**forall**', bind list, '&', expression ;

exists expression $=$ '**exists**', bind list, '&', expression ;

bind list $=$ multiple bind, { ',', multiple bind } ;

exists unique expression $=$ '**exists1**', bind, '&', expression ;

Semantics: There are three forms of quantified expressions: *universal* (written as **forall**), *existential* (written as **exists**), and *unique existential* (written as **exists1**). Each yields a boolean value **true** or **false**, as explained in the following.

The *universal quantification* has the form:

```
forall mbd1, mbd2, ..., mbdn & e
```

where each `mbdi` is a multiple bind `pi` **in set** `s` (or if it is a type bind `pi :  type`), and `e` is a boolean expression involving the pattern identifiers of the `mbdi`'s. It has the value **true** if `e` is **true** when evaluated in the context of every choice of bindings from `mbd1`, `mbd2, ..., mbdn` and **false** otherwise.

The *existential quantification* has the form:

```
exists mbd1, mbd2, ..., mbdn & e
```

where the `mbdi`'s and the `e` are as for a universal quantification. It has the value **true** if `e` is **true** when evaluated in the context of at least one choice of bindings from `mbd1`, `mbd2, ..., mbdn`, and **false** otherwise.

The *unique existential quantification* has the form:

```
exists1 bd & e
```

where `bd` is either a set bind or a type bind and `e` is a boolean expression involving the pattern identifiers of `bd`. It has the value **true** if `e` is **true** when evaluated in the context of exactly one choice of bindings, and **false** otherwise.

All quantified expressions have the lowest possible precedence. This means that the longest possible constituent expression is taken. The expression is continued to the right as far as it is syntactically possible.

Examples: An example of an existential quantification is given in the function shown below, `QualificationOk`. This function, taken from the specification of a nuclear tracking system in [Fitzgerald&98], checks whether a set of experts has a required qualification.

```
types

ExpertId = token;
Expert :: expertid : ExpertId
          quali : set of Qualification
inv ex == ex.quali <> {};
Qualification = <Elec> | <Mech> | <Bio> | <Chem>

functions

QualificationOK: set of Expert * Qualification -> bool
QualificationOK(exs,reqquali) ==
  exists ex in set exs & reqquali in set ex.quali
```

The function `min` gives us an example of a universal quantification:

```
min(s: set of nat) x: nat
pre s <> {}
post x in set s and
     forall y in set s \ {x} & y < x
```

We can use unique existential quantification to state the functional property satisfied by all maps `m`:

```
forall d in set dom m &
  exists1 r in set rng m & m(d) = r
```

6.6   The Iota Expression

Syntax: expression = …
                    | iota expression

| … ;

iota expression = '**iota**', bind, '&', expression ;

Semantics: An *iota expression* has the form:

```
iota bd & e
```

where `bd` is either a set bind or a type bind, and `e` is a boolean expression involving the pattern identifiers of `bd`. The **iota** operator can only be used if a unique value exists which matches the bind and makes the body expression `e` yield **true** (i.e. **exists1** `bd & e` must be **true**). The semantics of the iota expression is such that it returns the unique value which satisfies the body expression (`e`).

Examples: Using the values `sc1,...,sc4` defined by

```
sc1 = mk_Score (<France>, 3, 0, 0, 9);
sc2 = mk_Score (<Denmark>, 1, 1, 1, 4);
sc3 = mk_Score (<SouthAfrica>, 0, 2, 1, 2);
sc4 = mk_Score (<SaudiArabia>, 0, 1, 2, 1);
```

we have

```
iota x in set {sc1,sc2,sc3,sc4} & x.team = <France>   ≡  sc1
iota x in set {sc1,sc2,sc3,sc4} & x.points > 3        ≡  ⊥
iota x :  Score & x.points < x.won                    ≡  ⊥
```

Notice that the last example cannot be executed and that the last two expressions are undefined - in the former case because there is more than value satisfying the expression, and in the latter because no value satisfies the expression.

6.7   Set Expressions

Syntax: expression = …
       | set enumeration
       | set comprehension
       | set range expression
       | … ;

set enumeration = '{', [expression list], '}' ;

expression list = expression, { ',', expression } ;

set comprehension = '{', expression, '|', bind list,
['&', expression], '}' ;

set range expression = '{', expression, ',', '...', ',',
expression, '}' ;

Semantics: A *Set enumeration* has the form:

```
{e1, e2, e3, ..., en}
```

where `e1` up to `en` are general expressions. It constructs a set of the values of the enumerated expressions. The empty set must be written as `{}`.

The *set comprehension* expression has the form:

```
{e | mbd1, mbd2, ..., mbdn & P}
```

It constructs a set by evaluating the expression `e` on all the bindings for which the predicate `P` evaluates to **true**. A multiple binding can contain both set bindings and type bindings. Thus `mbdn` will look like `pat1` **in set** `s1, pat2 :   tp1, ...`**in set** `s2`, where `pati` is a pattern (normally simply an identifier), and `s1` and `s2` are sets constructed by expressions (whereas `tp1` is used to illustrate that type binds can also be used). Notice however that type binds can only be executed by the VDM interpreters in case the types can be statically declared as finite.

The *set range expression* is a special case of a set comprehension. It has the form

```
{e1, ..., e2}
```

where `e1` and `e2` are numeric expressions. The set range expression denotes the set of integers from `e1` to `e2` inclusive. If `e2` is smaller than `e1` the set range expression denotes the empty set.

Examples: Using the values `Europe={<France>,<England>,<Denmark>,<Spain>}` and `GroupC = {sc1,sc2,sc3,sc4}` (where `sc1,...,sc4` are as defined in the preceding example) we have

```
{<France>, <Spain>} subset Europe       ≡  true
{<Brazil>, <Chile>, <England>}          ≡  false
      subset Europe
{<France>, <Spain>, "France"}           ≡  false
      subset Europe
{sc.team | sc in set GroupC             ≡  {<France>,
      & sc.points > 2}                        <Denmark>}
{sc.team | sc in set GroupC             ≡  {<SouthAfrica>,
      & sc.lost > sc.won }                    <SaudiArabia>}
{2.718,...,3.141}                       ≡  {3}
{3.141,...,2.718}                       ≡  {}
{1,...,5}                               ≡  {1,2,3,4,5}
{ x | x:nat & x < 10 and x mod 2 = 0}   ≡  {0,2,4,6,8}
```

6.8   Sequence Expressions

Syntax: expression = …
　　　　　　　　　 | sequence enumeration
　　　　　　　　　 | sequence comprehension
　　　　　　　　　 | subsequence
　　　　　　　　　 | … ;

　　　 sequence enumeration = '[', [expression list], ']' ;

　　　 sequence comprehension = '[', expression, '|', set bind,
　　　　　　　　　　　　　　　　　　 ['&', expression], ']' ;

　　　 subsequence = expression,
　　　　　　　　　　　 '(', expression, ',', '...', ',', expression, ')' ;

Semantics: A *sequence enumeration* has the form:

```
    [e1, e2, ..., en]
```

where `e1` through `en` are general expressions. It constructs a sequence of the enumerated elements. The empty sequence must be written as `[]`.

A *sequence comprehension* has the form:

```
    [e | pat in set S & P]
```

where the expression `e` will use the identifiers from the pattern `pat` (normally this pattern will simply be an identifier, but the only real requirement is that exactly one pattern identifier

must be present in the pattern). `S` is a set of values (normally natural numbers). The bindings of the pattern identifier must be to some kind of numeric values which then are used to indicate the ordering of the elements in the resulting sequence. It constructs a sequence by evaluating the expression `e` on all the bindings for which the predicate `P` evaluates to **true**.

A *subsequence* of a sequence `l` is a sequence formed from consecutive elements of `l`; from index `n1` up to and including index `n2`. It has the form:

```
l(n1, ..., n2)
```

where `n1` and `n2` are positive integer expressions. If the lower bound `n1` is smaller than 1 (the first index in a non-empty sequence) the subsequence expression will start from the first element of the sequence. If the upper bound `n2` is larger than the length of the sequence (the largest index which can be used for a non-empty sequence) the subsequence expression will end at the last element of the sequence.

Examples: Given that `GroupA` is equal to the sequence

```
[ mk_Score(<Brazil>,2,0,1,6),
  mk_Score(<Norway>,1,2,0,5),
  mk_Score(<Morocco>,1,1,1,4),
  mk_Score(<Scotland>,0,1,2,1) ]
```

then:

```
[GroupA(i).team            ≡   [<Brazil>,
| i in set inds GroupA          <Norway>,
& GroupA(i).won <> 0]           <Morocco>]
[GroupA(i)                 ≡   [mk_Score(<Scotland>,0,1,2,1)]
| i in set inds GroupA
& GroupA(i).won = 0]
GroupA(1,...,2)            ≡   [mk_Score(<Brazil>,2,0,1,6),
                                mk_Score(<Norway>,1,2,0,5)]
[GroupA(i)                 ≡   []
| i in set inds GroupA
& GroupA(i).points = 9]
```

6.9   Map Expressions

Syntax: expression = …
 | map enumeration
 | map comprehension
 | … ;

map enumeration = '{', maplet, { ',', maplet }, '}'
                  | '{', '|->', '}' ;

maplet = expression, '|->', expression ;

map comprehension = '{', maplet, '|', bind list,
                      ['&', expression], '}' ;

Semantics: A *map enumeration* has the form:

```
    {d1 |-> r1, d2 |-> r2, ..., dn |-> rn}
```

where all the domain expressions `di` and range expressions `ri` are general expressions. The empty map must be written as `{|->}`.

A *map comprehension* has the form:

```
    {ed |-> er | mbd1, ..., mbdn & P}
```

where constructs `mbd1, ..., mbdn` are multiple bindings of variables from the expressions `ed` and `er` to sets (or types). The *map comprehension* constructs a mapping by evaluating the expressions `ed` and `er` on all the possible bindings for which the predicate `P` evaluates to **true**.

Examples: Given that `GroupG` is equal to the map

```
    { <Romania>  |-> mk_(2,1,0), <England>  |-> mk_(2,0,1),
      <Colombia> |-> mk_(1,0,2), <Tunisia>  |-> mk_(0,1,2)  }
```

then:

```
{ t |-> let mk_(w,d,-) = GroupG(t)    ≡  {<Romania>  |-> 7,
             in w * 3 + d                 <England>  |-> 6,
| t in set dom GroupG}                    <Colombia> |-> 3,
                                          <Tunisia>  |-> 1}
{ t |-> w * 3 + d                     ≡  {<Romania>  |-> 7,
| t in set dom GroupG, w,d,l:nat         <England>  |-> 6}
& mk_(w,d,l) = GroupG(t)
 and w > 1}
```

6.10 Tuple Constructor Expressions

Syntax: expression = ...
              | tuple constructor
              | ... ;

tuple constructor = '**mk_**', '(', expression, ',', expression list, ')' ;

Semantics: The *tuple constructor expression* has the form:

```
mk_(e1, e2, ..., en)
```

where `ei` is a general expression. It can only be used by the equality and inequality operators.

Examples: Using the map `GroupG` defined in the preceding example, we have:

mk_(2,1,0) **in set rng** GroupG	≡ **true**
mk_("Romania",2,1,0) **not in set rng** GroupG	≡ **true**
mk_(<Romania>,2,1,0) <> **mk_**("Romania",2,1,0)	≡ **true**

6.11 Record Expressions

Syntax: expression = ...
              | record constructor
              | record modifier
              | ... ;

record constructor = '**mk_**', name, '(', [expression list], ')' ;

record modifier = '**mu**', '(', expression, ',', record modification,
                     { ',', record modification } ')' ;

record modification = identifier, '|->', expression ;

Semantics: The *record constructor* has the form:

```
mk_T(e1, e2, ..., en)
```

where the type of the expressions (`e1`, `e2`, `...`, `en`) matches the type of the corresponding entrances in the composite type `T`.

The *record modification* has the form:

```
        mu (e, id1 |-> e1, id2 |-> e2, ..., idn |-> en)
```

where the evaluation of the expression `e` returns the record value to be modified. All the identifiers `idi` must be distinct named entrances in the record type of `e`.

Examples: If `sc` is the value **mk_**`Score(<France>,3,0,0,9)` then

```
        mu (sc, drawn |-> sc.drawn + 1, points |-> sc.points + 1)
        ≡ mk_Score(<France>,3,1,0,10)
```

Further examples are demonstrated in the function `win`. This function takes two teams and a set of scores. From the set of scores it locates the scores corresponding to the given teams (`wsc` and `lsc` for the winning and losing team respectively), then updates these using the **mu** operator. The set of teams is then updated with the new scores replacing the original ones.

```
  win : Team * Team * set of Score -> set of Score
  win (wt,lt,gp) ==
    let wsc = iota sc in set gp & sc.team = wt,
        lsc = iota sc in set gp & sc.team = lt
    in
      let new_wsc = mu (wsc, won |-> wsc.won + 1,
                             points |-> wsc.points + 3),
          new_lsc = mu (lsc, lost |-> lsc.lost + 1)
      in
        (gp \ {wsc,lsc}) union {new_wsc, new_lsc}
  pre forall sc1, sc2 in set gp &
          sc1 <> sc2 <=> sc1.team <> sc2.team
          and {wt,lt} subset {sc.team | sc in set gp}
```

6.12   Apply Expressions

Syntax: expression = …
                         | apply
                         | field select
                         | tuple select
                         | function type instantiation
                         | … ;

apply = expression, '(', [expression list], ')' ;

field select = expression, '.', identifier ;

tuple select = expression, '.#', numeral ;

function type instantiation = name, '[', type, { ',', type }, ']' ;

Semantics: The *field select expression* can be used for records and it has already been explained in section 3.2.5 so no further explanation will be given here.

The *apply* is used for looking up in a map, indexing in a sequence, and finally for calling a function. In section 3.2.3 it has already been shown what it means to look up in a map. Similarly in section 3.2.2 it is illustrated how indexing in a sequence is performed.

In the VDM languages an operation can also be called here. This is not allowed in standard VDM-SL and because this kind of operation call can modify the state such usage should be done with care in complex expressions. Note however that such operation calls are not allowed to throw exceptions.

With such operation calls the order of evaluation can become important. Therefore the type checker will allow the user to enable or disable operation calls inside expressions.

The tuple select expression is used to extract a particular component from a tuple. The meaning of the expression is if `e` evaluates to some tuple **mk_**`(v1,...,vN)` and `M` is an integer in the range `{1,...,N}` then `e.#M` yields `vM`. If `M` lies outside `{1,...,N}` the expression is undefined.

The *function type instantiation* is used for instantiating polymorphic functions with the proper types. It has the form:

```
pf [ t1, ..., tn ]
```

where `pf` is the name of a polymorphic function, and `t1, ..., tn` are types. The resulting function uses the types `t1, ..., tn` instead of the variable type names given in the function definition.

Examples: Recall that `GroupA` is a sequence (see page 56), `GroupG` is a map (see page 57) and `selection_sort` is a function (see page 46):

```
GroupA(1)                  ≡  mk_Score(<Brazil>,2,0,1,6)
GroupG(<Romania>)          ≡  mk_(2,1,0)
GroupG(<Romania>).#2       ≡  1
selection_sort([3,2,9,1,3]) ≡  [1,2,3,3,9]
```

As an example of the use of polymorphic functions and function type instantiation, we use the example functions from section 5:

```
    let emptyInt = empty_bag[int]
    in
      plus_bag[int](-1, emptyInt())

≡


    { -1 |-> 1 }
```

6.13 The New Expression (VDM++ and VDM-RT)

Syntax: expression = …
                | new expression ;

new expression = '**new**', name, '(', [expression list], ')' ;

Semantics: The *new expression* has the form:

```
    new classname(e1, e2, ..., en)
```

An object can be created (also called *instantiated*) from its class description using a *new expression*. The effect of a *new expression* is that a 'new', unique object as described in class classname is created. The value of the *new expression* is a reference to the new object.

If the *new expression* is invoked with no parameters, an object is created in which all instance variables take their "default" values (i.e. the values defined by their initialisation conditions). With parameters, the *new expression* represents a *constructor* (see Section 11.1) and creates customised instances (i.e. where the instance variables may take values which are different from their default values).

Examples: Suppose we have a class called Queue and that default instances of Queue are empty. Suppose also that this class contains a constructor (which will also be called Queue) which takes a single parameter which is a list of values representing an arbitrary starting queue. Then we can create default instances of Queue in which the actual queue is empty using the expression

```
    new Queue()
```

and an instance of Queue in which the actual queue is, say, e1, e2, e3 using the expression

```
new Queue([e1, e2, e3])
```

Using the class `Tree` defined on page 28 we create new `Tree` instances to construct `nodes`:

```
mk_node(new Tree(), x, new Tree())
```

6.14 The Self Expression (VDM++ and VDM-RT)

Syntax: expression = …
                    | self expression ;

self expression = '**self**' ;

Semantics: The *self expression* has the form:

```
self
```

The self expression returns a reference to the object currently being executed. It can be used to simplify the name space in chains of inheritance.

Examples: Using the class `Tree` defined on page 28 we can specify a subclass called `BST` which stores data using the binary search tree approach. We can then specify an operation which performs a binary search tree insertion:

```
Insert : int ==> ()
Insert (x) ==
  (dcl curr_node : Tree := self;

   while not curr_node.isEmpty() do
     if curr_node.rootval() < x
     then curr_node := curr_node.rightBranch()
     else curr_node := curr_node.leftBranch();
   curr_node.addRoot(x);
   )
```

This operation uses a self expression to find the root at which to being traversal prior to insertion. Further examples are given in section 12.9.

6.15   The Threadid Expression (VDM++ and VDM-RT)

Syntax: expression = …
                            | threadid expression ;

threadid expression = '**threadid**' ;

Semantics: The *threadid expression* has the form:

> **threadid**

The threadid expression returns a natural number which uniquely identifies the thread in which the expression is executed. Note that periodic threads gets a new threadid at the start of each new period.

Examples: Using **threadid**'s it is possible to provide a VDM++ base class that implements a Java-style wait-notify in VDM++ using permission predicates. Any object that should be available for the wait-notify mechanism must derive from this base class.

```
class WaitNotify

  instance variables
    waitset : set of nat := {};

  operations
    protected wait: () ==> ()
    wait() ==
      let p = threadid
      in (
        AddToWaitSet( p );
        Awake();
      );

    AddToWaitSet : nat ==> ()
    AddToWaitSet( p ) ==
      waitset := waitset union { p };

    Awake: () ==> ()
    Awake() ==
      skip;

    protected notify: () ==> ()
```

```
      notify() ==
        if waitset <> {}
        then let arbitrary_process in set waitset
              in waitset := waitset \ {arbitrary_process};

      protected notifyAll: () ==> ()
      notifyAll() ==
            waitset := {};

  sync
    mutex(notifyAll, AddToWaitSet, notify);
    per Awake => threadid not in set waitset;

end WaitNotify
```

In this example the **threadid** expression is used in two places:

- In the Wait operation for threads to register interest in this object.

- In the permission predicate for Awake. An interested thread should call Awake following registration using Wait. It will then be blocked until its threadid is removed from the waitset following another thread's call to notify.

6.16 The Lambda Expression

Syntax: expression = …
  | lambda expression
  | … ;

lambda expression = '**lambda**', type bind list, '&', expression ;

type bind list = type bind, { ',', type bind } ;

type bind = pattern, ':', type ;

Semantics: A *lambda expression* is of the form:

```
    lambda pat1 : T1, ..., patn : Tn & e
```

where the pati are patterns, the Ti are type expressions, and e is the body expression. The scope of the pattern identifiers in the patterns pati is the body expression. A lambda expression cannot be polymorphic, but apart from that, it corresponds semantically to an

explicit function definition as explained in chapter 5. A function defined by a lambda expression can be Curried by using a new lambda expression in the body of it in a nested way. When lambda expressions are bound to an identifier they can also define a recursive function.

Examples: An increment function can be defined by means of a lambda expression like:

```
Inc = lambda n : nat & n + 1
```

and an addition function can be Curried by:

```
Add = lambda a : nat & lambda b : nat & a + b
```

which will return a new lambda expression if it is applied to only one argument:

```
Add(5)  ≡  lambda b : nat & 5 + b
```

Lambda expression can be useful when used in conjunction with higher-order functions. For instance using the function set_filter defined on page 45:

```
set_filter[nat](lambda n:nat & n mod 2 = 0)({1,...,10})
≡ {2,4,6,8,10}
```

6.17   Narrow Expressions

Syntax: expression = …
                          | narrow expression
                          | … ;

narrow expression = '**narrow_**', '(', expression, ',', type, ')' ;

Semantics: The *narrow expression* converts the given expression value into the given type, returning a value of that type. It is legal to downcast a class to one of its subclasses, and it is legal to narrow an expression of a union type to one of its subtypes. However, a conversions between two completely unrelated types is a type error. Note that a narrow expression does not guarantee that its argument will be of the correct type at runtime, but using narrow gives extra type information to the specification.

Examples: In following examples, the Test() and Test'() operations should give the same results, but there is a type error in Test() which is resolved in Test' using a narrow expression.

```
class S
end S


class C1 is subclass of S


instance variables
public a : nat := 1;


end C1


class C2 is subclass of S


instance variables
public b : nat := 2;


end C2


class A


operations
public
Test: () ==>  seq of nat
Test() ==
 let list : seq of S = [ new C1(), new C2() ]
 in
   return [ let e = list(i)
            in cases true:
                 (isofclass(C1, e)) -> e.a,
                 (isofclass(C2, e)) -> e.b
               end | i in set inds list ];


public
Test': () ==>  seq of nat
Test'() ==
 let list : seq of S = [ new C1(), new C2() ]
 in
   return [ let e = list(i)
            in cases true:
                 (isofclass(C1, e)) -> narrow_(e, C1).a,
                 (isofclass(C2, e)) -> narrow_(e, C2).b
               end | i in set inds list ];
```

```
end A
```

```
class A

types
public C1 :: a : nat;
public C2 :: b : nat;
public S = C1 | C2;

operations
public
Test: () ==> nat
Test() ==
 let s : S = mk_C1(1)
 in
   let c : C1 = s
   in
     return c.a;

public
Test': () ==> nat
Test'() ==
 let s : S = mk_C1(1)
 in
   let c : C1 = narrow_(s, C1)
   in
     return c.a;
end A
```

6.18   Is Expressions

Syntax: expression = …
                    | general is expression
                    | … ;

         general is expression = is expression
                            | type judgement ;

         is expression = '**is_**', name, ' (', expression, ') '

| is basic type, '(', expression, ')' ;

is basic type = '**is_**', ('**bool**' | '**nat**' | '**nat1**' | '**int**'
         | '**rat**' | '**real**' | '**char**' | '**token**') ;

type judgement = '**is_**', '(', expression, ',', type, ')' ;

Semantics: The *is expression* can be used with values that are either basic or record values (tagged values belonging to some composite type). The is expression yields true if the given value belongs to the basic type indicated or if the value has the indicated tag. Otherwise it yields false.

A type judgement is a more general form which can be used for expressions whose types cannot be statically determined. The expression **is_**(e,t) is equal to true if and only if e is of type t.

Examples: Using the record type Score defined on page 24 we have:

```
is_Score(mk_Score(<France>,3,0,0,9))   ≡   true
is_bool(mk_Score(<France>,3,0,0,9))    ≡   false
is_real(0)                             ≡   true
is_nat1(0)                             ≡   false
```

An example of a type judgement:

```
Domain : map nat to nat | seq of (nat*nat) -> set of nat
Domain(m) ==
  if is_(m, map nat to nat)
  then dom m
  else {d | mk_(d,-) in set elems m}
```

In addition there are examples on page 27.

6.19   Base Class Membership (VDM++ and VDM-RT)

Syntax: expression = …
                     | isofbaseclass expression
                     | … ;

isofbaseclass expression = '**isofbaseclass**', '(', name, expression, ')' ;

Semantic: The function **isofbaseclass** when applied to an object reference expression and a class name name yields the boolean value **true** if and only if name is a root superclass in the inheritance chain of the object referenced to by expression, and **false** otherwise.

Examples: Suppose that `BinarySearchTree` is a subclass of `Tree`, `Tree` is not a subclass of any other class and `Queue` is not related by inheritance to either `Tree` or `BinarySearchTree`. Let `t` be an instance of `Tree`, `b` is an instance of `BinarySearchTree` and `q` is an instance of `Queue`. Then:

isofbaseclass(Tree, t)	$\equiv$	**true**
isofbaseclass(BinarySearchTree, b)	$\equiv$	**false**
isofbaseclass(Queue, q)	$\equiv$	**true**
isofbaseclass(Tree, b)	$\equiv$	**true**
isofbaseclass(Tree, q)	$\equiv$	**false**

6.20 Class Membership

Syntax   expression = ...
                  | isofclass expression
                  | ... ;

   isofclass expression = '**isofclass**', '(', name, expression, ')' ;

Semantics: The function **isofclass** when applied to an object reference expression and a class name name yields the boolean value **true** if and only if `expression` refers to an object of class name or to an object of any of the subclasses of name, and **false** otherwise.

Examples: Assuming the classes `Tree`, `BinarySearchTree`, `Queue`, and identifiers `t`, `b`, `q` as in the previous example, we have:

isofclass(Tree,t)	$\equiv$	**true**
isofclass(Tree,b)	$\equiv$	**true**
isofclass(Tree,q)	$\equiv$	**false**
isofclass(Queue,q)	$\equiv$	**true**
isofclass(BinarySearchTree,t)	$\equiv$	**false**
isofclass(BinarySearchTree,b)	$\equiv$	**true**

6.21 Same Base Class Membership (VDM++ and VDM-RT)

Syntax: expression = ...
                   | samebaseclass expression
                   | ... ;

   samebaseclass expression = '**samebaseclass**',
                                '(', expression, expression, ')' ;

Semantics: The function **samebaseclass** when applied to object references expression1 and expression2 yields the boolean value **true** if and only if the objects denoted by expression1 and expression2 are instances of classes that can be derived from the same root superclass, and **false** otherwise.

Examples: Assuming the classes Tree, BinarySearchTree, Queue, and identifiers t, b, q as in the previous example, suppose that AVLTree is another subclass of Tree, BalancedBST is a subclass of BinarySearchTree, a is an instance of AVLTree and bb is an instance of BalancedBST:

```
samebaseclass(a,b)   ≡  true
samebaseclass(a,bb)  ≡  true
samebaseclass(b,bb)  ≡  true
samebaseclass(t,bb)  ≡  false
samebaseclass(q,a)   ≡  false
```

6.22   Same Class Membership (VDM++ and VDM-RT)

Syntax: expression = …
                         | sameclass expression
                         | … ;

   sameclass expression = '**sameclass**',
                         '(', expression, expression, ')' ;

Semantics: The function **sameclass** when applied to object references expression1 and expression2 yields the boolean value **true** if and only if the objects denoted by expression1 and expression2 are instances of the same class, and **false** otherwise.

Examples: Assuming the classes Tree, BinarySearchTree, Queue, and identifiers t, b, q from section 6.19, and assuming b' is another instance of BinarySearchTree we have:

```
sameclass(b,t)   ≡  false
sameclass(b,b')  ≡  true
sameclass(q,t)   ≡  false
```

6.23   History Expressions (VDM++ and VDM-RT)

Syntax: expression = …
                         | act expression
                         | fin expression
                         | active expression
                         | req expression
                         | waiting expression
                         | … ;

act expression = '**#act**', '(', name, ')'
                  | '**#act**', '(', name list, ')' ;

fin expression = '**#fin**', '(', name, ')'
                  | '**#fin**', '(', name list, ')' ;

active expression = '**#active**', '(', name, ')'
                     | '**#active**', '(', name list, ')' ;

req expression = '**#req**', '(', name, ')'
                  | '**#req**', '(', name list, ')' ;

waiting expression = '**#waiting**', '(', name, ')'
 | '**#waiting**', '(', name list, ')' ;

Semantics: History expressions can only be used in permission predicates (see section 14.1). History expressions may contain one or more of the following expressions:

- **#act** (*operation name*). The number of times that *operation name* operation has been activated.

- **#fin** (*operation name*). The number of times that the *operation name* operation has been completed.

- **#active** (*operation name*). The number of *operation name* operations that are currently active.
  Thus: **#active** (*operation name*) = **#act** (*operation name*) - **#fin** (*operation name*).

- **#req** (*operation name*). The number of requests that has been issued for the *operation name* operation.

- **#waiting** (*operation name*). The number of outstanding requests for the *operation name* operation.
  Thus: **#waiting** (*operation name*) = **#req** (*operation name*) - **#act** (*operation name*).

For all of these operators, the name list version #*history op(op1,. . .,opN)* is simply shorthand for #*history op(op1)* + · · · +#*history op(opN)*.

Examples: Suppose at a point in the execution of a particular thread, three operations, A, B and C may be executed. A sequence of requests, activations and completions occur during this thread. This is shown graphically in figure 6.1.

Here we use the notation *rA* to indicate a request for an execution of operation A, *aA* indicates an activation of A, *fA* indicates completion of an execution of operation A, and likewise for operations B and C. The respective history expressions have the following values after the interval [**S**,**T**]:
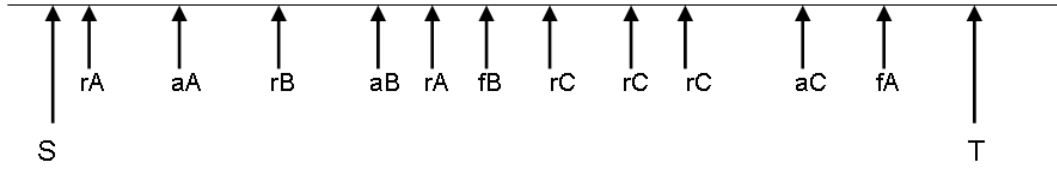
Figure 6.1: *History Expressions*

```
#act(A) = 1       #act(B) = 1       #act(C) = 1       #act(A,B,C) = 3
#fin(A) = 1       #fin(B) = 1       #fin(C) = 0       #fin(A,B,C) = 2
#active(A) = 0    #active(B) = 0    #active(C) = 1    #active(A,B,C) = 1
#req(A) = 2       #req(B) = 1       #req(C) = 3       #req(A,B,C) = 6
#waiting(A) = 1   #waiting(B) = 0   #waiting(C) = 2   #waiting(A,B,C) = 3
```

6.24   The Time Expression (VDM-RT)

Syntax: time expression = '**time**' ;

Semantics: This is simply an easy way to refer to the current time on a given CPU. The time is provided as a natural number, with a resolution of 1 nsec.

Examples: If for example one would like to log when a certain operation takes place one can create an operation such as `logEnvToSys` below.

```
public logEnvToSys: nat ==> ()
logEnvToSys (pev) == e2s := e2s munion {pev |-> time};
```

6.25   Literals and Names

Syntax: expression = …
                 | name
                 | old name
                 | symbolic literal
                 | … ;

name = identifier, ['`', identifier] ;

name list = name, { ',', name } ;

old name = identifier, '~' ;

Semantics: *Names* and *old names* are used to access definitions of functions, operations, values and state components. A *name* has the form:

```
id1`id2
```

where `id1` and `id2` are simple identifiers. If a name consists of only one identifier, the identifier is defined within scope, i.e. it is defined either locally as a pattern identifier or variable, or globally within the current module as a function, operation, value or global variable. Otherwise, the identifier `id1` indicates the module/class name where the construct is defined (see also section 13.1 and section 13.3.1 and appendix B.)

An *old name* is used to access the old value of global variables in the post condition of an operation definition (see chapter 11) and in the post condition of specification statements (see section 12.16). It has the form:

```
id˜
```

where `id` is a state component.

Symbolic literals are constant values of some basic type.

Examples: *Names* and *symbolic literals* are used throughout all examples in this document (see appendix B.2).

For an example of the use of *old names*, consider the VDM-SL state defined as:

```
state sigma of
  numbers : seq of nat
  index   : nat
inv mk_sigma(numbers, index) ==
    index not in set elems numbers
init s == s = mk_sigma([], 1)
end
```

For an example of the use of *old names*, consider the VDM++/VDM-RT instance variables defined as:

```
    instance variables
      numbers: seq of nat := [];
      index  : nat := 1;
    inv index not in set elems numbers;
```

We can define an operation that increases the variable index in an implicit manner:

```
    IncIndex()
    ext wr index : nat
    post index = index˜ + 1
```

The operation IncIndex manipulates the variable index, indicated with the **ext wr** clause. In the post condition, the new value of index is equal to the old value of index plus 1. (See more about operations in chapter 11).

For a simple example of module/class names, suppose that a function called build_rel is defined (and exported) in a module/class called CGRel as follows:

```
types

  Cg = <A> | <B> | <C> | <D> | <E> | <F> |
       <G> | <H> | <J> | <K> | <L> | <S>;
  CompatRel = map Cg to set of Cg

functions

  build_rel : set of (Cg * Cg) -> CompatRel
  build_rel (s) == {|->}
```

In another module/class we can access this function by in VDM-SL first importing the module CGRel then by using the following call

```
  CGRel`build_rel({mk_(<A>, <B>)})
```

Note that in VDM++ and VDM-RT build_rel function can additionally have an access modifier allowing access to it outside the defining class.

6.26 The Undefined Expression

Syntax: expression = ...
               | undefined expression ;

        undefined expression = '**undefined**' ;

Semantics: The *undefined expression* is used to state explicitly that the result of an expression is undefined. This could for instance be used if it has not been decided what the result

of evaluating the else-branch of an if-then-else expression should be. When an *undefined expression* is evaluated the VDM interpreters will terminate the execution and report that an undefined expression was evaluated.

Pragmatically use of undefined expressions differs from pre-conditions: use of a pre-condition means it is the caller's responsibility to ensure that the pre-condition is satisfied when the function is called; if an undefined expression is used it is the called function's responsibility to deal with error handling.

Examples: We can check that the type invariant holds before building `Score` values:

```
build_score : Team * nat * nat * nat * nat -> Score
build_score (t,w,d,l,p) ==
  if 3 * w + d = p
  then mk_Score(t,w,d,l,p)
  else undefined
```

6.27 The Precondition Expression

Syntax: expression = …
                    | precondition expression ;

         precondition expression = '**pre_**', '(', expression,
                            [{ ',', expression }], ')' ;

Semantics: Assuming `e` is of function type the expression **pre_**(e,e1,...,en) is true if and only if the pre-condition of `e` is true for arguments e1,...,em where m is the arity of the pre-condition of `e`. If `e` is not a function or m > n then the result is **true**. If `e` has no pre-condition then the expression equals true.

Examples: Consider the functions `f` and `g` defined below

```
f : nat * nat -> nat
f(m,n) == m div n
pre n <> 0;

g (n: nat) sqrt: nat
pre n >= 0
post sqrt * sqrt <= n and
     (sqrt+1) * (sqrt+1) > n
```

Then the expression

```
pre_(let h in set {f,g,lambda mk_(x,y): nat * nat & x div y}
       in h, 1,0,-1)
```

is equal to

- false if h is bound to f since this equates to **pre_f**(1,0);

- true if h is bound to g since this equates to **pre_g**(1);

- true if h is bound to **lambda mk_**(x,y):**nat** * **nat** & x **div** y since there is no pre-condition defined for this function.

Note that however h is bound, the last argument (−1) is never used.

Chapter 7

Patterns

Syntax: pattern bind = pattern | bind ;

        pattern = pattern identifier
              | match value
              | set enum pattern
              | set union pattern
              | seq enum pattern
              | seq conc pattern
              | map enumeration pattern
              | map muinon pattern
              | tuple pattern
              | record pattern
              | object pattern ;

        pattern identifier = identifier | '–' ;

        match value = symbolic literal
                 | '(', expression, ')' ;

        set enum pattern = '{', [pattern list], '}' ;

        set union pattern = pattern, '**union**', pattern ;

        seq enum pattern = '[', [pattern list], ']' ;

        seq conc pattern = pattern, 'ˆ', pattern ;

        map enumeration pattern = '{', [maplet pattern list], '}'
                      | '{', '|–>', '}' ;

        maplet pattern list = maplet pattern, { ',', maplet pattern } ;

maplet pattern = pattern, '|−>', pattern ;

map muinon pattern = pattern, '**munion**', pattern ;

tuple pattern = '**mk_(**', pattern, ',', pattern list, ')' ;

record pattern = '**mk_**', name, '(', [pattern list], ')' ;

object pattern = '**obj_**', identifier, '(', [field pattern list], ')'[1];

field pattern list = field pattern, { ',', field pattern } ;

field pattern = identifier, '|−>', pattern ;

pattern list = pattern, { ',', pattern } ;

Semantics: A pattern is always used in a context where it is matched to a value of a particular type. Matching consists of checking that the pattern can be matched to the value, and binding any pattern identifiers in the pattern to the corresponding values, i.e. making the identifiers denote those values throughout their scope. In some cases where a pattern can be used, a bind can be used as well (see next chapter). If a bind is used it simply means that additional information (a type or a set expression) is used to constrain the possible values which can match the given pattern.

Matching is defined as follows

1. A *pattern identifier* fits any type and can be matched to any value. If it is an identifier, that identifier is bound to the value; if it is the don't-care symbol '−', no binding occurs.

2. A *match value* can only be matched against the value of itself; no binding occurs. If a match value is not a literal like e.g. 7 or <RED> it must be an expression enclosed in parentheses in order to discriminate it to a pattern identifier.

3. A *set enumeration pattern* fits only set values. The patterns are matched to distinct elements of a set; all elements must be matched.

4. A *set union pattern* fits only set values. The two patterns are matched to a partition of two subsets of a set. In the VDM interpreters the two subsets will always be chosen such that they are non-empty and disjoint.

5. A *sequence enumeration pattern* fits only sequence values. Each pattern is matched against its corresponding element in the sequence value; the length of the sequence value and the number of patterns must be equal.

---

[1]**Note:** object pattern is only be used in VDM++ and VDM-RT. object pattern is explained in 7.1

6. A *sequence concatenation pattern* fits only sequence values. The two patterns are matched against two subsequences which together can be concatenated to form the original sequence value. In the VDM interpreters the two subsequences will always be chosen so that they are non-empty.

7. A *map enumeration pattern* fits only map values.

8. A *maplet pattern list* are matched to distinct elements of a map; all elements must be matched.

9. A *map munion pattern* fits only map values. The two patterns are matched to a partition of two sub maps of a map. In the VDM interpreters the two sub maps will always be chosen such that they are non-empty and disjoint.

10. A *tuple pattern* fits only tuples with the same number of elements. Each of the patterns are matched against the corresponding element in the tuple value.

11. A *record pattern* fits only record values with the same tag. Each of the patterns are matched against the field of the record value. All the fields of the record must be matched.

12. A *object pattern* is explained in 7.1

Examples: The simplest kind of pattern is the pattern identifier. An example of this is given in the following let expression:

```
let top = GroupA(1)
in top.sc
```

Here the identifier `top` is bound to the head of the sequence `GroupA` and the identifier may then be used in the body of the let expression.

In the following examples we use match values:

```
let a = <France>
in cases GroupA(1).team:
     <Brazil> -> "Brazil are winners",
     (a)      -> "France are winners",
     others   -> "Neither France nor Brazil are winners"
   end;
```

Match values can only match against their own values, so here if the team at the head of `GroupA` is `<Brazil>` then the first clause is matched; if the team at the head of `GroupA` is `<France>` then the second clause is matched. Otherwise the **others** clause is matched. Note here that the use of brackets around `a` forces `a` to be considered as a match value.

Set enumerations match patterns to elements of a set. For instance in

```
  let {sc1, sc2, sc3, sc4} = elems GroupA
  in
    sc1.points + sc2.points + sc3.points + sc4.points;
```

the identifiers `sc1`, `sc2`, `sc3` and `sc4` are bound to the four elements of `GroupA`. Note that the choice of binding is loose – for instance `sc1` may be bound to [any] element of **elems** `GroupA`. In this case if **elems** `GroupA` does not contain precisely four elements, then the expression is not well-formed.

A set union pattern can be used to decompose a set for recursive function calls. An example of this is the function `set2seq` which converts a set into a sequence (with arbitrary order):

```
set2seq[@elem] : set of @elem -> seq of @elem
set2seq(s) ==
  cases s:
    {}           -> [],
    {x}          -> [x],
    s1 union s2 -> (set2seq[@elem](s1))^(set2seq[@elem](s2))
  end
```

In the third cases alternative we see the use of a set union pattern. This binds `s1` and `s2` to arbitrary subsets of `s` such that they partition `s`. The VDM interpreters always ensures a disjoint partition.

Sequence enumeration patterns can be used to extract specific elements from a sequence. An example of this is the function `promoted` which extracts the first two elements of a sequence of scores and returns the corresponding pair of teams:

```
promoted : seq of Score -> Team * Team
promoted([sc1,sc2]^-) == mk_(sc1.team,sc2.team);
```

Here `sc1` is bound to the head of the argument sequence, and `sc2` is bound to the second element of the sequence. If `promoted` is called with a sequence with fewer than two elements then a runtime error occurs. Note that as we are not interested in the remaining elements of the list we use a don't care pattern for the remainder.

The preceding example also demonstrated the use of sequence concatenation patterns. Another example of this is the function `quicksort` which implements a standard quicksort algorithm:

```
quicksort : seq of nat -> seq of nat
quicksort (l) ==
```

```
  cases l:
    [] -> [],
    [x] -> [x],
    [x,y] -> if x < y then [x,y] else [y,x],
    -^[x]^- ->
           quicksort([l(i) | i in set inds l & l(i) < x])
                    ^ [l(i) | i in set inds l & l(i) = x] ^
           quicksort([l(i) | i in set inds l & l(i) > x])
  end
```

Here, in the second cases clause a sequence concatenation pattern is used to decompose `l` into an arbitrary pivot element and two subsequences. The pivot is used to partition the list into those values less than the pivot and those values greater, and these two partitions are recursively sorted.

Maplet pattern match patterns to elements of a maplet.

```
let {a |-> b} = {1 |-> 2} in mk_(a,b) = mk_(1,2)
```

Maplet pattern list match patterns to elements of each maplet in a map.

```
let {1 |-> a,a |-> b,b |-> c} = {1 |-> 4,2 |-> 3,4 |-> 2} in
c = 3
```

Map munion pattern can be used to decompose a map for recursive function calls. Following `map2seq` function converts a map to a seq of maplet.

```
public map2seq[@T1, @T2] :
    map @T1 to @T2 -> seq of (map @T1 to @T2)
map2seq(m) ==
  cases m:
    ({|->})      -> [],
    {- |-> -}    -> [m],
    m1 munion m2 ->
        map2seq[@T1, @T2] (m1) ^  map2seq[@T1, @T2] (m2)
  end;
```

Here, in the third cases clause a map munion pattern is used to decompose `m` into two maps.

Tuple patterns can be used to bind tuple components to identifiers. For instance since the function `promoted` defined above returns a pair, the following value definition binds the

winning team of `GroupA` to the identifier `Awinner`:

```
values

   mk_(Awinner,-) = promoted(GroupA);
```

Record patterns are useful when several fields of a record are used in the same expression. For instance the following expression constructs a map from team names to points score:

```
{ t |-> w * 3 + l | mk_Score(t,w,l,-,-) in set elems GroupA}
```

The function `print_Expr` on page 50 also gives several examples of record patterns.

7.1 Object Pattern (VDM++ and VDM-RT)

Syntax: object pattern = '**obj_**', identifier, '(', [field pattern list], ')' ;

field pattern list = field pattern, { ',', field pattern } ;

field pattern = identifier, '|->', pattern ;

Semantics: An object pattern matches object references. An object is matched to the class identified by the class name led by the prefix `obj_`. If the class of the object is class `identifier` or any of its subclasses, separate instance variables named in an object pattern are matched against the object from left to right. It is not necessary to enumerate all the instance variables of a referred class in an object pattern.

Object patterns may appear only where access to the instance variables named in the pattern would be permitted. So private instance variables can only be matched within operations of the same class. This also implies an object pattern may not appear in function definitions, preconditions, postconditions or invariants. The type checker will validate accessibility of each instance variable in the pattern. The value of referred instance variables may or may not change during the matching.

Examples: Using an object pattern, an object can be matched by the values of its instance variables. The following simple example evaluates to `<STUDENT>` if and only if the object reference person points at an instance of the class Student or any of its subclasses, otherwise `<NOT_STUDENT>`.

```
cases person:
   obj_Student() -> <STUDENT>,
   others -> <NOT_STUDENT>
```

```
end
```

The object pattern `obj_Student()` does not only match an instance of the class Student. If a class Undergraduate is a subclass of Student, the object pattern `obj_Student()` also matches any instance of the class Undergraduate.

The following example matches the instance variable name against `"John"`.

```
cases person:
    obj_Student(name |-> "John") -> <JOHN>,
    others -> <NOT_A_STUDENT_OR_NOT_JOHN>
end
```

An object pattern can also bind the value of instance variable. For example, the following expression binds n to the `name` instance variable of `person` if person is an instance of the class Student.

```
cases person:
    obj_Student(name |-> n) -> n,
    others -> ""
end
```

An object pattern may contain another object pattern. The following expression gives the name and department concatenated by `"@"` for Students, and add the title for Professors.

```
cases person:
    obj_Student
        (name |-> n, dept |-> obj_Department(name->dname))
                    -> n^"@"^dname,
    obj_Professor(name |-> n) -> "Prof. "^n,
    others -> ""
end
```

An object pattern may appear in conjunction with other kinds of patterns. The following operation takes a sequence of `Professor` objects and returns titled names sorted into initials. For example, `professorDict([new Professor("Smith"), new Professor("Scott"), new Professor("Adams")])` evaluates to

```
{'A' |-> {"Prof. Adams"},
 'S' |-> {"Prof. Scott", "Prof. Smith"}}.
```

The first pattern of the cases expression consists of a list concatenation pattern. The first element of the list is matched against an object pattern whose instance variable `name` is also matched against another sequence concatenation pattern of the initial character `initial` and the `remaining`.

```
professorDict: seq of Professor
                        ==> map char to set of seq of char
professorDict(ps) == return cases ps:
  [obj_Professor(name |-> [initial]^remaining)]^rest ->
    let
      dict = professorDict(rest),
      name = "Prof. "^[initial]^remaining,
      names = if initial in set dom dict
                  then dict(initial)
                  else {}
    in
  dict ++ {initial |-> names union {name}},
  [-]^rest -> professorDict(rest),[] -> {|->}
  end;
```

Chapter 8

Bindings

Syntax: bind = set bind | type bind ;

         set bind = pattern, '**in set**', expression ;

         type bind = pattern, ':', type ;

         bind list = multiple bind, { ',', multiple bind } ;

         multiple bind = multiple set bind
                       | multiple type bind ;

         multiple set bind = pattern list, '**in set**', expression ;

         multiple type bind = pattern list, ':', type ;

Semantics: A *bind* matches a pattern to a value. In a *set bind* the value is chosen from the set defined by the set expression of the bind. In a *type bind* the value is chosen from the type defined by the type expression. *Multiple bind* is the same as *bind* except that several patterns are bound to the same set or type. Notice that type binds **can** only be executed by the VDM interpreters in case the type can be deduced to be finite statically. This would require the VDM interpreters to search through infinite domains like the natural numbers.

Examples: Bindings are mainly used in quantified expressions and comprehensions which can be seen from these examples:

```
forall i, j in set inds list & i < j => list(i) <= list(j)

{ y | y in set S & y > 2 }

{ y | y: nat & y > 3 }
```

```
occurs : seq1 of char * seq1 of char -> bool
occurs (substr,str) ==
  exists i,j in set inds str & substr = str(i,...,j);
```

Chapter 9

Value (Constant) Definitions

The VDM languages supports the definition of constant values. A value definition corresponds to a constant definition in traditional programming languages.

Syntax: value definitions = '**values**', [access value definition],
                            { ';', access value definition }, [';'] ;

          access value definition = [access], value definition ;

          value definition = pattern, [':', type], '=', expression ;

Semantics: The value definition has the form:

```
values
  access pat1 = e1;
  ...
  access patn = en
```

where the `access` part only can be used in VDM++ and VDM-RT.

The global values (defined in a value definition) can be referenced at all levels in a VDM specification. However, in order to be able to execute a specification these values must be defined before they are used in the sequence of value definitions. This "declaration before use" principle is only used by the VDM interpreters for value definitions. Thus for instance functions can be used before they are declared. In standard VDM-SL there are not any restrictions on the order of the definitions at all. It is possible to provide a type restriction as well, and this can be useful in order to obtain more exact type information.

Details of the VDM++ and VDM-RT access specifiers can be found in section 13.3.3.

Examples: The example below, taken from [Fitzgerald&98] assigns token values to identifiers `p1` and `eid2`, an `Expert` record value to `e3` and an `Alarm` record value to `a1`.

```
types

Period = token;
ExpertId = token;
Expert :: expertid : ExpertId
          quali : set of Qualification
inv ex == ex.quali <> {};
Qualification = <Elec> | <Mech> | <Bio> | <Chem>;
Alarm :: alarmtext : seq of char
         quali : Qualification

values

public p1: Period = mk_token("Monday day");
private eid2 : ExpertId = mk_token(145);
protected e3 : Expert = mk_Expert(eid2, { <Mech>, <Chem> });
a1 : Alarm = mk_Alarm("CO2 detected", <Chem>)
```

As this example shows, a value can depend on other values which are defined previous to itself. The access modifiers **private**, **protected** and **public** can only be used in VDM++ and VDM-RT. A top-level VDM-SL specification can consist of specifications from a number of files or modules (see section 13.1). It is good practice not to let a value depend on values defined in other modules as the ordering is important.

Chapter 10

Declaration of Modifiable State Components

Syntactically the definition of state components that can be modified using VDM operations differ in VDM-SL compared with VDM++ and VDM-RT. Since VDM-SL is module based the state definition is similar to a monolitic record like construct. On the other hand VDM++ and VDM-RT are object-oriented and thus state components needs to be more flexible in order to enable inheritence of such definitions and thus they are defined in terms of instance variables. In the two sections in this chapter the two different ways of defining states is presented.

10.1   Instance Variables (VDM++ and VDM-RT)

Both an object instantiated from a class description and the class itself can have an internal state, also called the *instance variables* of the object or class. In the case of objects, we also refer to this state as the global state of the object.

Syntax: instance variable definitions = '**instance**', '**variables**',
         [instance variable definition,
         { ';', instance variable definition }] ;

    instance variable definition = access assignment definition
        | invariant definition ;

    access assignment definition = ([access], ['**static**'])
        | (['**static**'], [access]),
        assignment definition ;

    assignment definition = identifier, ':', type, [':=', expression] ;

    invariant definition = '**inv**', expression ;

Semantics: The section describing the internal state is preceded by the keyword **instance variables**. A list of instance variable definitions and/or invariant definitions follows. Each instance variable definition consists of an instance variable name with its corresponding type indication and may also include an initial value and access and **static** specifiers. Details of the access and **static** specifiers can be found in section 13.3.3.

It is possible to restrict the values of the instance variables by means of invariant definitions. Each invariant definition, involving one or more instance variables, may be defined over the values of the instance variables of objects of a class. All instance variables in the class including those inherited from superclasses are visible in the invariant expression. Each invariant definition must be a boolean expression that limits the values of the instance variables to those where the expression is true. All invariant expressions must be true during the entire lifetime of each object of the class.

The overall invariant expression of a class is all the invariant definitions of the class and its superclasses combined by logical **and** in the order that they are defined in 1) the superclasses and 2) the class itself.

Example: The following examples show instance variable definitions. The first class specifies one instance variable:

```
class GroupPhase

types

GroupName = <A> | <B> | <C> | <D> | <E> | <F> | <G> | <H>;
Team = ... -- as on page 24
Score::team : Team
       won    : nat
       drawn  : nat
       lost   : nat
       points : nat;

instance variables
gps : map GroupName to set of Score;
inv forall gp in set rng gps &
     (card gp = 4 and
      forall sc in set gp & sc.won + sc.lost + sc.drawn <= 3)

end GroupPhase
```

10.2 The State Definition (VDM-SL)

If global variables are desired in a VDM-SL specification, it is possible to make a state definition. The components of the state definition can be considered the collection of global variables which can be referenced inside operations. A state in a module is initialised before any of the operation definitions (using that state) in a module can be used by the VDM interpreters.

Syntax: state definition = '**state**', identifier, '**of**', field list,
　　　　　　　　[invariant], [initialisation], '**end**', [';'] ;

　　　　invariant = '**inv**', invariant initial function ;

　　　　initialisation = '**init**', invariant initial function ;

　　　　invariant initial function = pattern, '==', expression ;

Semantics: The state definition has the form:

```
state ident of
  id1 : type1
  ...
  idn : typen
inv  pat1 == invpred
init pat2 == initpred
end
```

A state identifier `idn` is declared of a specific type `typen`. The invariant `invpred` is a boolean expression denoting a property which must hold for the state `ident` at all times. `initpred` denotes a condition which must hold initially. It should be noticed that in order to use the VDM interpreters, it is necessary to have an initialisation predicate (if any of the operations using the state are to be executed). In addition the body of this initialisation predicate must be a binary equality expression with the name (which also must be used as the pattern) of the entire state on the left-hand side of the equality and the right-hand side must evaluate to a record value of the correct type. This enables the VDM interpreters to evaluate the `initpred` condition. A simple example of an initialisation predicate is shown below:

```
state St of
  x: nat
  y: nat
  l: seq1 of nat
init s == s = mk_St(0,0,[1])
end
```

In the specification of both the invariant and the initial value the state must be manipulated as a whole, and this is done by referring to it as a record tagged with the state name (see the example). When a field in the state is manipulated in some operation, the field must however be referenced to directly by the field name without pre-fixing it with the state name.

Examples: In the following example we create one state variable:

```
types

GroupName = <A> | <B> | <C> | <D> | <E> | <F> | <G> | <H>

state GroupPhase of
  gps : map GroupName to set of Score
inv mk_GroupPhase(gps) ==
  forall gp in set rng gps &
    (card gp = 4 and
    forall sc in set gp & sc.won + sc.lost + sc.drawn <= 3)
init gp ==
  gp = mk_GroupPhase({<A> |-> init_sc({<Brazil>, <Norway>,
                                        <Morocco>, <Scotland>}),
                     ...})
end

functions

init_sc : set of Team -> set of Score
init_sc (ts) ==
  { mk_Score (t,0,0,0,0) | t in set ts }
```

In the invariant we state that each group has four teams, and no team plays more than three games. Initially no team has played any games.

Chapter 11

Operation Definitions

Operations have already been mentioned in chapter 4. The general form is described here and for VDM++ and VDM-RT special operations called *constructors* which are used for constructing instances of a class are described in section 11.1. Note that the **async** keyword can only be used in the VDM-RT dialect.

Syntax: operation definitions = '**operations**', [access operation definition],
                                    { ';', access operation definition } , [';'] ;

            access operation definition = (['**async**'] [access], ['**static**'])
                                 | (['**async**'] ['**static**'], [access]),
                                      operation definition ;

            operation definition = explicit operation definition
                                | implicit operation definition
                                | extended explicit operation definition ;

            explicit operation definition = identifier, ':', operation type,
                                    identifier, parameters,
                                    '==',
                                    operation body,
                                    ['**pre**', expression],
                                    ['**post**', expression] ;

            implicit operation definition = identifier, parameter types,
                                    [identifier type pair list],
                                    implicit operation body ;

            implicit operation body = [externals],
                                    ['**pre**', expression],
                                    '**post**', expression,
                                    [exceptions] ;

extended explicit operation definition = identifier,
                                          parameter types,
                                          [identifier type pair list],
                                          '==', operation body,
                                          [externals],
                                          ['**pre**', expression],
                                          ['**post**', expression],
                                          [exceptions] ;

operation type = discretionary type, '==>', discretionary type ;

discretionary type = type | '()' ;

parameters = '(', [pattern list], ')' ;

pattern list = pattern, { ',', pattern } ;

operation body = statement
                  | '**is not yet specified**'
                  | '**is subclass responsibility**' ;

externals = '**ext**', var information, { var information } ;

var information = mode, name list, [':', type] ;

mode = '**rd**' | '**wr**' ;

name list = identifier, { ',', identifier } ;

exceptions = '**errs**', error list ;

error list = error, { error } ;

error = identifier, ':', expression, '->', expression ;

Semantics: Operations in VDM are per default synchronous but if the keyword "**async**" is used in VDM-RT in front of an operation definition it means that that operation will be treated as an asynchronous operation. This means that the operation cannot have a return type and the thread calling an asynchronous operation will continue its own execution after having requested the invocation of the asynchronous operation. Note that constructors cannot be declared asynchronous. In both VDM++ and VDM-RT the details of the access and **static** specifiers can be found in section 13.3.3. Note that a static operation may not call non-static operations, and self expressions cannot be used in the definition of a static operation.

The following example of an explicit operation updates the VDM-SL state GroupPhase and the VDM++ instance variables of class GroupPhase when one team beats another.

```
Win : Team * Team ==> ()
Win (wt,lt) ==
  let gp in set dom gps be st
        {wt,lt} subset {sc.team | sc in set gps(gp)}
  in gps := gps ++ { gp |->
                          {if sc.team = wt
                           then mu(sc, won |-> sc.won + 1,
                                        points |-> sc.points + 3)
                           elseif sc.team = lt
                           then mu(sc, lost |-> sc.lost + 1)
                           else sc
                     | sc in set gps(gp)}}
pre exists gp in set dom gps &
        {wt,lt} subset {sc.team | sc in set gps(gp)};
```

An explicit operation consists of a statement (or several composed using a block statement), as described in chapter 12. The statement may access any state/instance variables it wishes, reading and writing to them as it sees fit.

An implicit operation is specified using an optional pre-condition, and a mandatory post-condition. For example we could specify the `Win` operation implicitly:

```
Win (wt,lt: Team)
ext wr gps : map GroupName to set of Score
pre exists gp in set dom gps &
        {wt,lt} subset {sc.team | sc in set gps(gp)}
post exists gp in set dom gps &
        {wt,lt} subset {sc.team | sc in set gps(gp)}
        and gps = gps~ ++
                      { gp |->
                        {if sc.team = wt
                         then mu(sc, won |-> sc.won + 1,
                                     points |-> sc.points + 3)
                         elseif sc.team = lt
                         then mu(sc, lost |-> sc.lost + 1)
                         else sc
                     | sc in set gps(gp)}};
```

The externals field lists the state/instance variables that the operation will manipulate. The state/instance variables listed after the reserved word **rd** can only be read whereas the operation can both read and write the variables listed after **wr**.

In VDM-SL these pre- and post-conditions the VDM interpreters also creates new functions as with the pre- and post-conditions of operation definitions. However, if a specification contains a global state, the state is also part of the newly created functions. Thus, functions with the following signatures are created for operations with pre- and/or post-conditions[1]:

```
pre_Op : InType * State +> bool

post_Op : InType * OutType * State * State +> bool
```

with the following exceptions:

- If the operation does not take any arguments, the InType part of the signature is left out in both the **pre**_Op and **post**_Op signatures.
- If the operation does not return a value, the OutType part is left out in the **post**_Op signature.
- If the specification does not define a state, the State part(s) of both signatures are left out.

In the **post**_Op signature, the first State part is for the old state, whereas the second State part is for the state after the operation call.

For instance, consider the following specifications:

```
module A

definitions

state St of
  n : nat
end

operations

Op1 (a : nat) b :nat
pre a > 0
post b = 2 * a;

Op2 () b : nat
post b = 2;
```

---

[1]However, you should remember that these pre and post condition predicates for an operation are simply boolean functions and the state components are thus not changed by calling such a predicate.

```
Op3 ()
post true

end A
```

```
module B

definitions

operations

Op1 (a : nat) b : nat
pre a > 0
post b = 2 * a;

Op2 () b : nat
post b = 2;

Op3 ()
post true

end B
```

For **module A** we could then quote the pre and post conditions defined in this specification as illustrated below

Quote expression	Explanation
pre_Op1(1,**mk**_St(2))	a bound to 1 in state St with n bound to 2
post_Op1(1,2,**mk**_St(1), **mk**_St(2))	a bound to 1, b bound to 2, state before with n bound to 1, state after with n bound to 2
post_Op2(2,**mk**_St(1), **mk**_St(2))	b bound to 2, state before with n bound to 1, state after with n bound to 2
post_Op3(**mk**_St(1), **mk**_St(2))	state before with n bound to 1, state after with n bound to 2

For **module B** we can quote the pre and post conditions defined in this specification as illustrated below

Quote expression	Explanation
pre_Op1(1)	a bound to 1
post_Op1(1,2)	a bound to 1, b bound to 2
post_Op2(2)	b bound to 2
post_Op3()	No binding at all

The exceptions clause can be used to describe how an operation should deal with error situations. The rationale for having the exception clause is to give the user the ability to separate the exceptional cases from the normal cases. The specification using exceptions does not give any commitment as to how exceptions are to be signalled, but it gives the means to show under which circumstances an error situation can occur and what the consequences are for the result of calling the operation.

The exception clause has the form:

```
errs COND1: c1 -> r1
     ...
     CONDn: cn -> rn
```

The condition names `COND1, ..., CONDn` are identifiers which describe the kind of error which can be raised[2]. The condition expressions `c1, ..., cn` can be considered as pre-conditions for the different kinds of errors. Thus, in these expressions the identifiers from the arguments list and the variables from the externals list can be used (they have the same scope as the pre-condition). The result expressions `r1, ..., rn` can correspondingly be considered as post-conditions for the different kinds of errors. In these expressions the result identifier and old values of global variables (which can be written to) can also be used. Thus, the scope corresponds to the scope of the post-condition.

An operation definition making use of an **errs** clause essentially gets an effective pre-condition which is a disjunction of the original pre-condition and all the condition expressions `c1, ..., cn`. The effective post-condition in these cases becomes a disjunction of the conjuncts (`orig_pre and orig_post`) and the `c1 and r1,...,cn and rn`.

Superficially there appears to be some redundancy between exceptions and pre-conditions here. However there is a conceptual distinction between them which dictates which should be used and when. The pre-condition specifies what callers to the operation must ensure for correct behaviour; the exception clauses indicate that the operation being specified takes responsibility for error handling when an exception condition is satisfied. Hence normally exception clauses and pre-conditions do not overlap.

The next VDM-SL example of an operation uses the following state definition:

```
    state qsys of
```

---

[2]Notice that these names are purely of mnemonic value, i.e. semantically they are not important.

```
      q : Queue
    end
```

The next VDM++/VDM-RT example of an operation uses the following instance variable definition:

```
    instance variables
      q : Queue
```

This example shows how exceptions with an implicit definition can be used:

```
    DEQUEUE() e: [Elem]
    ext wr q : Queue
    post q˜ = [e] ˆ q
    errs QUEUE_EMPTY: q = [] -> q = q˜ and e = nil
```

This is a dequeue operation which uses a global variable q of type Queue to get an element e of type Elem out of the queue. The exceptional case here is that the queue in which the exception clause specifies how the operation should behave is empty.

Note that the VDM interpreters for VDM-SL models creates a function here:

```
    post_DEQUEUE: [Elem] * qsys * qsys +> bool
```

11.1 Constructors (VDM++ and VDM-RT)

Constructors are operations which have the same name as the class in which they are defined and which create new instances of that class. Their return type must therefore be the same class name, and if a return value is specified this should be **self** though this can optionally be omitted.

Multiple constructors can be defined in a single class using operation overloading as described in section 13.3.1.

Chapter 12

Statements

In this chapter the different kind of statements will be described one by one. Each of them will be described by means of:

- A syntax description in BNF.

- An informal semantics description.

- An example illustrating its usage.

12.1 Let Statements

Syntax: statement = let statement
                  | let be statement
                  | … ;

        let statement = '`let`', local definition, { ',', local definition },
                       '`in`', statement ;

        let be statement = '`let`', multiple bind, ['`be`', '`st`', expression], '`in`',
                         statement ;

        local definition = value definition
                     | function definition ;

        value definition = pattern, [':', type], '=', expression ;

    where the "function definition" component is described in chapter 5.

Semantics: The *let statement* and the *let-be-such-that statement* are similar to the corresponding *let* and *let-be-such-that expressions* except that the *in* part is a statement instead of an expression. Thus it can be explained as follows:

A simple *let statement* has the form:

```
    let p1 = e1, ..., pn = en in s
```

where p1, ..., pn are patterns, e1, ..., en are expressions which match the cor-
responding patterns pi, and s is a statement, of any type, involving the pattern identifiers
of p1, ..., pn. It denotes the evaluation of the statement s in the context in which the
patterns p1, ..., pn are matched against the corresponding expressions e1, ...,
en.

More advanced let statements can also be made by using local function definitions. The
semantics of doing that is simply that the scope of such locally defined functions is restricted
to the body of the let statement.

A *let-be-such-that statement* has the form

```
    let mb be st e in s
```

where mb is a multi-binding of one or more patterns (mostly just one pattern) to a set value
(or a type), e is a boolean expression, and s is a statement, involving the pattern identifiers of
the patterns from mb. The **be st** e part is optional. The expression denotes the evaluation
of the statement s in the context where all the patterns from mb has been matched against
an element in the set (or type) from mb[1]. If the **be st** expression e is present, only such
bindings where e evaluates to true in the matching context are used.

Examples: An example of a **let be st** statement is provided in the operation GroupWinner
from the class GroupPhase which returns the winning team in a given group:

```
GroupWinner : GroupName ==> Team
GroupWinner (gp) ==
  let sc in set gps(gp) be st
      forall sc' in set} gps(gp) \ {sc} &
        (sc.points > sc'.points) or
        (sc.points = sc'.points and sc.won > sc'.won)
  in
    return sc.team
```

The companion operation GroupRunnerUp gives an example of a simple let statement as
well:

```
GroupRunnerUp_expl : GroupName ==> Team
GroupRunnerUp_expl (gp) ==
```

---

[1]Remember that only the set bindings can be executed by means of the VDM interpreters.

```
    def t = GroupWinner(gp)
in let sct = iota sc in set gps(gp) & sc.team = t
   in
      let sc in set gps(gp) \ {sct} be st
          forall sc' in set gps(gp) \ {sc,sct} &
          (sc.points > sc'.points) or
          (sc.points = sc'.points and sc.won > sc'.won)
      in
         return sc.team
```

Note the use of the **def** statement (section 12.2) here; this is used rather than a **let** statement since the right-hand side is an operation call, and therefore is not an expression.

12.2   The Define Statement

Syntax: statement = …
$\qquad$ | $\quad$ def statement
$\qquad$ | $\quad$ … ;

$\quad$ def statement = '**def**', equals definition,
$\qquad\qquad\qquad$ { ';', equals definition },[';'], '**in**',
$\qquad\qquad\qquad$ statement ;

$\quad$ equals definition = pattern bind, '=', expression ;

Semantics: A *define statement* has the form:

```
    def pb1 = e1;
       ...
       pbn = en
    in
      s
```

The *define statement* corresponds to a *define expression* except that it is also allowed to use operation calls on the right-hand sides. Thus, operations that change the state can also be used here, and if there are more than one definition they are evaluated in the order in which they are presented. It denotes the evaluation of the statement s in the context in which the patterns (or binds) pb1, ..., pbn are matched against the values returned by the corresponding expressions or operation calls e1, ..., en[2].

---

[2]If binds are used it simply means that the values which can match the pattern are further constrained by the type or set expression as it is explained in section 7.

Examples: Given the following sequences:

```
secondRoundWinners = [<A>,<B>,<C>,<D>,<E>,<F>,<G>,<H>];
secondRoundRunnersUp = [<B>,<A>,<D>,<C>,<F>,<E>,<H>,<G>]
```

The operation `SecondRound`, in VDM++ from class `GroupPhase` returns the sequence of pairs representing the second round games gives an example of a **def** statement:

```
SecondRound : () ==> seq of (Team * Team)
SecondRound () ==
  def winners = { gp |-> GroupWinner(gp)
                  | gp in set dom gps };
      runners_up = { gp |-> GroupRunnerUp(gp)
                     | gp in set dom gps}
  in
    return ([mk_(winners(secondRoundWinners(i)),
                 runners_up(secondRoundRunnersUp(i)))
             | i in set {1,...,8}])
```

12.3 The Block Statement

Syntax: statement = …
                  | block statement
                  | … ;

      block statement = ' (', { dcl statement },
                     statement, { ';', statement }, [';'], ')' ;

      dcl statement = '**dcl**', assignment definition,
                 { ',', assignment definition }, ';' ;

      assignment definition = identifier, ':', type, [':=', expression] ;

Semantics: The *block statement* corresponds to block statements from traditional high-level imperative programming languages. It enables the use of locally defined variables (by means of the declare statement) which can be modified inside the body of the block statement. It simply denotes the ordered execution of what the individual statements prescribe. The first statement in the sequence that returns a value causes the evaluation of the sequence statement to terminate. This value is returned as the value of the block statement. If none of the statements in the block returns a value, the evaluation of the block statement is terminated

when the last statement in the block has been evaluated. When the block statement is left the values of the local variables are discharged. Thus, the scope of these variables is simply inside the block statement.

Examples: In the context of a VDM-SL state definition

```
state St of
  x: nat
  y: nat
  l: seq1 of nat
end
```

or in the context of a VDM++ instance variables

```
instance variables
  x: nat;
  y: nat;
  l: seq1 of nat;
```

the operation `Swap` uses a block statement to swap the values of variables `x` and `y`:

```
Swap : () ==> ()
Swap () ==
  (dcl temp: nat := x;
   x := y;
   y := temp
  )
```

12.4 The Assignment Statement

Syntax: statement = …
                    | general assign statement
                    | … ;

general assign statement = assign statement
                    | multiple assign statement ;

assign statement = state designator, ':=', expression ;

state designator = name
                   | field reference
                   | map or sequence reference ;

field reference = state designator, '.', identifier ;

map or sequence reference = state designator, '(', expression, ')' ;

multiple assign statement = '**atomic**', '(' assign statement, ';',
                              assign statement,
                              [{ ';', assign statement }] ')' ;

Semantics: The *assignment statement* corresponds to a generalisation of assignment statements from traditional high level programming languages. It is used to change the value of the global or local state. Thus, the assignment statement has side-effects on the state. However, in order to be able to simply change a part of the state, the left-hand side of the assignment can be a state designator. A state designator is either simply the name of a variable, a reference to a field of a variable, a map reference of a variable, or a sequence reference of a variable. In this way it is possible to change the value of a small component of the state. For example, if a state component is a map, it is possible to change a single entry in the map.

An assignment statement has the form:

```
sd := ec
```

where `sd` is a state designator, and `ec` is either an expression or a call of an operation. The assignment statement denotes the change to the given state component described at the right-hand side (expression or operation call). If the right-hand side is a state changing operation then that operation is executed (with the corresponding side effect) before the assignment is made.

Multiple assignment is also possible. This has the form:

```
atomic (sd1 := ec1;
          ...;
          sdN := ecN
        )
```

All of the expressions or operation calls on the right hand sides are executed or evaluated, and then the results are bound to the corresponding state designators. The right-hand sides are executed in the order given in the statement, and normal invariant processing and thread switching and statement durations can occur. But once all of the right-hand values have been obtained, they are assigned to the left-hand variables in one atomic step, which occurs

without invariant checking, thread switching or extra duration. It is as if the statement is as follows:

```
let t1 = ec1,
    ...
    tN = ecN in
(
    -- turn off invariants, threading and durations
    sd1 := t1;
    ...
    sdN := tN;
    -- turn on invariants, threading and durations
    -- and check that invariants hold.
);
```

Examples: The operation in the previous example (Swap) illustrated normal assignment. The operation Win_sd, a refinement of Win on page 95 illustrates the use of state designators to assign to a specific map key:

```
Win_sd : Team * Team ==> ()
Win_sd (wt,lt) ==
  let gp in set dom gps be st
      {wt,lt} subset {sc.team | sc in set gps(gp)}
  in
    gps(gp) := { if sc.team = wt
                 then mu(sc, won |-> sc.won + 1,
                            points |-> sc.points + 3)
                 elseif sc.team = lt
                 then mu(sc, lost |-> sc.lost + 1)
                 else} sc
               | sc in set gps(gp)}
pre exists gp in set dom gps &
              {wt,lt} subset {sc.team | sc in set gps(gp)}
```

The operation SelectionSort is a state based version of the function selection_sort on page 46. It demonstrates the use of state designators to modify the contents of a specific sequence index, using the VDM-SL state St or the VDM++ instance variables defined on page 105.

```
functions
```

```
min_index : seq1 of nat -> nat
min_index(l) ==
  if len l = 1
  then 1
  else let mi = min_index(tl l)
       in if l(mi+1) < hd l
          then mi+1
          else 1

operations

SelectionSort : nat ==> ()
SelectionSort (i) ==
  if i < len l
  then (dcl temp: nat;
        dcl mi : nat := min_index(l(i,...,len l)) + i - 1;
        temp := l(mi);
        l(mi) := l(i);
        l(i) := temp;
        SelectionSort(i+1)
       );
```

The following VDM++ example illustrates multiple assignment.

```
class C

instance variables
size : nat;
l : seq of nat;
inv size = len l

operations
add1 : nat ==> ()
add1 (x) ==
( l := [x] ^ l;
  size := size + 1);

add2 : nat ==> ()
add2 (x) ==
  atomic (l := [x] ^ l;
          size := size + 1)
```

```
end C
```

Here, in add1 the invariant on the class's instance variables is broken, whereas in add2 using the multiple assignment, the invariant is preserved.

12.5   Conditional Statements

Syntax: statement =   …
                | if statement
                | cases statement
                | … ;

if statement = '**if**', expression, '**then**', statement,
                 { elseif statement }, ['**else**', statement] ;

elseif statement = '**elseif**', expression, '**then**', statement ;

cases statement = '**cases**', expression, ':',
                    cases statement alternatives,
                    [',', others statement], '**end**' ;

cases statement alternatives = cases statement alternative,
                                 { ',', cases statement alternative } ;

cases statement alternative = pattern list, '->', statement ;

others statement = '**others**', '->', statement ;

Semantics: The semantics of the *if statement* corresponds to the *if expression* described in section 6.4 except for the alternatives which are statements (and that the **else** part is optional)[3].

The semantics for the *cases statement* corresponds to the *cases expression* described in section 6.4 except for the alternatives which are statements.

Examples: Assuming functions clear_winner and winner_by_more_wins and operation RandomElement with the following signatures:

```
clear_winner : set of Score -> bool
winner_by_more_wins : set of Score -> bool
RandomElement : set of Team ==> Team
```

---

[3]If the **else** part is omitted semantically it is like using **else skip**.

then the operation `GroupWinner_if` demonstrates the use of a nested if statement (the iota expression is presented on page 52):

```
GroupWinner_if : GroupName ==> Team
GroupWinner_if (gp) ==
  if clear_winner(gps(gp))
  -- return unique score in gps(gp) which has more points
  -- than any other score
  then return ((iota sc in set gps(gp) &
                  forall sc' in set gps(gp) \ {sc} &
                    sc.points > sc'.points).team)
  elseif winner_by_more_wins(gps(gp))
  -- return unique score in gps(gp) with maximal points
  -- & has won more than other scores with maximal points
  then return ((iota sc in set gps(gp) &
          forall sc' in set gps(gp) \ {sc} &
            (sc.points > sc'.points) or
            (sc.points = sc'.points and
             sc.won > sc'.won)).team)
  -- no outright winner, so choose random score
  -- from joint top scores
  else RandomElement ( {sc.team | sc in set gps(gp) &
                          forall sc' in set gps(gp) &
                            sc'.points <= sc.points} );
```

Alternatively, we could use a cases statement with match value patterns for this operation:

```
GroupWinner_cases : GroupName ==> Team
GroupWinner_cases (gp) ==
  cases true:
    (clear_winner(gps(gp))) ->
        return ((iota sc in set gps(gp) &
                  forall sc' in set gps(gp) \ {sc} &
                    sc.points > sc'.points).team),
    (winner_by_more_wins(gps(gp))) ->
        return ((iota sc in set gps(gp) &
                  forall sc' in set gps(gp) \ {sc} &
                    (sc.points > sc'.points) or
                    (sc.points = sc'.points and
                       sc.won > sc'.won)).team),
    others -> RandomElement ( {sc.team | sc in set gps(gp) &
```

110

```
                                          forall sc' in set gps(gp) &
                                            sc'.points <= sc.points} )
end
```

12.6   For-Loop Statements

Syntax: statement = …
  | sequence for loop
  | set for loop
  | index for loop
  | … ;

sequence for loop = '**for**', pattern bind, '**in**', expression,
      '**do**', statement ;

set for loop = '**for**', '**all**', pattern, '**in set**', expression,
      '**do**', statement ;

index for loop = '**for**', identifier, '=', expression, '**to**', expression,
      ['**by**', expression], '**do**', statement ;

Semantics: There are three kinds of *for-loop statements*. The for-loop using an index is known
from most high-level programming languages. In addition, there are two for-loops for
traversing sets and sequences. These are especially useful if access to all elements from
a set (or sequence) is needed one by one.

An *index for-loop statement* has the form:

```
for id = e1 to e2 by e3 do
  s
```

where id is an identifier, $e1$ and $e2$ are integer expressions indicating the lower and upper
bounds for the loop, $e3$ is an integer expression indicating the step size, and s is a statement
where the identifier id can be used. It denotes the evaluation of the statement s as a sequence
statement where the current context is extended with a binding of id. Thus, the first time
s is evaluated id is bound to the value returned from the evaluation of the lower bound $e1$
and so forth until the upper bound is reached i.e. until $s > e2$. Note that $e1, e2$ and $e3$
are evaluated before entering the loop.

A *set for-loop statement* has the form:

```
      for all e in set S do
      s
```

where S is a set expression. The statement s is evaluated in the current environment extended with a binding of e to subsequent values from the set S.

A *sequence for-loop statement* has the form:

```
      for e in l do
      s
```

where l is a sequence expression. The statement s is evaluated in the current environment extended with a binding of e to subsequent values from the sequence l.

Examples: The operation Remove demonstrates the use of a *sequence-for* loop to remove all occurences of a given number from a sequence of numbers:

```
Remove : (seq of nat) * nat ==> seq of nat
Remove (k,z) ==
(dcl nk : seq of nat := [];
 for elem in k do
   if elem <> z
   then nk := nk^[elem];
 return nk
);
```

A *set-for* loop can be exploited to return the set of winners of all groups:

```
GroupWinners: () ==> set of Team
GroupWinners () ==
(dcl winners : set of Team := {};
 for all gp in set dom gps do
   (dcl winner: Team := GroupWinner(gp);
    winners := winners union {winner}
   );
 return winners
);
```

An example of a *index-for* loop is the classic bubblesort algorithm:

```
BubbleSort : seq of nat ==> seq of nat
BubbleSort (k) ==
  (dcl sorted_list : seq of nat := k;
   for i = len k to 1 by -1 do
     for j = 1 to i-1 do
       if sorted_list(j) > sorted_list(j+1)
       then (dcl temp:nat := sorted_list(j);
              sorted_list(j) := sorted_list(j+1);
              sorted_list(j+1) := temp
            );
   return sorted_list
   )
```

12.7 The While-Loop Statement

Syntax: statement = …
  | while loop
  | … ;

while loop = '**while**', expression, '**do**', statement ;

Semantics: The semantics for the *while statement* corresponds to the while statement from traditional programming languages. The form of a *while loop* is:

```
while e do
  s
```

where e is a boolean expression and s a statement. As long as the expression e evaluates to **true** the body statement s is evaluated.

Examples: The *while loop* can be illustrated by the following example which uses Newton's method to approximate the square root of a real number r within relative error e.

```
SquareRoot : real * real ==> real
SquareRoot (r,e) ==
  (dcl x:real := 1,
       nextx: real := r;
   while abs (x - nextx) >= e * x do
     ( x := nextx;
       nextx := ((r / x) + x) / 2;
```

```
      );
    return nextx
  );
```

12.8   The Nondeterministic Statement

Syntax: statement = ...
                    | nondeterministic statement
                    | ... ;

            nondeterministic statement = '||', '(', statement,
                                            { ',', statement }, ')' ;

Semantics: The *nondeterministic statement* has the form:

```
    || (stmt1, stmt2, ..., stmtn)
```

and it represents the execution of the component statements `stmti` in an arbitrary (non-deterministic) order. However, it should be noted that the component statements are not executed simultaneously. Notice that the VDM interpreters will use an underdetermined[4] semantics even though this construct is called a non-deterministic statement.

Examples: Using the VDM-SL state definition

```
state St of
  x:nat
  y:nat
  l:seq1 of nat
end
```

or the VDM++ instance variables

```
instance variables
  x:nat;
  y:nat;
  l:seq1 of nat;
```

---

[4]Even though the user of the VDM interpreters does not know the order in which these statements are executed they are always executed in the same order unless the seed option is used.

we can use the non-deterministic statement to effect a bubble sort:

```
Sort: () ==> ()
Sort () ==
  while x < y do
    ||(BubbleMin(), BubbleMax());
```

Here `BubbleMin` "bubbles" the minimum value in the subsequence `l(x,...,y)` to the head of the subsequence and `BubbleMax` "bubbles" the maximum value in the subsequence `l(x,...,y)` to the last index in the subsequence. `BubbleMin` works by first iterating through the subsequence to find the index of the minimum value. The contents of this index are then swapped with the contents of the head of the list, `l(x)`.

```
BubbleMin : () ==> ()
BubbleMin () ==
  (dcl z:nat := x;
   dcl m:\keyw{nat} := l(z);
   -- find min val in l(x..y)
   for i = x to y do
     if l(i) < m
     then ( m := l(i);
            z := i);
   -- move min val to index x
   (dcl temp:nat;
    temp := l(x);
    l(x) := l(z);
    l(z) := temp;
    x := x+1));
```

`BubbleMax` operates in a similar fashion. It iterates through the subsequence to find the index of the maximum value, then swaps the contents of this index with the contents of the last element of the subsequence.

```
BubbleMax : () ==> ()
BubbleMax () ==
  (dcl z:nat := x;
   dcl m:nat := l(z);
   -- find max val in l(x..y)
   for i = x to y do
     if l(i) > m
     then ( m := l(i);
```

```
                z := i);
      -- move max val to index y
     (dcl temp:nat;
      temp := l(y);
      l(y)  := l(z);
      l(z)  := temp;
      y := y-1));
```

12.9 The Call Statement

Syntax: statement = …
           | call statement
           | … ;

For VDM-SL call statements are defined as:

    call statement = name, ' (', [expression list], ') ' ;

For VDM++ and VDM-RT call statements are defined as:

    call statement = [object designator, '.'], name,
                  ' (', [expression list], ') ', ;

    object designator = name
                   | self expression
                   | new expression
                   | object field reference
                   | object apply ;

    object field reference = object designator, '.', identifier ;

    object apply = object designator, ' (', [expression list], ') ' ;

Semantics: In VDM-SL the *call statement* has the form:

```
opname(param1, param2, ..., paramn)
```

In VDM++ and VDM-RT the *call statement* can additionally have the form:

```
object.opname(param1, param2, ..., paramn)
```

The *call statement* calls an operation, opname, (in a VDM++ and VDM-RT context it can also be on a specific object, object), and returns the result of evaluating the operation. Because operations can manipulate global variables a *call statement* does not necessarily have to return a value as function call do.

In VDM++ and VDM-RT if an *object designator* is specified it must yield an object reference to an object of a class in which the operation opname is defined, and then the operation must be specified as public. If no *object designator* is specified the operation will be called in the current object. If the operation is defined in a superclass, it must have been defined as public or protected.

Examples: In VDM-SL the operation ResetStack given below does not have any parameter and does not return a value whereas the operation PopStack returns the top element of the stack.

```
    ResetStack();
    ...
    top := PopStack();
```

where PopStack could be defined as:

```
    PopStack: () ==> Elem
    PopStack() ==
      def res = hd stack in
        (stack := tl stack;
         return res)
    pre stack <> []
    post stack~ = [RESULT] ^ stack
```

where stack is a global variable.

In VDM++ and VDM-RT this Stack example can be made like:

```
  class Stack

  instance variables
    stack: seq of Elem := [];

  operations

    public Reset: () ==> ()
    Reset() ==
      stack := [];
```

```
   public Pop: () ==> Elem
   Pop() ==
     def res = hd stack in
       (stack := tl stack;
        return res)
   pre stack <> []
   post stack~ = [RESULT] ^ stack

 end Stack
```

In the example the operation Reset does not have any parameters and does not return a value whereas the operation Pop returns the top element of the stack. The stack could be used as follows:

```
 ( dcl stack := new Stack();
   stack.Reset();
   ....
   top := stack.Pop();
 )
```

Inside class Stack the operations can be called as shown below:

```
 Reset();
 ....
 top := Pop();
```

Or using the **self** reference:

```
 self.Reset();
 top := self.Pop();
```

12.10   The Return Statement

Syntax: statement = …
              | return statement
              | … ;

          return statement = 'return', [expression] ;

Semantics: The *return statement* returns the value of an expression inside an operation. The value is evaluated in the given context. If an operation does not return a value, the expression must be omitted. A *return statement* has the form:

```
    return e
```

or

```
    return
```

where expression `e` is the return value of the operation.

Examples: In the following example `OpCall` is an operation call whereas `FunCall` is a function call. As the *if statement* only accepts statements in the two branches `FunCall` is "converted" to a statement by using the *return statement*.

```
    if test
    then OpCall()
    else return FunCall()
```

For instance in VDM++, we can extend the `stack` class from the previous section with an operation which examines the top of the stack:

```
public Top : () ==> Elem
Top() ==
  return (hd stack)
pre stack <> [];
```

12.11   Exception Handling Statements

Syntax: statement = ...
                    | always statement
                    | trap statement
                    | recursive trap statement
                    | exit statement
                    | ... ;

always statement = '**always**', statement, '**in**', statement ;

```
trap statement =  'trap',  pattern bind,  'with',  statement,  'in',
                       statement ;

recursive trap statement =  'tixe',  traps,  'in',  statement ;

traps =  '{',  pattern bind,  '|->',  statement,
           { ',',  pattern bind,  '|->',  statement }, '}' ;

exit statement =  'exit', [ expression ] ;
```

Semantics: The exception handling statements are used to control exception errors in a specification. This means that we have to be able to signal an exception within a specification. This can be done with the *exit statement*, and has the form:

```
exit e
```

or

```
exit
```

where e is an expression which is optional. The expression e can be used to signal what kind of exception is raised.

The *always statement* has the form:

```
always s1 in
s2
```

where s1 and s2 are statements. First statement s2 is evaluated, and regardless of any exceptions raised, statement s1 is also evaluated. The result value of the complete *always statement* is determined by the evaluation of statement s1: if this raises an exception, this value is returned, otherwise the result of the evaluation of statement s2 is returned.

The *trap statement* only evaluates the handler statement, s1, when certain conditions are fulfilled. It has the form:

```
trap pat with s1 in s2
```

where pat is a pattern or bind used to select certain exceptions, s1 and s2 are statements. First, we evaluate statement s2, and if no exception is raised, the result value of the complete *trap statement* is the result of the evaluation of s2. If an exception is raised, the value of

s2 is matched against the pattern `pat`. If there is no matching, the exception is returned as result of the complete *trap statement*, otherwise, statement s1 is evaluated and the result of this evaluation is also the result of the complete *trap statement*.

The *recursive trap statement* has the form:

```
tixe {
  pat1 |-> s1,
  ...
  patn |-> sn
} in s
```

where `pat1, ..., patn` are patterns or binds, `s, s1, ..., sn` are statements. First, statement `s` is evaluated, and if no exception is raised, the result is returned as the result of the complete *recursive trap statement*. Otherwise, the value is matched in order against each of the patterns `pati`. When a match cannot be found, the exception is returned as the result of the *recursive trap statement*. If a match is found, the corresponding statement `si` is evaluated. If this does not raise an exception, the result value of the evaluation of `si` is returned as the result of the *recursive trap statement*. Otherwise, the matching starts again, now with the new exception value (the result of the evaluation of `si`).

Examples: In many programs, we need to allocate memory for a single operation. After the operation is completed, the memory is not needed anymore. This can be done with the *always statement*:

```
( dcl mem : Memory;
  always Free(mem) in
  ( mem := Allocate();
    Command(mem, ...)
  )
)
```

In the above example, we cannot act upon a possible exception raised within the body statement of the *always statement*. By using the *trap statement* we can catch these exceptions:

```
trap pat with ErrorAction(pat) in
( dcl mem : Memory;
  always Free(mem) in
  ( mem := Allocate();
    Command(mem, ...)
  )
)
```

Now all exceptions raised within the *always statement* are captured by the *trap statement*. If we want to distinguish between several exception values, we can use either nested *trap statements* or the *recursive trap statement*:

```
DoCommand : () ==> int
DoCommand () ==
( dcl mem : Memory;
  always Free(mem) in
  ( mem := Allocate();
    Command(mem, ...)
  )
);

Example : () ==> int
Example () ==
tixe
{ <NOMEM> |-> return -1,
  <BUSY>  |-> DoCommand(),
  err     |-> return -2 }
in
  DoCommand()
```

In operation `DoCommand` we use the *always statement* in the allocation of memory, and all exceptions raised are captured by the *recursive trap statement* in operation `Example`. An exception with value `<NOMEM>` results in a return value of −1 and no exception raised. If the value of the exception is `<BUSY>` we try to perform the operation `DoCommand` again. If this raises an exception, this is also handled by the *recursive trap statement*. All other exceptions result in the return of the value −2.

12.12   The Error Statement

Syntax: statement = ...
            | error statement
            | ... ;

            error statement = '**error**' ;

Semantics: The *error statement* corresponds to the undefined expression. It is used to state explicitly that the result of a statement is undefined and because of this an error has occurred. When an *error statement* is evaluated the VDM interpreters will terminate the execution of the specification and report that an *error statement* was evaluated.

Pragmatically use of error statements differs from pre-conditions as was the case with undefined expressions: use of a pre-condition means it is the caller's responsibility to ensure that the pre-condition is satisfied when the operation is called; if an error statement is used it is the called operation's responsibility to deal with error handling.

Examples: The operation `SquareRoot` on page 114 does not exclude the possibility that the number to be square rooted might be negative. If we do not wish this to be a pre-condition we can remedy this in the operation `SquareRootErr`:

```
SquareRootErr : real * real ==> real
SquareRootErr (r,e) ==
  if r < 0
  then error
  else
    (dcl x:real := 1;
     dcl nextx:real := r;
     while abs (x - nextx) >= e * x do
       ( x := nextx;
         nextx := ((r / x) + x) / 2;
       );
     return nextx
    )
```

12.13 The Identity Statement

Syntax: statement = …
                          | identity statement ;

identity statement = '**skip**' ;

Semantics: The *identity statement* is used to signal that no evaluation takes place.

Examples: In the operation `Remove` in section 12.6 the behaviour of the operation within the **for** loop if `elem=z` is not explicitly stated. `Remove2` below does this.

```
Remove2 : (seq of nat) * nat ==> seq of nat
Remove2 (k,z) ==
  (dcl nk : seq of nat := [];
   for elem in k do
     if elem <> z
     then nk := nk^[elem]
```

```
        else skip;
    return nk
);
```

Here, we explicitly included the **else**-branch to illustrate the *identity statement*, however, in most cases the **else**-branch will not be included and the *identity statement* is implicitly assumed.

12.14 Start and Start List Statements (VDM++ and VDM-RT)

Syntax: statement = …
                | start statement
                | start list statement ;

         start statement = '**start**', '(', expression, ')' ;

         start list statement = '**startlist**', '(', expression, ')' ;

Semantics: The *start* and *start list* statements have the form:

```
start(aRef)
startlist(aRef_s)
```

If a class description includes a thread (see chapter 15), each object created from this class will have the ability to operate as a stand-alone virtual machine, or in other terms: the object has its own processing capability. In this situation, a *new expression* creates the 'process' leaving it in a waiting state. For such objects VDM++ and VDM-RT has a mechanism to change the waiting state into an active state[5] in terms of a predefined operation, which can be invoked through a *start statement*.

The explicit separation of object creation and start provides the possibility to complete the initialisation of a (concurrent) system *before* the objects start exhibiting their described behaviour, in this way avoiding problems that may arise when objects are referred to that are not yet created and/or connected.

A syntactic variant of the start statement is available to start up a number of active objects in arbitrary order: the *start list statement*. The parameter aRef_s to **startlist** must be a set of object references to objects instantiated from classes containing a thread.

Examples: Consider the specification of an operating system. A component of this would be the daemons and other processes started up during the boot sequence. From this perspective, the following definitions are relevant:

---

[5]When an object is in an active state, its behaviour can be described using a thread (see chapter 15).

```
types

   runLevel = nat;

   Process = Kerneld | Ftpd | Syslogd | Lpd | Httpd

instance variables
   pInit : map runLevel to set of Process
```

where `Kerneld` is an object reference type specified elsewhere, and similarly for the other processes listed.

We can then model the boot sequence as an operation:

```
bootSequence : runLevel ==> ()
bootSequence(rl) ==
  for all p in set pInit(rl) do
    start(p);
```

Alternatively we could use the **startlist** statement here:

```
bootSequenceList : runLevel ==> ()
bootSequenceList(rl) ==
  startlist(pInit(rl))
```

12.15 Stop and Stop List Statements (VDM++ and VDM-RT)

Syntax: statement = …
| stop statement
| stop list statement ;

stop statement = '**stop**', '(', expression, ')' ;

stop list statement = '**stoplist**', '(', expression, ')' ;

Semantics: The *stop* and *stop list* statements have the following form:

```
stop(aRef)
stoplist(aRef_s)
```

The stop statement can be used to terminate a thread or to prevent a periodic thread from being invoked again, whereby a thread is identified by the object reference *aRef*. The stop list statement is available to stop a number of active objects, in arbitrary order. The parameter `aRef_s` to **stoplist** must be a set of object references. A few additional remarks can be made:

1. '**self**' can be passed as an argument to `stop` or `stoplist`, effectively killing the current thread

2. Note that *thread identifiers* can *not* be passed as an argument

3. in contrast to `start` and `startlist`, `stop` and `stoplist` will block until the thread(s), identified by their object identifiers that are passed as an argument, are dead; if '**self**' is part of the set, it will be stopped last

4. passing the object identifier of a task that has never been started or is already dead, will yield a run-time error

5. objects can be restarted by calling `start` and `startlist` again, which will create a fresh and unique *thread identifier* for each thread.

6. it is only possible to stop threads that are running on the same CPU

12.16   The Specification Statement

Syntax: statement = ...
                                    | specification statement ;

          specification statement = '[', implicit operation body, ']' ;

Semantics: The specification statement can be used to describe a desired effect a statement in terms of a pre- and a post-condition. Thus, it captures the abstraction of a statement, permitting it to have an abstract (implicit) specification without being forced to an operation definition. The specification statement is equivalent with the body of an implicitly defined operation (see chapter 11). Thus specification statements can not be executed.

Examples: We can use a specification statement to specify a bubble maximum part of a bubble sort:

```
Sort2 : () ==> ()
Sort2 () ==
  while x < y do
    || (BubbleMin(),
        [ext wr l : seq1 of nat
              wr y : nat
```

126

```
            rd x : nat
       pre x < y
       post y < y~ and
             permutation(l~(x,...,y~),l(x,...,y~)) and
             forall i in set {x,...,y} & l(i) < l(y~)]
       )
```

(`permutation` is an auxiliary function taking two sequences which returns true iff one sequence is a permutation of the other.)

12.17 The Duration Statement (VDM-RT)

Syntax: statement = …
                   | duration statement ;

duration statement = '**duration**', '(', expression, ')', statement ;

Semantics: The duration statement is a runtime directive to the VDM interpreters telling it that when incrementing the internal clock for the enclosed statement, the value (an expression that must yield a natural number as its return type, otherwise causing a run-time error) given in the duration statement should be used instead of the increment which would normally be computed for that statement. Thus the duration statement provides a mechanism to override the VDM interpreter's default execution time computation. Note that the execution of the expression is done in zero time and cannot cause thread switches.

Example: First a simple example:

```
while n < 10 do
  duration(10) n := n + 1;
```

In this example, assuming that this loop is not executed in the context of an enclosing duration statement, on each iteration of the loop the VDM interpreters will increment its internal clock by 10 time units (nanoseconds), rather than computing the amount of time required to execute the statement `n := n + 1`.

If duration statements are nested, the outermost one takes precedence and the remainder are ignored. For instance

```
duration(30)
( n := 1;
  while n < 10 do
```

```
        duration(10) n := n + 1;
  )
```

The outer duration statement takes precedence, so assuming this is not executed in the context of an enclosing duration statement, the VDM interpreters would increment its internal clock by 30 time units when executing this statement.

Note that nesting can occur due to operation calls. Consider the following example:

```
op1 : nat ==> nat
op1(m) ==
  duration (20) return m + 1;

op2 : () ==> nat
op2() ==
( dcl n : nat := 3;
  duration(10)  n := op1(1);
  return n)
```

When executing op2, if the call to op1 is executed, the duration statement in op1 will be overridden by the duration statement in the environment of the call. Thus in op2 following execution of the statement n := op1(1); the internal clock is incremented by 10 time units only.

```
  ( n := 1; while n < 10 do duration(n) n := n + 1 )
```

The final example demonstrates the use of a general expression as the argument to the duration statement, whereby the consecutive executions of the body of the while loop take more time as n increases.

12.18   The Cycles Statement (VDM-RT)

Syntax: statement = …
                   | cycles statement ;

          cycles statement = '**cycles**', '(', expression, ')', statement ;

Semantics: The cycles statement is a runtime directive to the VDM interpreters telling it that when incrementing the internal clock for the enclosed statement, the value (an expression that must yield a natural number as its return type, otherwise causing a run-time error) given in the cycles statement should be used as an indication of how many clock cycles that the

enclosed statement should be incremented by instead of the increment which would normally be computed for that statement. Thus the cycles statement provides a mechanism to override the VDM interpreter's default execution time computation similar to the duration statement but in a way that is relative to the speed of the CPU that the computation is carried out on. Note that the execution of the expression is done in zero time and cannot cause thread switches.

Example: First a simple example:

```
while n < 10 do
  cycles(1000) n := n + 1;
```

In this example, assuming that this loop is not executed in the context of an enclosing cycles statement, on each iteration of the loop the VDM interpreters will increment its internal clock by the time it will take to process 1000 instructions on the given CPU (relative to its capacity), rather than computing the amount of time required to execute the statement n := n + 1.

If cycles statements are nested, the outermost one takes precedence and the remainder are ignored. For instance

```
cycles(3000)(
      n := 1;
      while n < 10 do
         cycles(1000) n := n + 1;
      )
```

The outer cycles statement takes precedence, so assuming this is not executed in the context of an enclosing cycles statement, the interpreter would increment its internal clock by the time it takes to process 3000 instructions on the given CPU when executing this statement.

Note that nesting can occur due to operation calls. Consider the following example:

```
    op1 : nat ==> nat
    op1(m) ==
      cycles (2000) return m + 1;

    op2 : () ==> nat
    op2() ==
    (dcl n : nat := 3;
     cycles(1000)  n := op1(1);
     return n)
```

When executing `op2`, if the call to `op1` is executed, the cycles statement in `op1` will be overridden by the cycles statement in the environment of the call. Thus in `op2` following execution of the statement `n := op1(1);` the internal clock is incremented by the time it takes to process 1000 instructions on the given `CPU` only.

```
( n := 1; while n < 10 do cycles(n) n := n + 1 )
```

The final example demonstrates the use of a general expression as the argument to the cycles statement, whereby the consecutive executions of the body of the while loop take more time as n increases.

Chapter 13

Top-level Specification in VDM

The top-level specification structure differs significantly between the VDM-SL appraoch and the VDM++ and VDM-RT approach. In VDM-SL the ISO standard prescribes a flat-language but here a modular extension is also enabled using imports and exprts primitives. In VDM++ and VDM-RT structuring is done using object-oriented classes that can inherit constructs between them controlled by access modifiers. These two different approaches are explained in the two sections in this chapter.

13.1   Top-level Specification in VDM-SL

In the previous chapters all the VDM-SL constructs such as types, expressions, statements, functions and operations have been described. A number of these constructs can constitute a top-level VDM-SL specification. A top-level specification can be created in two ways:

1. The specification is split into a number of modules which are specified separately, but can depend on each other.

2. The specification is specified in a flat manner, i.e. no modules are used (note that in VDM-10 it is possible to have access to standard modules also from a flat VDM-SL specification).

Thus, a complete specification, or document, has the following syntax.

Syntax: document = any module, { any module }
                           | definition block, { definition block } ;

            any module = module ;

13.1.1   A Flat Specification

As said, a flat specification does not use modules. This means that all constructs can be used throughout the specification. In the flat case, a document has a syntax of:

document = …
 | definition block, { definition block } ;

definition block = type definitions
 | state definition
 | value definitions
 | function definitions
 | operation definitions
 | traces definitions ;

Thus, a flat specification is made up of several *definition* blocks. However, only one state definition is allowed. The following is an example of a flat top-level specification:

```
values

  st1 = mk_St([3,2,-9,11,5,3])

state St of
  l:seq1 of nat
end

functions

  min_index : seq1 of nat -> nat
  min_index(l) ==
    if len l = 1
    then 1
    else let mi = min_index(tl l)
         in
           if l(mi+1) < hd l
           then mi+1
           else 1

operations

  SelectionSort : nat ==> ()
  SelectionSort (i) ==
    if i < len l
    then (dcl temp: nat;
          dcl mi : nat := min_index(l(i,...,len l)) + i - 1;

          temp := l(mi);
          l(mi) := l(i);
```

```
        l(i) := temp;
        SelectionSort(i+1)
      )
```

13.1.2 A Structured Specification

As an extension to the standard VDM-SL language, it is possible to structure an VDM-SL specification using modules. In this section, the use of modules to create the top-level specification will be described. With the structuring facilities offered by VDM-SL it is possible to:

- Export constructs from a module.

- Import constructs from a module.

- Rename constructs upon import.

- Define a state in a module.

The Layout of a Module

Before the actual facilities are described, the general layout of a module is described. A module consists of three parts: a *module declaration*, an *interface section*, and a *definitions section*. It is possible to leave out the definitions part in the early development of a module specification.

In the module declaration, the module is named. The name must be a unique module name within the complete specification. The second part, the interface section, defines the relation of a module with other modules and consists of two sections. These sections are:

- An *imports section*. In the imports section, all the constructs that are going to be used from other modules are described. If constructs are going to be renamed it has to be done in the imports section.

- An *exports section*. Here all the constructs that are going to be used in other modules are defined. If no exports section is present the module cannot be used from any other modules.

The third part of a module declaration, the definitions section, contains all the definitions of the module. Thus, in general, the syntax of a module is:

Syntax: module = '**module**', identifier, interface,
                         [module body], '**end**', identifier ;

           module body = '**definitions**', definition block, { definition block } ;

To illustrate the use of modules, the example flat top-level specification are rewritten with some minor modifications. Some unimportant parts of the flat specification are left out for clarity.

The Exports Section

Syntax: interface = [import definition list],
      export definition ;

    export definition = '**exports**', export module signature ;

    export module signature = '**all**'
            | export signature,
             { export signature } ;

    export signature = export types signature
          | values signature
          | export functions signature
          | operations signature ;

    export types signature = '**types**', type export,
            { ';', type export }, [';'] ;

    type export = ['**struct**'], name ;

    values signature = '**values**', value signature,
          { ';', value signature }, [';'] ;

    value signature = name list, ':', type ;

    export functions signature = '**functions**' function export,
             { ';', function export } ;

    function export = name list, [type variable list], ':',
         function type ;

    functions signature = '**functions**' function signature,
           { ';', function signature }, [';'] ;

    function signature = name list, ':', function type ;

    operations signature = '**operations**' operation signature,
           { ';', operation signature }, [';'] ;

    operation signature = name list, ':', operation type ;

Semantics: The exports section must be used to make constructs visible to other modules. Some or all of the defined constructs from a module can be exported. In the latter case, the keyword **all** is used. However, imported constructs are not exported from the module. If only part of the constructs are exported, the visible constructs with the appropriate signatures are stated.

Normally, if a construct is visible to another module, that construct can be considered to be defined inside the module. However, with types and operations there are some exceptions:

Types: If a type T is defined in module A and this type is also going to be used in module B, the type from module A has to be exported. This can be done in two ways:

1. The name of the type is exported.
2. The structure of the type is exported.

If only the name of the type is exported, the other module cannot create values of type T. This means that the exporting module (A) must provide functions and/or operations to directly create and manipulate values of type T by means of the constructors related to the representation of T.

If we export the structure of the type by using the keyword **struct**, the other module can create and manipulate values of type T (it can also use **mk_** keyword and the **is_** keyword for this type if it is a record type).

If the type also defines an invariant, the invariant predicate function is only exported if the structure of the type is exported.

Operations: In a module, a state that is global for the module can be defined. All operations within the module can manipulate that state. If operations are exported from a module, they manipulate the state in the exporting module, i.e. the state in the module where they are defined.

If an exported function or an operation defines a pre- and/or post-condition, the corresponding predicate functions (see chapter 5) are also exported.

Examples: Consider a model of a bank account. An account is characterised by the name of the holder, the account number, the bank branch at which the account is maintained, the balance, and an encrypted PIN code for the ATM card. We might model this as follows:

```
module BankAccount

exports types digit; account
        functions digval: digit -> nat;
                  withdrawal: account * real -> account;
                  isPin: account * nat -> bool;
                  requestWithdrawal: account * nat -> bool
definitions
```

```
types

  digit = nat
  inv d == d < 10;

  account::  holder : seq1 of char
             number : seq1 of digit
             branchcode : seq1 of digit
             balance: real
             epin: nat
  inv mk_account(holder, number, branchcode,-,-) ==
    len number = 8 and len branchcode = 6

functions

    digval : digit -> nat
    digval(d) == d;

    deposit: account * real -> account
    deposit(acc,r) ==
      mu(acc,balance |-> acc.balance + r);

    withdrawal : account * real -> account
    withdrawal (acc,r) ==
      mu(acc,balance |-> acc.balance - r);

    isPin : account * nat -> bool
    isPin(acc,ep) ==
      ep = acc.epin;

    requestWithdrawal : account * nat -> bool
    requestWithdrawal (acc,amt) ==
      acc.balance > amt

end BankAccount
```

In this module we export two types and five functions. Note that since we have enumerated the entities we are exporting, but have not exported `digit` or `account` using the **struct** keyword, the internals of `account` values may not be accessed by other modules, neither may the invariant for `digit`. If such access is necessary, the types should be exported with the **struct** keyword, or all constructs in the module should be exported using the

exports all clause.

The module `Keypad` given below models the keypad interface of an ATM machine. The state variable maintains a buffer of data typed at the keypad by the user.

```
module Keypad

imports
from BankAccount types digit

exports all

definitions

state buffer of
  data : seq of BankAccount'digit
end

operations

    DataAvailable : () ==> bool
    DataAvailable () ==
      return(data <> []);

    ReadData : () ==> seq of BankAccount'digit
    ReadData () ==
      return (data);

    WriteData : seq of BankAccount'digit ==> ()
    WriteData (d) ==
      data := data^d

end Keypad
```

In this module all constructs are exported. Since the only entities defined are the state and operations on it, this means that all of the operations may be accessed by an importing module. The state is not accessible to importing modules, but remains private to this module. However the state constructor **mk_**Keypad`buffer is accessible.

The Imports Section

Syntax: interface = [import definition list],

export definition ;

import definition list = '**imports**', import definition,
{ ',', import definition } ;

import definition = '**from**', identifier, import module signature ;

import module signature = '**all**'
| import signature,
{ import signature } ;

import signature = import types signature
| import values signature
| import functions signature
| import operations signature ;

import types signature = '**types**', type import,
{ ';', type import }, [';'] ;

type import = name, ['**renamed**', name]
| type definition, ['**renamed**', name] ;

import values signature = '**values**', value import,
{ ';', value import }, [';'] ;

value import = name, [':', type], ['**renamed**', name] ;

import functions signature = '**functions**', function import,
{ ';', function import }, [';'] ;

function import = name, [[type variable list], ':', function type],
['**renamed**', name] ;

import operations signature = '**operations**', operation import,
{ ';', operation import }, [';'] ;

operation import = name, [':', operation type],
['**renamed**', name] ;

Semantics: The imports section is used to state what constructs are used from other modules with the restriction that only visible constructs can be imported. If all the visible constructs from a module are going to be used, the keyword **all** is used, unless one or more constructs are going to be renamed. With renaming, an imported construct is given a new name which can be used instead of the original name preceded by the exporting module name. In general this has the form:

```
    name renamed new_name
```

where `name` is the name of the imported construct, and `new_name` is the new name for the construct. This way, more meaningful names can be given to constructs. Note that in the importing module it is not possible to refer to `DefModule'name` (where `DefModule` is the name of the defining module) any longer but only to `newname`.

It is possible to include type information in the imports section, such that this information will only be used by the static semantics check of the complete module. If no type information is given, the static semantics can also find this information in the exporting module.

When a type which has been exported with the **struct** keyword (with its structure) is imported the importing module may only make use of this structure if it repeats the type definition from the exporting module in its type import. In case such a type is a composite type and it is also renamed this has the consequence that the tag is renamed as well.

Examples: We can model an ATM card as consisting of a card number and an expiry date. This requires the `digit` type defined in the module `BankAccount`. It also uses the function `digval` from the same module.

```
module ATMCard

imports
from BankAccount types digit
                 functions digval renamed atmc_digval

exports all

definitions

types

    digit = BankAccount'digit;

    atmc:: cardnumber : seq1 of digit
          expiry : digit * digit * digit * digit
    inv mk_atmc(cardnumber, mk_(m1,m2,-,-)) ==
        atmc_digval(m1) * 10 + atmc_digval(m2) <= 12 and
        len cardnumber >= 8

functions

    getCardnumber : atmc -> seq1 of digit
```

```
      getCardnumber (atmc) ==
        atmc.cardnumber

end ATMCard
```

Here the invariant on the type `atmc` states that expiry dates must represent valid dates, and card numbers must be at least 8 digits long. Note that since `digit` is not exported with the **struct** keyword from the module `BankAccount`, we cannot access the invariant for `digit` in module `ATMCard`. However this notwithstanding, all values of type `digit` manipulated in `ATMCard` must satisfy the invariant.

13.2   Top-level Specification in VDM++ and VDM-RT

In the previous chapters VDM constructs such as types, expressions, statements, functions and operations have been described. A number of these constructs can constitute the definitions inside a class definition. A top-level specification, or document, is composed by one or more class definitions. Note that only in VDM-RT it is possible to have a **system** class.

Syntax: document = class | system ,{ class | system } ;

13.3   System (VDM-RT)

In order to be able to describe distributed systems in VDM-RT includes a notion of a system that describes how different parts of the system modelled are deployed to different Core Processing Units (CPUs) and communication busses connecting the CPUs together. Syntactically the system is described exactly like ordinary classes described below in Section 13.3.1, except that the keyword "**system**" instead of the keyword "**class**".

Syntax: system = '**system**', identifier,
                 [class body],
                 '**end**', identifier ;

         class body = definition block, { definition block } ;

         definition block = type definitions
                           | value definitions
                           | function definitions
                           | operation definitions
                           | instance variable definitions
                           | synchronization definitions
                           | thread definitions ;

Semantics: Each system description has the following parts:

- A system header with the system name.
- An optional *system body*.
- A system tail.

The system name as given in the system header is the defining occurrence of the name of the class. A system name is globally visible, i.e. visible in all other classes/systems in the specification.

The system name in the class header must be the same as the system name in the system tail. Furthermore, defining system names must be unique throughout the specification.

The special thing about the system is that it can make use of special implicitly defined classes called CPU and BUS. It is not possible to create instances of the system, but instances made of CPU and BUS will be created at initialisation time. Note that CPU and BUS cannot be used outside the system definition.

The instances of CPU and BUS must be made as instance variables and the definition must use constructors. The constructor for the CPU class takes two parameters: the first one indicate the primary scheduling policy used for the CPU whereas the second parameter provides the capacity of the CPU (indicated as number of instructions Per Second or Hz – NB. the (time unit) step size of time is 1 nanosecond). The constructor for the BUS class takes three parameters. The first one indicates the kind of bus, the second one the capacity of the bus (its band width in bytes per second) and finally the third parameter gives a set of CPU instances connected together by the given BUS instance.

The currently supported primary scheduling policies for the CPU are:

<FP>: Fixed Priority

<FCFS>: First Come First Served

The currently supported primary scheduling policy for the BUS is:

<FCFS>: First Come First Served

The CPU class have member operations called deploy and setPriority. The deploy operation takes one significant parameter which must be an object that is declared as a static instance variable inside the system[1]. The semantics of the deploy operation is that execution of all functionality inside this object will take place on the CPU that it has been deployed to. The setPriority operation takes two parameters where the first must be the name of a public operation that has been deployed to the CPU and the second parameter is a natural number. The semantics of the setPriority operation is that the given operation is assigned the given priority (the second parameter). This will be used when fixed priority

---

[1]It is also allowed to take a string as a second parameter for future extensions but that is ignored at the moment.

scheduling is used on the given CPU. Per default operations that are not explicitly assigned a priority using the `setPriority` operation are assigned a default priority of 1.

The system "class" is limited in the way that it can only contain:

Instance variables: The only instances that can be declared in the system "class" is of the special classes `CPU` and `BUS` as well as static instances of the different system components that one wish to allocate to different CPU's.

Constructor: The actual deployment of instances to CPU's and setting of priorities for the different operations is set inside the constructor which is the only operation that can be placed in the system "class". The only kind of statements that can be used inside this constructor is a block statement with a sequence of invocations of the special `deploy` and `setPriority` operations.

In addition there are limitations with respect to the use of static declarations for instances that are deployed to different CPU's. Basically the user should ensure that only one instance is deployed to a CPU if the class the instance comes from contains any static operations or functions. In case a static instance variable is used it is accessed directly (without any communication over the busses, so this in essence not proper from a distribution standpoint. Thus, all instance variables of instances to be deployed should only be accessed through the use of operations.

Example: The system class could for example be defined as:

```
system Simple

instance variables
  static public a : A := new A();
  static public b : B := new B();
  -- define the first CPU with fixed priority scheduling
  -- and 22E6 Hz
  CPU1 : CPU := new CPU (<FP>, 22E6);

  static public c : C := new C();
  -- define the second CPU with fixed priority scheduling
  -- and 11E6 Hz
  CPU2 : CPU := new CPU (<FP>, 11E6);

  -- create a communication bus that links the three
  -- CPU's together
  BUS1 : BUS := new BUS (<CSMACD>, 72E3, {CPU1, CPU2})

operations
  public Simple: () ==> Simple
```

```
  Simple () ==
    ( -- deploy a on CPU1
      CPU1.deploy(a);
      -- deploy b on CPU1
      CPU1.deploy(b);
      -- deploy c on CPU2
      CPU2.deploy(c,"CT");
      -- "CT" is a label here which is ignored
    );

end Simple
```

where A, B and C all are defined as classes.

13.3.1 Classes

Compared to the standard VDM-SL language, VDM++ and VDM-RT have been extended with classes. In this section, the use of classes to create and structure a top-level specification will be described. With the object oriented facilities offered by VDM++ and VDM-RT it is possible to:

- Define classes and create objects.

- Define associations and create links between objects.

- Make generalisation and specialisation through inheritance.

- Describe the functional behaviour of the objects using functions and operations.

- Describe the dynamic behaviour of the system through threads and synchronisation constraints.

Before the actual facilities are described, the general layout of a class is described.

Syntax: class = '**class**', identifier, [inheritance clause],
             [class body],
             '**end**', identifier ;

        inheritance clause = '**is subclass of**', identifier, ',', { identifier } ;

        class body = definition block, { definition block } ;

```
definition block  =  type definitions
                   |  value definitions
                   |  function definitions
                   |  operation definitions
                   |  instance variable definitions
                   |  synchronization definitions
                   |  thread definitions
                   |  traces definitions  ;
```

Semantics: Each class description has the following parts:

- A class header with the class name and an optional *inheritance clause*.

- An optional *class body*.

- A class tail.

The class name as given in the class header is the defining occurrence of the name of the class. A class name is globally visible, i.e. visible in all other classes in the specification.

The class name in the class header must be the same as the class name in the class tail. Furthermore, defining class names must be unique throughout the specification.

The (optional) class body may consist of:

- A set of *value definitions* (constants).

- A set of *type definitions*.

- A set of *function definitions*.

- A set of *instance variable definitions* describing the internal state of an object instantiated from the class. State invariant expressions are encouraged but are not mandatory.

- A set of *operation definitions* that can act on the internal state.

- A set of the *synchronization definitions*, specified either in terms of permission predicates or using mutex constraints.

- A set of *thread definitions* that describe the thread of control for active objects.

- A set of *traces* that are used to indicate the sequences of operation calls for which test cases are desired to be produced automatically.

In general, all constructs defined within a class must have a unique name, e.g. it is not allowed to define an operation and a type with the same name. However, it is possible to *overload* function and operation names (i.e. it is possible to have two or more functions with the same name and two or more operations with the same name) subject to the restriction that the types of their input parameters should not overlap. That is, it should be possible using static type checking alone to determine uniquely and unambiguously which function/operation definition corresponds to each function/operation call. Note that this applies not only to

functions and operations defined in the local interface of a class but also to those inherited from superclasses. Thus, for example, in a design involving multiple inheritance a class C may inherit a function from a class A and a function with the same name from a class B and all calls involving this function name must be resolvable in class C.

13.3.2   Inheritance

The concept of inheritance is essential to object orientation. When one defines a class as a subclass of an already existing class, the definition of the subclass introduces an extended class, which is composed of the definitions of the superclass together with the definitions of the newly defined subclass.

Through inheritance, a subclass inherits from the superclass:

- Its instance variables. This also includes all invariants and their restrictions on the allowed modifications of the state.

- Its operation and function definitions.

- Its value and type definitions.

- Its synchronization definitions as described in section 14.2.

- Its thread definitions as described in chapter 15.

A name conflict occurs when two constructs of the same kind and with the same name are inherited from different superclasses. Name conflicts must be explicitly resolved through *name qualification*, i.e. prefixing the construct with the name of the superclass and a '-sign (back-quote).

Example: In the first example, we see that inheritance can be exploited to allow a class definition to be used as an abstract interface which subclasses must implement:

```
class Sort

  instance variables
    protected data : seq of int

  operations

    initial_data : seq of int ==> ()
    initial_data (l) ==
      data := l;

    sort_ascending : () ==> ()
    sort_ascending () == is subclass responsibility;
```

```
end Sort

class SelectionSort is subclass of Sort

  functions

    min_index : seq1 of nat -> nat
    min_index(l) ==
      if len l = 1
      then 1
      else let mi = min_index(tl l)
           in
             if l(mi+1) < hd l
             then mi+1
             else 1

  operations

    sort_ascending : () ==> ()
    sort_ascending () == selectSort(1);

    selectSort : nat ==> ()
    selectSort (i) ==
      if i < len data
      then (dcl temp: nat;
            dcl mi: nat := min_index(data(i,...,len data)) +
                           i - 1;

            temp := data(mi);
            data(mi) := data(i);
            data(i) := temp;
            selectSort(i+1)
            )

end SelectionSort
```

Here the class Sort defines an abstract interface to be implemented by different sorting algorithms. One implementation is provided by the SelectionSort class.

The next example clarifies how name space clashes are resolved.

```
    class A
```

```
    instance variables}
      i: int := 1;
      j: int := 2;
    end A

    class B is subclass of A
    end B

    class C is subclass of A
    instance variables
      i: int := 3;
    end C

    class D is subclass of B,C
    operations
      GetValues: () ==> seq of int
      GetValues() ==
        return [
            A`i, -- equal to 1
            B`i, -- equal to 1 (A`i)
            C`i, -- equal to 3
            j   -- equal to 2 (A`j)
          ]
      end D
```

In the example objects of class D have 3 instance variables: A`i, A`j and C`j. Note that objects of class D will have only one copy of the instance variables defined in class A even though this class is a common super class of both class B and C. Thus, in class D the names B`j, C`j, D`j and j are all referring to the same variable, A`j. It should also be noticed that the variable name i is ambiguous in class D as it refers to different variables in class B and class C.

13.3.3   Interface and Availability of Class Members

In VDM++ and VDM-RT definitions inside a class are distinguished between:

Class attribute: an attribute of a class for which there exists exactly one incarnation no matter how many instances (possibly zero) of the class may eventually be created. Class attributes in VDM++ and VDM-RT correspond to **static** class members in languages like C++ and Java. Class (static) attributes can be referenced by prefixing the name of the attribute with the name of the class followed by a '-sign (back-quote), so that, for example, ClassName`val refers to the value val defined in class ClassName.

Instance attribute: an attribute for which there exists one incarnation for each instance of the class. Thus, an instance attribute is only available in an object and each object has its own copy of its instance attributes. Instance (non-static) attributes can be referenced by prefixing the name of the attribute with the name of the object followed by a dot, so that, for example, `object.op()` invokes the operation `op` in the object denoted by `object` (provided that `op` is visible to `object`).

Functions, operations, instance variables and constants[2] in a class may be either class attributes or instance attributes. This is indicated by the keyword **static**: if the declaration is preceded by the keyword **static** then it represents a class attribute, otherwise it denotes an instance attribute.

Other class components are by default always either class attributes or instance attributes as follows:

- Type definitions are always class attributes.

- Thread definitions are always instance attributes. Thus, each active object has its own thread(s).

- Synchronization definitions are always instance attributes. Thus, each object has its own "history" when it has been created.

In addition, the interface or accessibility of a class member may be explicitly defined using an access specifier: one of **public**, **private** or **protected**. The meaning of these specifiers is:

public: Any class may use such members

protected: Only subclasses of the current class may use such members

private: No other class may use such members - they may only be used in the class in which they are specified.

The default access to any class member is **private**. That is, if no access specifier is given for a member it is private.

This is summarized in table 13.1. A few provisos apply here:

- Granting access to instance variables (i.e. through a **public** or **protected** access specifier) gives both read and write access to these instance variables.

- Public instance variables may be read (but not written) using the dot (for object instance variables) or back-quote (for class instance variables) notation e.g. a public instance variable `v` of an object `o` may be accessed as `o.v`.

- Access specifiers may only be used with type, value, function, operation and instance variable definitions; they cannot be used with thread or synchronization definitions.

---

[2]In practice, constants will generally be static – a non-static constant would represent a constant whose value may vary from one instance of the class to another which would be more naturally represented by an instance variable.

	public	protected	private
Within the class	√	√	√
In a subclass	√	√	×
In an arbitrary external class	√	×	×

Table 13.1: Summary of Access Specifier Semantics

- It is not possible to convert a class attribute into an instance attribute, or vice-versa.

- For inherited classes, the interface to the subclass is the same as the interface to its superclasses extended with the new definitions within the subclass.

- Access to an inherited member cannot be made more restrictive e.g. a public instance variable in a superclass cannot be redeclared as a private instance variable in a subclass.

Example In the example below use of the different access specifiers is demonstrated, as well as the default access to class members. Explanation is given in the comments within the definitions.

```
class A

types
public Atype = <A> | <B> | <C>

values
public Avalue = 10;

functions
public compare : nat -> Atype
compare(x) ==
  if x < Avalue
  then <A>
  elseif x = Avalue
  then <B>
  else <C>

instance variables
public v1: nat;
private v2: bool := false;
protected v3: real := 3.14;

operations
```

```
protected AInit : nat * bool * real ==> ()
AInit(n,b,r) ==
  (v1 := n;
   v2 := b;
   v3 := r)
end A


class B is subclass of A


instance variables
v4 : Atype --inherited from A


operations


BInit: () ==> ()
BInit() ==
  (AInit(1,true,2.718); --OK: can access protected members
                        --in superclass
   v4 := compare(v1);   --OK since v1 is public
   v3 := 3.5;           --OK since v3 protected and this
                        --is a subclass of A
   v2 := false          --illegal since v2 is private to A
  )


end B

class C

instance variables
a: A := new A();
b: B := new B();

operations

CInit: () ==> A`Atype--types are class attributes
CInit() ==
  (a.AInit(3,false,1.1);
                    --illegal since AInit is protected
   b.BInit();       --illegal since BInit is (by default)
                    --private
   let - = a.compare(b.v3) in skip;
                    --illegal since C is not subclass
```

```
                            --of A so b.v3 is not available
    return b.compare(B`Avalue)
                            --OK since compare is a public instance
                            --attribute and Avalue is public class
                            --attribute in B
  )

end C
```

Chapter 14

Synchronization Constraints (VDM++ and VDM-RT)

In general a complete system contains objects of a passive nature (which only react when their operations are invoked) and active objects which 'breath life' into the system. These active objects behave like virtual machines with their own processing thread of control and after start up they do not need interaction with other objects to continue their activities. In another terminology a system could be described as consisting of a number of active clients requesting services of passive or active servers. In such a parallel environment the server objects need synchronization control to be able to guarantee internal consistency, to be able to maintain their state invariants. Therefore, in a parallel world, a passive object needs to behave like a Hoare monitor with its operations as entries.

If a sequential system is specified (in which only one thread of control is active at a time) only a special case of the general properties is used and no extra syntax is needed. However, in the course of development from specification to implementation more differences are likely to appear.

The following default synchronization rules for each object apply in VDM++ and VDM-RT:

- operations are to be viewed as though they are atomic, from the point of the caller;

- operations which have no corresponding permission predicate are subject to no restrictions at all;

- synchronization constraints apply equally to calls within an object (i.e. one operation within an object calls another operation within that object) and outside an object (i.e. an operation from one object calls an operation in another object);

- operation invocations have the semantics of a rendez-vous (as in Ada, see [Ada LRM]) in case two active objects are involved. Thus if an object O_1 calls an operation o in object O_2, if O_2 is currently unable to start operation o then O_1 blocks until the operation may be executed. Thus invocation occurs when both the calling object and the called object are ready. (Note here a slight difference from the semantics of Ada: in Ada both parties to the rendez-vous are active objects; in VDM++ and VDM-RT only the calling party is active).

The synchronization definition blocks of the class description provide the user with ways to override the defaults described above.

Syntax: synchronization definitions = '**sync**', [synchronization] ;

synchronization = permission predicates ;

Semantics: Synchronization is specified in VDM++ and VDM-RT using permission predicates.

14.1   Permission Predicates

The following gives the syntax used to state rules for accepting the execution of concurrently callable operations. Some notes are given explaining these features.

Syntax: permission predicates = permission predicate, { '**;**',
                                          permission predicate } ;

permission predicate = '**per**', name, '**=>**', expression
                                 | mutex predicate ;

mutex predicate = '**mutex**', ' (', '**all**', ') '
                             | '**mutex**', ' (', name list ') ' ;

Semantics: Permission to accept execution of a requested operation depends on a guard condition in a (deontic) permission predicate of the form:

> **per** *operation name => guard condition*

The use of implication to express the permission means that truth of the guard condition (expression) is a necessary but not sufficient condition for the invocation. The permission predicate is to be read as stating that if the guard condition is false then there is non-permission. Expressing the permission in this way allows further similar constraints to be added without risk of contradiction through inheritance for the subclasses. There is a default for all operations:

> **per** *operation name =>* **true**

but when a permission predicate for an operation is specified this default is overridden.

Guard conditions can be conceptually divided into:

- a *history guard* defining the dependence on events in the past;
- an *object state guard*, which depends on the instance variables of the object, and

154

- a *queue condition guard*, which depends on the states of the queues formed by operation invocations (messages) awaiting service by the object.

These guards can be freely mixed. **Note** that there is no *syntactic* distinction between these guards - they are all expressions. However they may be distinguished at the semantic level.

A mutex predicate allows the user to specify either that all operations of the class are to be executed mutually exclusive, or that a list of operations are to be executed mutually exclusive to each other. Operations that appear in one mutex predicate are allowed to appear in other mutex predicates as well, and may also be used in the usual permission predicates. Each mutex predicate will implicitly be translated to permission predicates using history guards for each operation mentioned in the name list. For instance,

```
sync
  mutex(opA, opB);
  mutex(opB, opC, opD);
  per opD => someVariable > 42;
```

would be translated to the following permission predicates:

```
sync
  per opA => #active(opB) = 0;
  per opB => #active(opA) = 0 and
             #active(opC) + #active(opD) = 0;
  per opC => #active(opB) + #active(opD) = 0;
  per opD => #active(opB) + #active(opC) = 0 and
             someVariable > 42;
```

Note that it is only permitted to have one "stand-alone" permission predicate for each operation. It is also important to note that if permission predicates are made over operations that are overloaded (see Section 13.3.1) then it will incorporate all of their history counters as the same operation. The **#active** operator is explained below.

A **mutex(all)** constraint specifies that all of the operations specified in that class *and any superclasses* are to be executed mutually exclusively.

14.1.1 History guards

Semantics: A history guard is a guard which depends on the sequence of earlier invocations of the operations of the object expressed in terms of history expressions (see section 6.23). History expressions denotes the number of activations and completions of the operations, given as functions

#act and **#fin**, respectively.

#act: operation name $\rightarrow \mathbb{N}$
#fin: operation name $\rightarrow \mathbb{N}$

Furthermore, a derived function **#active** is available such that **#active**(A) = **#act**(A) - **#fin**(A), giving the number of currently active instances of A. Another history function – **#req** – is defined in section 14.1.3.

Examples: Consider a Web server that is capable of supporting 10 simultaneous connections and can buffer a further 100 requests. In this case we have one instance variable, representing the mapping from URLs to local filenames:

```
instance variables
  site_map : map URL to Filename := {|->}
```

The following operations are defined in this class (definitions omitted for brevity):

ExecuteCGI:	URL ==> File	Execute a CGI script on the server
RetrieveURL:	URL ==> File	Transmit a page of html
UploadFile:	File * URL ==> ()	Upload a file onto the server
ServerBusy:	() ==> File	Transmit a "server busy" page
DeleteURL:	URL ==> ()	Remove an obsolete file

Since the server can support only 10 simultaneous connects, we can only permit an execute or retrieve operation to be activated if the number already active is less than 10:

```
    per RetrieveURL => #active(RetrieveURL) +
                       #active(ExecuteCGI) < 10;
    per ExecuteCGI  => #active(RetrieveURL) +
                       #active(ExecuteCGI) < 10;
```

14.1.2 The object state guard

Semantics: The object state guard is a boolean expression which depends on the values of one (or more) instance variable(s) of the object itself. Object state guards differ from operation pre-conditions in that a call to an an operation whose permission predicate is false results in the caller blocking until the predicate is satisfied, whereas a call to an operation whose pre-condition is false means the operation's behaviour is unspecified.

Examples: Using the web server example again, we can only allow file removal if some files already exist:

```
        per DeleteURL   => dom site_map <> {}
```

Constraints for safe execution of the operations `Push` and `Pop` in a stack object can be expressed using an object state guard as:

```
per Push => length < maxsize;
per Pop => length > 0
```

where `maxsize` and `length` are instance variables of the stack object.

It is often possible to express such constraints as a consequence of the history, for example the empty state of the stack:

```
    length = 0 <=> #fin(Push) = #fin(Pop)
```

However, the size is a property which is better regarded as a property of the particular stack instance, and in such cases it is more elegant to use available instance variables which store the effects of history.

14.1.3 Queue condition guards

Semantics: A queue condition guard acts on requests waiting in the queues for the execution of the operations. This requires use of a third history function **#req** such that **#req**(A) counts the number of messages which have been received by the object requesting execution of operation A. Again it is useful to introduce the function **#waiting** such that: **#waiting**(A) = **#req**(A) - **#act**(A), which counts the number of items in the queue.

Examples: Once again, with the web server we can only activate the `ServerBusy` operation if 100 or more connections are waiting:

```
per ServerBusy  => #waiting(RetrieveURL) +
                   #waiting(ExecuteCGI) >= 100;
```

The most important use of such expressions containing queue state functions is for expressing priority between operations. The protocol specified by:

```
per B => #waiting(A) = 0
```

gives priority to waiting requests for activation of `A`. There are, however, many other situations when operation dispatch depends on the state of waiting requests. Full description

of the queuing requirements to allow specification of operation selection based on request arrival times or to describe 'shortest job next' behaviour will be a future development.

Note that #**req**(A) have value 1 at the time of evaluation of the permission predicate for the first invocation of operation A. That is,

```
per A => #req(A) = 0
```

would always block.

14.1.4 Evaluation of Guards

Using the previous example, consider the following situation: the web server is handling 10 RetrieveURL requests already. While it is dealing with these requests, two further RetrieveURL requests (from objects O_1 and O_2) and one ExecuteCGI request (from object O_3) are received. The permission predicates for these two operations are false since the number of active RetrieveURL operations is already 10. Thus these objects block.

Then, one of the active RetrieveURL operations reaches completion. The permission predicate so far blocking O_1, O_2 and O_3 will become "true" simultaneously. This raises the question: which object is allowed to proceed? Or even all of them?

Guard expressions are only reevaluated when an event occurs (in this case the completion of a RetrieveURL operation). In addition to that the test of a permission predicate by an object and its (potential) activation is an atomic operation. This means, that when the first object evaluates its guard expression, it will find it to be true and activate the corresponding operation (RetrieveURL or ExecuteCGI in this case). The other objects evaluating their guard expressions afterwards will find that **#active**(RetrieveURL) + **#active**(ExecuteCGI) = 10 and thus remain blocked. *Which object is allowed to evaluate the guard expression first is undefined.*

It is important to understand that the guard expression need only evaluate to **true** at the time of the activation. In the example as soon as O_1, O_2 or O_3's request is activated its guard expression becomes false again.

14.2 Inheritance of Synchronization Constraints

Synchronization constraints specified in a superclass are inherited by its subclass(es). The manner in which this occurs depends on the kind of synchronization.

14.2.1 Mutex constraints

Mutex constraints from base classes and derived classes are simply added. If the base class and derived class have the mutex definitions M_A and M_B, respectively, then the derived class simply

has both mutex constraints M_A, and M_B. The binding of operation names to actual operations is always performed in the class where the constraint is defined. Therefore a **mutex(all)** constraint defined in a superclass and inherited by a subclass only makes the operations from the base class mutually exclusive and does not affect operations of the derived class.

Inheritance of mutex constraints is completely analogous to the inheritance scheme for permission predicates. Internally mutex constraints are always expanded into appropriate permission predicates which are added to the existing permission predicates as a conjunction. This inheritance scheme ensures that the result (the final permission predicate) is the same, regardless of whether the mutex definitions are expanded in the base class and inherited as permission predicates or are inherited as mutex definitions and only expanded in the derived class.

The intention for inheriting synchronization constraints in the way presented is to ensure, that any derived class at least satisfies the constraints of the base class. In addition to that it must be possible to strengthen the synchronization constraints. This can be necessary if the derived class adds new operations as in the following example:

```
class A
operations

  writer: () ==> ()
  writer() == is not yet specified

  reader: () ==> ()
  reader() == is not yet specified

  sync
  per reader => #active(writer) = 0;
  per writer => #active(reader, writer) = 0;
end A

class B is subclass of A
operations

  newWriter: () ==> ()
  newWriter() == is not yet specified

  sync
  per reader => #active(newWriter) = 0;
  per writer => #active(newWriter) = 0;
  per newWriter => #active(reader, writer, newWriter) = 0;

end B
```

Class A implements reader and writer operations with the permission predicates specifying the multiple readers-single writer protocol. The derived class B adds newWriter. In order to ensure deterministic behaviour B also has to add permission predicates for the inherited operations.

The actual permission predicates in the derived class are therefore:

```
per reader => #active(writer)=0 and #active(newWriter)=0;
per writer => #active(reader, writer)=0 and #active(newWriter)=0;
per newWriter => #active(reader, writer, newWriter)=0;
```

A special situation arises when a subclass overrides an operation from the base class. The overriding operation is treated as a new operation. It has no permission predicate (and in particular inherits none) unless one is defined in the subclass.

The semantics of inheriting mutex constraints for overridden operations is completely analogous: newly defined overriding operations are not restricted by mutex definitions for equally named operations in the base class. The **mutex(all)** shorthand makes all inherited and locally defined operations mutually exclusive. Overridden operations (defined in a base class) are not affected. In other words, all operations, that can be called with an unqualified name ("locally visible operations") will be mutex to each other.

Chapter 15

Threads (VDM++ and VDM-RT)

Objects instantiated from a class with a *thread* part are called *active* objects. The scope of the instance variables and operations of the current class is considered to extend to the thread specification. Note that from a tool perspective the thread for the expression a user would like to evaluate in relation to a VDM model is called a debug thread. This thread has a special role in the sense that when it is finished the entire execution is completed (and thus all other threads ready to be scheduled in or running will be thrown away and aborted). If a session where a series of expressions are being evaluated this is not true (in this case the other threads will be continued when the next expression is executed, see [Larsen&13] for more details about "sessions"). Thus if one would like to ensure a specific number of such other threads to be completed before stopping the execution one needs to block the debug thread using a synchonisation as explained in Chapter 14.

Syntax: thread definitions = '**thread**', [thread definition] ;

thread definition = periodic thread definition
          | procedural thread definition ;

periodic thread definition = periodic obligation
          | sporadic obligation ;

Subclasses inherit threads from superclasses. If a class inherits from several classes only one of these may declare its own thread (possibly through inheritance). Furthermore, explicitly declaring a thread in a subclass will override any inherited thread.

15.1   Periodic Thread Definitions (VDM-RT)

The periodic obligation can be regarded as the way of describing repetitive activities in a class. As the 'period' implies a explicit notion of time, this construct is only available in VDM-RT.

Syntax(VDM-RT): Time is explicit in VDM-RT, using a discrete clock with a 1 nsec resolution, where the periodic obligation looks like:
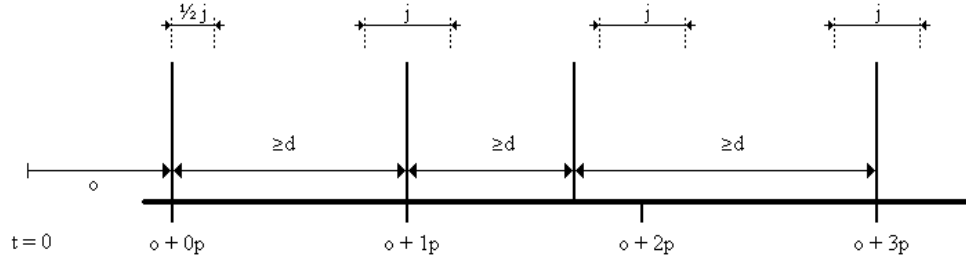
Figure 15.1: Period (p), jitter (j), delay (d) and offset (o)

$$\text{periodic obligation} = \text{'}\mathbf{periodic}\text{'},$$
$$\text{'('}, \text{expression}, \text{expression}, \text{expression}, \text{expression}, \text{')'},$$
$$\text{'('}, \text{name}, \text{')'} \text{ ;}$$

Semantics (VDM-RT): The type of the expressions should all yield a natural number (as we use a natural number valued wall-clock in the VDM-RT interpreter), otherwise a run-time error will occur. Note that the evaluation of the expressions also causes time to elapse whenever the '**start**' or '**start_list**' statement is executed. The expressions all denote a time value with a resolution of 1 nsec. For each periodic obligation, four different numbers are used. They are, in order of appearance (also illustrated in Figure 15.1):

1. **period:** This is a non-negative, non-zero value that describes the length of the time interval between two adjacent events in a strictly periodic event stream (where jitter = 0). Hence, a value of 1E9 denotes a period of 1 second.

2. **jitter:** This is a non-negative value that describes the amount of time variance that is allowed around a single event. We assume that the interval is balanced [-j, j]. Note that jitter is allowed to be bigger than the period to characterize so-called event bursts.

3. **delay:** This is a non-negative value smaller than the period which is used to denote the minimum inter arrival distance between two adjacent events.

4. **offset:** This is a non-negative value which is used to denote the absolute time value at which the first period of the event stream starts. Note that the first event occurs in the interval [offset, offset + jitter].

Given a defined time resolution ΔT, a thread with a periodic obligation invokes the mentioned operation at the beginning of each time interval with length *period*. This creates the periodic execution of the operation simulating the discrete equivalent of continuous relations which have to be maintained between instance variables, parameter values and possibly other external values obtained through operation invocations. It is not possible to dynamically change the length of the interval.

Periodic obligations are intended to describe e.g. analogue physical relations between values in formulas (e.g. transfer functions) and their discrete event simulation. It is a requirement

on the implementation to guarantee that the execution time of the operation is at least smaller than the used periodic time length. If other operations are present the user has to guarantee that the fairness criteria for the invocation of these other operations are maintained by reasoning about the time slices used internally and available for external invocations.

Note that a periodic thread is *neither* created *nor* started when an instance of the corresponding class is created. Instead, as with procedural threads, '**start**' (or '**start_list**') statements should be used with periodic threads.

Examples: Consider a timer class which periodically increments its clock in its own thread. It provides operations for starting, and stopping timing, and reading the current time.

```
class Timer
values
  PERIOD : nat = 1000
```

The Timer has two instance variables the current time and a flag indicating whether the Timer is active or not (the current time is only incremented if the Timer is active).

```
instance variables
  curTime : nat := 0;
  acti    : bool := false;
```

The Timer provides straightforward operations which need no further explanation.

```
operations
  public Start : () ==> ()
  Start() ==
    (acti := true;
     curTime := 0);

  public Stop : () ==> ()
  Stop() ==
    acti := false;

  public GetTime : () ==> nat
  GetTime() ==
    return curTime;

  IncTime: () ==> ()
  IncTime() ==
    if acti
```

```
    then curTime := curTime + 100;
```

The Timer's thread ensures that the current time is incremented. The period with which this is done is 1000 time units (nanoseconds). The allowed jitter is 10 time units and the minimal distance between two instances is 200 time units and finally no offset has been used.

```
thread
  periodic (PERIOD, PERIOD/100, PERIOD/5, 0) (IncTime)

end Timer
```

15.2   Sporadic Thread Definitions (VDM-RT)

The sporadic obligation can be regarded as the way of describing stochastic activities in a class.

Syntax: (VDM-RT) sporadic obligation = '**sporadic**',
'(', expression, expression, expression, ')',
'(', name, ')' ;

Semantics: (VDM-RT) The type of the expressions should all yield a natural number (as we use a natural number valued wall-clock in the VDM-RT interpreter), otherwise a run-time error will occur. Note that the evaluation of the expressions also causes time to elapse whenever the '**start**' or '**start_list**' statement is executed. The expressions all denote a time value with a resolution of 1 nsec. For each sporadic obligation, three different numbers are used. They are, in order of appearance:

1. **delay:** This value is used to denote the minimum inter arrival distance between two adjacent thread invocations.

2. **bound:** This value, greater than **delay**, is the maximum inter arrival distance between two adjacent thread invocations.

3. **offset:** This is a non-negative value which is used to denote the absolute time value at which the first period starts, randomly in the interval $[offset, offset + bound]$.

Given these definitions, and assuming the last thread was invoked at t_0 then the next invocation is randomly scheduled in the interval $[t_0 + delay, t_0 + bound]$.

Examples: Analogous to the example in '**periodic**' for the Timer class, the sporadic definition could be as follows:

```
class Timer
...

thread
  sporadic (100, 1000, 0) (IncTime)

end Timer
```

15.3   Procedural Thread Definitions (VDM++ and VDM-RT)

A procedural thread provides a mechanism to explicitly define the external behaviour of an active object through the use of *statements*, which are executed when the object is started (see section 12.14).

Syntax: procedural thread definition =   statement ;

Semantics: A procedural thread is scheduled for execution following the application of a start statement to the object owning the thread. The statements in the thread are then executed sequentially, and when execution of the statements is complete, the thread dies. Synchronization between multiple threads is achieved using permission predicates on shared objects.

Examples: The example below demonstrates procedural threads by using them to compute the factorial of a given integer concurrently.

```
class Factorial

instance variables
  result : nat := 5;
operations

public factorial : nat ==> nat
factorial(n) ==
  if n = 0
  then return 1
  else (dcl m : Multiplier;
        m := new Multiplier();
        m.calculate(1,n);
        start(m);
        result:= m.giveResult();
        return result
```

```
        )

end Factorial

class Multiplier

instance variables
   i : nat1;
   j : nat1;
   k : nat1;
   result : nat1

operations

public calculate : nat1 * nat1 ==> ()
calculate (first, last) ==
   (i := first; j := last);

doit : () ==> ()
doit() ==
( if i = j
   then result := i
   else (dcl p : Multiplier;
         dcl q : Multiplier;
         p := new Multiplier();
         q := new Multiplier();
         start(p);
         start(q);
         k := (i + j) div 2;
         -- division with rounding down
         p.calculate(i,k);
         q.calculate(k+1,j);
         result := p.giveResult() *  q.giveResult ()
        )
);

public giveResult : () ==> nat1
giveResult() ==
   return result;

sync
-- cyclic constraints allowing only the
```

```
-- sequence calculate; doit; giveResult

per doit => #fin (calculate) > #act(doit);
per giveResult => #fin (doit) > #act (giveResult);
per calculate => #fin (giveResult) = #act (calculate)

thread
  doit();

end Multiplier
```

Chapter 16

Trace Definitions

In order to automate the testing process VDM-10 contains a notation enabling the expression of the traces that one would like to have tested exhaustively. Such traces are used to express combinations of sequences of operations that wish to be tested in all possible combinations. In a sense this is similar to model checking limitations except that this is done with real and not symbolic values. However, errors in test cases are filtered away so other test cases with the same prefix will be skipped automatically.

Syntax: traces definitions = '**traces**', [named trace], { ';', named trace } ;

named trace = identifier, { '/', identifier }, ':', trace definition list ;

trace definition list = trace definition term, { ';', trace definition term } ;

trace definition term = trace definition
          | trace definition term, '|', trace definition ;

trace definition = trace binding definition
        | trace repeat definition ;

trace binding definition = trace let def binding
         | trace let best binding ;

trace let def binding = '**let**', local definition, { ',', local definition },
        '**in**', trace definition ;

trace let best binding = '**let**' multiple bind, ['**be**', '**st**', expression],
        '**in**', trace definition ;

trace repeat definition = trace core definition, [trace repeat pattern] ;

```
trace repeat pattern  =  '*'
                      |  '+'
                      |  '?'
                      |  '{',  numeric literal, '}'
                      |  '{',  numeric literal, ',' numeric literal, '}'  ;


trace core definition  =  trace apply expression
                       |   trace concurrent expression
                       |   trace bracketed expression  ;


trace apply expression  =  call statement  ;


trace concurrent expression  =  '||', '(',  trace definition,
                                 ',',  trace definition,
                                 { ',',  trace definition }, ')'  ;


trace bracketed expression  =  '(',  trace definition list, ')'  ;
```

Semantics: Semantically the trace definitions provided in a class have no effect. These definitions are simply used to enhance testing of a VDM specification using principles from combinatorial testing (also called all-pairs testing). So each trace definition can be considered as a regular expression describing the test sequences in which different operations should be executed to test the VDM specification. Inside the trace definitions, bindings may appear and for each possible such binding a particular test case can be automatically derived. So one trace definition expand into a set of test cases. In this sense a test case is a sequence of operation calls executed after each other. Between each test case the VDM specification is initialised so they become entirely independent. From a static semantics perspective it is important to note that the expressions used inside trace definitions must be executed in the expansion process. This means that it cannot directly refer to instance variables, because these could be changed during the execution.

So here it makes sense to explain what kind of expansion the different kinds of trace definitions gives rise to.

The *trace definition lists* simply use a semicolon ("**;**") and this simply result in sequencing between the *trace definition terms* used inside it.

In the *trace definition term* it is possible to introduce alternatives using the bar ("**|**") operator. This result in test cases for all alternatives.

The *trace binding definition* exists in two forms where the *trace let def binding* simply enables the binding introduced to be used after the '**in**' in the same way as in let-expressions. Alternatively the *trace let best binding* can be used and this will expand to test cases with all the different possible bindings.

The *trace repeat definition* is used to introduce the possibility of having repetitions of the operation calls used in the trace. The different kinds of repeat patterns have the following meanings:

- '$*$' means 0 to n occurences (n is tool specific).
- '+' means 1 to n occurences (n is tool specific).
- '?' means 0 or 1 occurences.
- '{', n, '}' means n occurences.
- '{', n, ',' m '}' means between n and m occurences.

The *trace core definitions* have three possibilities. These are ordinary operation calls, trace concurrency expressions and bracketed trace definitions respectively. The *trace concurrency expressions* are similar to the *nondeterministic statements* in the sense that the trace definition lists inside it will be executed in all possible permulations of the elements. This is particular useful for concurrent VDM++ models where potential deadlocks can occur under some circumstances.

Examples: In an example like the one below test cases will be generated in all possible combination starting with a call of Reset followed by one to four Pushes of values onto the stack followed again by one to three Pops from the stack.

```
class Stack

instance variables
    stack : seq of int := [];

operations

    public Reset : () ==> ()
    Reset () ==
      stack := [];

    public Pop : () ==> int
    Pop() ==
      def res = hd stack in
        (stack := tl stack;
         return res)
    pre stack <> []
    post stack~ = [RESULT] ^ stack;

    public Push: int ==> ()
```

```
      Push(elem) ==
        stack := stack ^ [elem];

      public Top : () ==> int
      Top() ==
        return (hd stack);

end Stack
class UseStack

instance variables

  s : Stack := new Stack();

traces

  PushBeforePop : s.Reset();
                  (let x in set {1,2} in s.Push(x)){1,4};
                  s.Pop(){1,3}

end UseStack
```

References

[Ada LRM] *Reference Manual for the Ada Programming Language*. Technical Report, United States Government (Department of Defence), American National Standards Institute, 1983.

[Bicarregui&94] Juan Bicarregui and John Fitzgerald and Peter Lindsay and Richard Moore and Brian Ritchie. *Proof in VDM: A Practitioner's Guide. FACIT*, Springer-Verlag, 1994. 245 pages. ISBN 3-540-19813-X.

This book is a tutorial on the process of formal reasoning in VDM. It discusses how to go about about building proofs and provides the most complete set of proof rules for VDM-SL to date.

[Bjørner&78a] D. Bjørner and C.B. Jones, editors. *The Vienna Development Method: The Meta-Language*. Volume 61 of *Lecture Notes in Computer Science*, Springer-Verlag, 1978.

This was the first monograph on *Meta-IV*. See also entries: [Bjørner78b], [Bjørner78c], [Lucas78], [Jones78a], [Jones78b], [Henhapl&78]

[Bjørner78b] D. Bjørner. Programming in the Meta-Language: A Tutorial. *The Vienna Development Method: The Meta-Language*, 24–217, 1978.

An informal introduction to *Meta-IV*

[Bjørner78c] D. Bjørner. Software Abstraction Principles: Tutorial Examples of an Operating System Command Language Specification and a PL/I-like On-Condition Language Definition. *The Vienna Development Method: The Meta-Language*, 337–374, 1978.

Exemplifies so called **exit** semantics uses of *Meta-IV* to slightly non-trivial examples.

[Dawes91]   John Dawes. *The VDM-SL Reference Guide*. Pitman, 1991. 217 pages. ISBN 0-273-03151-1.

This is the best reference manual for the complete VDM-SL which is being standardised. It refers to a draft version of the standard which is quite close to the current version of the standard.

[Fitzgerald&05]   John Fitzgerald and Peter Gorm Larsen and Paul Mukherjee and Nico Plat and Marcel Verhoef. *Validated Designs for Object–oriented Systems*. Springer, New York, 2005.

[Fitzgerald&08a] J. S. Fitzgerald and P. G. Larsen and M. Verhoef. Vienna Development Method. *Wiley Encyclopedia of Computer Science and Engineering*, 2008. 11 pages. edited by Benjamin Wah, John Wiley & Sons, Inc..

[Fitzgerald&08b] John Fitzgerald and Peter Gorm Larsen and Shin Sahara. VDMTools: Advances in Support for Formal Modeling in VDM. *ACM Sigplan Notices*, 43(2):3–11, February 2008. 8 pages.

[Fitzgerald&98]   J.S. Fitzgerald and C.B. Jones. Proof in the Validation of a Formal Model of a Tracking System for a Nuclear Plant. In J.C. Bicarregui, editor, *Proof in VDM: Case Studies*, Springer-Verlag, 1998.

[Henhapl&78]   W. Henhapl, C.B. Jones. A Formal Definition of ALGOL 60 as described in the 1975 modified Report. In *The Vienna Development Method: The Meta-Language*, pages 305–336, Springer-Verlag, 1978.

One of several examples of ALGOL 60 descriptions.

[ISOVDM96]   P. G. Larsen and B. S. Hansen and H. Brunn and N. Plat and H. Toetenel and D. J. Andrews and J. Dawes and G. Parkin and others. Information technology – Programming languages, their environments and system software interfaces – Vienna Development Method – Specification Language – Part 1: Base language. December 1996.

[Jones78a]   C.B. Jones. The Meta-Language: A Reference Manual. In *The Vienna Development Method: The Meta-Language*, pages 218–277, Springer-Verlag, 1978.

[Jones78b]   C.B. Jones. The Vienna Development Method: Examples of Compiler Development. In Amirchachy and Neel, editors, *Le Point sur la Compilation*, INRIA Publ. Paris, 1979.

[Jones90]   Cliff B. Jones. *Systematic Software Development Using VDM*. Prentice-Hall International, Englewood Cliffs, New Jersey, second edition, 1990. 333 pages. ISBN 0-13-880733-7.

This book deals with the Vienna Development Method. The approach explains formal (functional) specifications and verified design with an emphasis on the study of proofs in the development process.

[Larsen&09]   Peter Gorm Larsen and John Fitzgerald and Sune Wolff. Methods for the Development of Distributed Real-Time Embedded Systems using VDM. *Intl. Journal of Software and Informatics*, 3(2-3), October 2009.

[Larsen&10]   Peter Gorm Larsen and Nick Battle and Miguel Ferreira and John Fitzgerald and Kenneth Lausdahl and Marcel Verhoef. The Overture Initiative – Integrating Tools for VDM. *SIGSOFT Softw. Eng. Notes*, 35(1):1–6, January 2010. 6 pages.

[Larsen&13]   Peter Gorm Larsen and Kenneth Lausdahl and Peter Jørgensen and Joey Coleman and Sune Wolff and Nick Battle. *Overture VDM-10 Tool Support: User Guide*. Technical Report TR-2010-02, The Overture Initiative, www.overturetool.org, April 2013. 130 pages.

[Lucas78]   P. Lucas. On the Formalization of Programming Languages: Early History and Main Approaches. In *The Systematic Development of Compiling Algorithm*, INRIA Publ. Paris, 1978.

An historic overview of the (VDL and other) background for VDM.

[Mukherjee&00]   Paul Mukherjee and Fabien Bousquet and Jérôme Delabre and Stephen Paynter and Peter Gorm Larsen. Exploring Timing Properties Using VDM++ on an Industrial Application. In J.C. Bicarregui and J.S. Fitzgerald, editors, *Proceedings of the Second VDM Workshop*, September 2000. Available at www.vdmportal.org.

[Paulson91]   Lawrence C. Paulson. *ML for the Working Programmer*. Cambridge University Press, 1991.

[Verhoef&06]   Marcel Verhoef and Peter Gorm Larsen and Jozef Hooman. Modeling and Validating Distributed Embedded Real-Time Systems with VDM++. In Jayadev Misra and Tobias Nipkow and Emil Sekerinski, editors, *FM 2006: Formal Methods*, pages 147–162, Springer-Verlag, 2006.

[Verhoef&07]   Marcel Verhoef and Peter Gorm Larsen. Interpreting Distributed System Architectures Using VDM++ – A Case Study. In Brian Sauser and Gerrit Muller, editors, *5th Annual Conference on Systems Engineering Research*, March 2007. Available at http://www.stevens.edu/engineering/cser/.

[Verhoef08]   Marcel Verhoef. *Modeling and Validating Distributed Embedded Real-Time Control Systems*. PhD thesis, Radboud University Nijmegen, 2009. The right one is [Verhoef09].

                ISBN 978-90-9023705-3

[Verhoef09]   Marcel Verhoef. *Modeling and Validating Distributed Embedded Real-Time Control Systems*. PhD thesis, Radboud University Nijmegen, 2009.

                ISBN 978-90-9023705-3

Appendix A

The Syntax of the VDM Languages

This appendix specifies the complete syntax for the VDM languages.

A.1 VDM-SL Document

document = any module, { any module }
             | definition block, { definition block } ;

any module = module ;

A.1.1 Modules

This entire subsection is not present in the current version of the VDM-SL standard.

module = '**module**', identifier, interface,
           [module body], '**end**', identifier ;

interface = [import definition list],
             export definition ;

import definition list = '**imports**', import definition,
                      { ',', import definition } ;

import definition = '**from**', identifier, import module signature ;

import module signature = '**all**'
                         | import signature, { import signature } ;

import signature = import types signature
                    | import values signature
                    | import functions signature
                    | import operations signature ;

import types signature = '**types**', type import,
{ ';', type import }, [';'] ;

type import = name, ['**renamed**', name]
| type definition, ['**renamed**', name] ;

import values signature = '**values**', value import,
{ ';', value import }, [';'] ;

value import = name, [':', type], ['**renamed**', name] ;

import functions signature = '**functions**', function import,
{ ';', function import }, [';'] ;

function import = name, [[type variable list], ':', function type],
['**renamed**', name] ;

import operations signature = '**operations**', operation import,
{ ';', operation import }, [';'] ;

operation import = name, [':', operation type], ['**renamed**', name] ;

export definition = '**exports**', export module signature ;

export module signature = '**all**'
| export signature,
{ export signature } ;

export signature = export types signature
| values signature
| export functions signature
| operations signature ;

export types signature = '**types**', type export,
{ ';', type export }, [';'] ;

type export = ['**struct**'], name ;

values signature = '**values**', value signature,
{ ';', value signature }, [';'] ;

value signature = name list, ':', type ;

export functions signature = '**functions**' function export,
{ ';', function export } ;

178

function export = name list, [type variable list], ':',
                         function type ;

functions signature = '**functions**' function signature,
                         { ';', function signature }, [';'] ;

function signature = name list, ':', function type ;

operations signature = '**operations**' operation signature,
                         { ';', operation signature }, [';'] ;

operation signature = name list, ':', operation type ;

A.2 VDM++ and VDM-RT Document

document = class | system , { class | system } ;

A.3 System (VDM-RT)

system = '**system**', identifier,
            [class body],
            '**end**', identifier ;

A.3.1 Classes

class = '**class**', identifier, [inheritance clause],
          [class body],
          '**end**', identifier ;

inheritance clause = '**is subclass of**', identifier, ',', { identifier } ;

A.4 Definitions

class body = definition block, { definition block } ;

module body = '**definitions**', definition block, { definition block } ;

definition block = type definitions
                       | state definition
                       | value definitions
                       | function definitions

| operation definitions
| instance variable definitions
| synchronization definitions
| thread definitions
| traces definitions ;

A.4.1   Type Definitions

type definitions = '**types**', [access type definition] ,
                        { ';', access type definition }, [';'] ;

access type definition = ([access], ['**static**']) | (['**static**'], [access]),
                            type definition ;

The access part is only possible in VDM++ and VDM-RT.

access = '**public**'
       | '**private**'
       | '**protected**' ;

type definition = identifier, '=', type, [invariant]
                 | identifier, '::', field list, [invariant] ;

type = bracketed type
      | basic type
      | quote type
      | composite type
      | union type
      | product type
      | optional type
      | set type
      | seq type
      | map type
      | partial function type
      | type name
      | type variable ;

bracketed type = '(', type, ')' ;

basic type = '**bool**' | '**nat**' | '**nat1**' | '**int**' | '**rat**'
            | '**real**' | '**char**' | '**token**' ;

quote type = quote literal ;

composite type = '**compose**', identifier, '**of**', field list, '**end**' ;

field list = { field } ;

field = [identifier, ':'], type
    | [identifier, ':−'], type ;

union type = type, '|', type, { '|', type } ;

product type = type, '∗', type, { '∗', type } ;

optional type = '[', type, ']' ;

set type = '**set of**', type ;

seq type = seq0 type
    | seq1 type ;

seq0 type = '**seq of**', type ;

seq1 type = '**seq1 of**', type ;

map type = general map type
    | injective map type ;

general map type = '**map**', type, '**to**', type ;

injective map type = '**inmap**', type, '**to**', type ;

function type = partial function type
    | total function type ;

partial function type = discretionary type, '−>', type ;

total function type = discretionary type, '+>', type ;

discretionary type = type
    | '(', ')' ;

type name = name ;

type variable = type variable identifier ;

invariant = '**inv**', invariant initial function ;

invariant initial function = pattern, '==', expression ;

A.4.2   The VDM-SL State Definition

state definition = '**state**', identifier, '**of**', field list,
[invariant], [initialisation], '**end**', ['**;**'] ;

invariant = '**inv**', invariant initial function ;

initialisation = '**init**', invariant initial function ;

invariant initial function = pattern, '==', expression ;

A.4.3   Value Definitions

value definitions = '**values**', [access value definition],
{ '**;**', access value definition }, ['**;**'] ;

access value definition = [access], value definition ;

value definition = pattern, ['**:**', type], '=', expression ;

A.4.4   Function Definitions

function definitions = '**functions**', [access function definition],
{ '**;**', access function definition }, ['**;**'] ;

access function definition = [access], function definition ;

function definition = explicit function definition
| implicit function definition
| extended explicit function definition ;

explicit function definition = identifier, [type variable list], '**:**',
function type,
identifier, parameters list,
'==', function body,
['**pre**', expression] ,
['**post**', expression],
['**measure**', name] ;

implicit function definition = identifier, [type variable list],
parameter types,
identifier type pair list,
['**pre**', expression],
'**post**', expression ;

182

In VDM-SL extended explicit function definition looks like:

extended explicit function definition = identifier, [type variable list],
                                 parameter types,
                                 identifier type pair list,
                                 '==', function body,
                                 ['**pre**', expression],
                                 ['**post**', expression] ;

In VDM++ and VDM-RT extended explicit function definition looks like:

extended explicit function definition = identifier, [type variable list],
                                 parameter types,
                                 identifier type pair list,
                                 '==', function body,
                                 ['**pre**', expression],
                                 ['**post**', expression] ;

type variable list = '[', type variable identifier,
                    { ',', type variable identifier }, ']' ;

identifier type pair = identifier, ':', type ;

parameter types = '(', [pattern type pair list], ')' ;

identifier type pair list = identifier, ':', type,
                            { ',', identifier, ':', type } ;

pattern type pair list = pattern list, ':', type,
                      { ',', pattern list,':', type } ;

parameters list = parameters, { parameters } ;

parameters = '(', [pattern list], ')' ;

function body = expression
              | '**is subclass responsibility**'
              | '**is not yet specified**' ;

A.4.5   Operation Definitions

operation definitions = '**operations**', [access operation definition],
                          { ';', access operation definition }, [';'] ;

access operation definition = (['**async**'] [access], ['**static**'])
                                | (['**async**'] ['**static**'], [access]),
                                    operation definition ;

operation definition = explicit operation definition
                         | implicit operation definition
                         | extended explicit operation definition ;

explicit operation definition = identifier, ':', operation type,
                                   identifier, parameters,
                                   '==', operation body,
                                   ['**pre**', expression],
                                   ['**post**', expression] ;

implicit operation definition = identifier, parameter types,
                                   [identifier type pair list],
                                   implicit operation body ;

implicit operation body = [externals],
                            ['**pre**', expression],
                            '**post**', expression,
                            [exceptions] ;

extended explicit operation definition = identifier, parameter types,
                                            [identifier type pair list],
                                            '==', operation body,
                                            [externals],
                                            ['**pre**', expression],
                                            ['**post**', expression],
                                            [exceptions] ;

operation type = discretionary type, '==>', discretionary type ;

operation body = statement
                   | '**is subclass responsibility**'
                   | '**is not yet specified**' ;

externals = '**ext**', var information, { var information } ;

var information = mode, name list, [':', type] ;

mode = '**rd**' | '**wr**' ;

exceptions = '**errs**', error list ;

error list = error, { error } ;

error = identifier, ':', expression, '->', expression ;

A.4.6   Instance Variable Definitions (VDM++ and VDM-RT)

instance variable definitions = '**instance**', '**variables**',
              [instance variable definition,
              { ';', instance variable definition }] ;

instance variable definition = access assignment definition
              | invariant definition ;

access assignment definition = ([access], ['**static**'])
              | (['**static**'], [access]),
              assignment definition ;

invariant definition = '**inv**', expression ;

A.4.7   Synchronization Definitions (VDM++ and VDM-RT)

synchronization definitions = '**sync**', [synchronization] ;

synchronization = permission predicates ;

permission predicates = permission predicate,
              { ';', permission predicate } ;

permission predicate = '**per**', name, '**=>**', expression
              | mutex predicate ;

mutex predicate = '**mutex**', '(', '**all**', ')'
           | '**mutex**', '(', name list ')' ;

A.4.8   Thread Definitions (VDM++ and VDM-RT)

thread definitions = '**thread**', [thread definition] ;

thread definition = periodic thread definition
    | procedural thread definition ;

periodic thread definition = periodic obligation
    | sporadic obligation ;

For VDM-RT where time is explicit, it looks like:

periodic obligation = '**periodic**',
    '(', expression, expression, expression, expression, ')',
    '(', name, ')' ;

sporadic obligation = '**sporadic**',
    '(', expression, expression, expression, ')',
    '(', name, ')' ;

For both VDM++ and VDM-RT, we can define:

procedural thread definition = statement ;

A.4.9   Trace Definitions

traces definitions = '**traces**', [named trace], { ';', named trace } ;

named trace = identifier, { '/', identifier }, ':', trace definition list ;

trace definition list = trace definition term, { ';', trace definition term } ;

trace definition term = trace definition
    | trace definition term, '|', trace definition ;

trace definition = trace binding definition
    | trace repeat definition ;

trace binding definition = trace let def binding
    | trace let best binding ;

;

trace let def binding = '**let**', local definition, { ',', local definition },
    '**in**', trace definition ;

trace let best binding = '**let**' multiple bind, ['**be**', '**st**', expression],
　　　　　　　　　　'**in**', trace definition ;

trace repeat definition = trace core definition, [trace repeat pattern] ;

trace repeat pattern = '$\star$'
　　　　　　　| '+'
　　　　　　　| '?'
　　　　　　　| '{', numeric literal, '}'
　　　　　　　| '{', numeric literal, ',' numeric literal, '}' ;

trace core definition = trace apply expression
　　　　　　　| trace concurrent expression
　　　　　　　| trace bracketed expression ;

trace apply expression = call statement ;

trace concurrent expression = '||', '(', trace definition,
　　　　　　　　　　',', trace definition,
　　　　　　　　　　{ ',', trace definition }, ')' ;

trace bracketed expression = '(', trace definition list, ')' ;

A.5   Expressions

expression list = expression, { ',', expression } ;

expression = bracketed expression
　　　　　　| let expression
　　　　　　| let be expression
　　　　　　| def expression
　　　　　　| if expression
　　　　　　| cases expression
　　　　　　| unary expression
　　　　　　| binary expression
　　　　　　| quantified expression
　　　　　　| iota expression
　　　　　　| set enumeration
　　　　　　| set comprehension
　　　　　　| set range expression
　　　　　　| sequence enumeration
　　　　　　| sequence comprehension
　　　　　　| subsequence

| map enumeration
| map comprehension
| tuple constructor
| record constructor
| record modifier
| apply
| field select
| tuple select
| function type instantiation
| lambda expression
| new expression
| self expression
| threadid expression
| general is expression
| undefined expression
| precondition expression
| isofbaseclass expression
| isofclass expression
| samebaseclass expression
| sameclass expression
| act expression
| fin expression
| active expression
| req expression
| waiting expression
| time expression
| name
| old name
| symbolic literal ;

A.5.1 Bracketed Expressions

bracketed expression = ‘ (’, expression, ‘) ’ ;

A.5.2 Local Binding Expressions

let expression = ‘**let**’, local definition, { ‘ , ’, local definition },
                   ‘**in**’, expression ;

let be expression = ‘**let**’, multiple bind, [‘**be**’, ‘**st**’, expression], ‘**in**’,
                       expression ;

def expression = ‘**def**’, pattern bind, ‘=’, expression,

$$\{ \; `;', \; \text{pattern bind}, \; `=', \; \text{expression} \}, [\; `;' \;],$$
$$`\textbf{in}', \; \text{expression} \; ;$$

A.5.3   Conditional Expressions

if expression = '**if**', expression, '**then**', expression,
                  { elseif expression },
                  '**else**', expression ;

elseif expression = '**elseif**', expression, '**then**', expression ;

cases expression = '**cases**', expression, '**:**',
                     cases expression alternatives,
                     [',', others expression], '**end**' ;

cases expression alternatives = cases expression alternative,
                                  { ',', cases expression alternative } ;

cases expression alternative = pattern list, '−>', expression ;

others expression = '**others**', '−>', expression ;

A.5.4   Unary Expressions

unary expression = prefix expression
                     | map inverse ;

prefix expression = unary operator, expression ;

unary operator = unary plus
                   | unary minus
                   | arithmetic abs
                   | floor
                   | not
                   | set cardinality
                   | finite power set
                   | distributed set union
                   | distributed set intersection
                   | sequence head
                   | sequence tail
                   | sequence length
                   | sequence elements
                   | sequence indices

```
                    |   sequence reverse
                    |   distributed sequence concatenation
                    |   map domain
                    |   map range
                    |   distributed map merge ;
```

unary plus = '+' ;

unary minus = '−' ;

arithmetic abs = '**abs**' ;

floor = '**floor**' ;

not = '**not**' ;

set cardinality = '**card**' ;

finite power set = '**power**' ;

distributed set union = '**dunion**' ;

distributed set intersection = '**dinter**' ;

sequence head = '**hd**' ;

sequence tail = '**tl**' ;

sequence length = '**len**' ;

sequence elements = '**elems**' ;

sequence indices = '**inds**' ;

sequence reverse = '**reverse**' ;

distributed sequence concatenation = '**conc**' ;

map domain = '**dom**' ;

map range = '**rng**' ;

distributed map merge = '**merge**' ;

map inverse = '**inverse**', expression ;

A.5.5   Binary Expressions

binary expression = expression, binary operator, expression ;

binary operator =   arithmetic plus
        | arithmetic minus
        | arithmetic multiplication
        | arithmetic divide
        | arithmetic integer division
        | arithmetic rem
        | arithmetic mod
        | less than
        | less than or equal
        | greater than
        | greater than or equal
        | equal
        | not equal
        | or
        | and
        | imply
        | logical equivalence
        | in set
        | not in set
        | subset
        | proper subset
        | set union
        | set difference
        | set intersection
        | sequence concatenate
        | map or sequence modify
        | map merge
        | map domain restrict to
        | map domain restrict by
        | map range restrict to
        | map range restrict by
        | composition
        | iterate ;

arithmetic plus = '+' ;

arithmetic minus = '−' ;

arithmetic multiplication = '*' ;

arithmetic divide = '/' ;

arithmetic integer division = '**div**' ;

arithmetic rem = '**rem**' ;

arithmetic mod = '**mod**' ;

less than = '<' ;

less than or equal = '<=' ;

greater than = '>' ;

greater than or equal = '>=' ;

equal = '=' ;

not equal = '<>' ;

or = '**or**' ;

and = '**and**' ;

imply = '=>' ;

logical equivalence = '<=>' ;

in set = '**in set**' ;

not in set = '**not in set**' ;

subset = '**subset**' ;

proper subset = '**psubset**' ;

set union = '**union**' ;

set difference = '\' ;

set intersection = '**inter**' ;

sequence concatenate = '^' ;

map or sequence modify = '++' ;

map merge = '**munion**' ;

map domain restrict to = '<:' ;

map domain restrict by = '<-:' ;

map range restrict to = ':>' ;

map range restrict by = ':->' ;

composition = '**comp**' ;

iterate = '**' ;

A.5.6   Quantified Expressions

quantified expression = all expression
| exists expression
| exists unique expression ;

all expression = '**forall**', bind list, '&', expression ;

exists expression = '**exists**', bind list, '&', expression ;

exists unique expression = '**exists1**', bind, '&', expression ;

A.5.7   The Iota Expression

iota expression = '**iota**', bind, '&', expression ;

A.5.8   Set Expressions

set enumeration = '{', [expression list], '}' ;

set comprehension = '{', expression, '|', bind list,
['&', expression], '}' ;

set range expression = '{', expression, ',', '...', ',',
expression, '}' ;

A.5.9   Sequence Expressions

sequence enumeration = '[', [expression list], ']' ;

sequence comprehension = '[', expression, '|', set bind,
['&', expression], ']' ;

subsequence = expression, '(', expression, ',', '...', ',',
expression, ')' ;

A.5.10   Map Expressions

map enumeration = '{', maplet, { ',', maplet }, '}'
| '{', '|->', '}' ;

maplet = expression, '|->', expression ;

map comprehension = '{', maplet, '|', bind list,
['&', expression], '}' ;

A.5.11   The Tuple Constructor Expression

tuple constructor = 'mk_', '(', expression, ',', expression list, ')' ;

A.5.12   Record Expressions

record constructor = '**mk_**',[1] name, '(', [expression list], ')' ;

record modifier = 'mu', '(', expression, ',',
                    record modification,
                    { ',', record modification }, ')' ;

record modification = identifier, '|->', expression ;

A.5.13   Apply Expressions

apply = expression, '(', [expression list], ')' ;

field select = expression, '.', identifier ;

tuple select = expression, '.#', numeral ;

function type instantiation = name, '[', type, { ',', type }, ']' ;

A.5.14   The Lambda Expression

lambda expression = '**lambda**', type bind list, '&', expression ;

A.5.15   The narrow Expression

narrow expression = '**narrow_**', '(', expression, ',', type, ')' ;

A.5.16   The New Expression (VDM++ and VDM-RT)

new expression = '**new**', name, '(', [expression list], ')' ;

A.5.17   The Self Expression (VDM++ and VDM-RT)

self expression = '**self**' ;

---

[1]**Note:** no delimiter is allowed

A.5.18   The Threadid Expression (VDM++ and VDM-RT)

threadid expression = '**threadid**' ;

A.5.19   The Is Expression

general is expression = is expression
                      | type judgement ;

is expression = '**is_**',[2] name, '(', expression, ')'
             | is basic type, '(', expression, ')' ;

type judgement = '**is_**', '(', expression, ',', type, ')' ;

A.5.20   The Undefined Expression

undefined expression = '**undefined**' ;

A.5.21   The Precondition Expression

pre-condition expression = '**pre_**', '(', expression,
                            [{ ',', expression }], ')' ;

A.5.22   Base Class Membership (VDM++ and VDM-RT)

isofbaseclass expression = '**isofbaseclass**', '(', name, expression, ')' ;

A.5.23   Class Membership (VDM++ and VDM-RT)

isofclass expression = '**isofclass**', '(', name, expression, ')' ;

A.5.24   Same Base Class Membership (VDM++ and VDM-RT)

samebaseclass expression = '**samebaseclass**', '(', expression,
                            expression, ')' ;

A.5.25   Same Class Membership (VDM++ and VDM-RT)

sameclass expression = '**sameclass**', '(', expression,
                        expression, ')' ;

---

[2]**Note:** no delimiter is allowed

A.5.26 History Expressions (VDM++ and VDM-RT)

act expression = '**#act**', '(', name, ')'
            | '**#act**', '(', name list, ')' ;

fin expression = '**#fin**', '(', name, ')'
            | '**#fin**', '(', name list, ')' ;

active expression = '**#active**', '(', name, ')'
              | '**#active**', '(', name list, ')' ;

req expression = '**#req**', '(', name, ')'
            | '**#req**', '(', name list, ')' ;

waiting expression = '**#waiting**', '(', name, ')'
              | '**#waiting**', '(', name list, ')' ;

A.5.27 Time Expressions (VDM-RT)

time expression = '**time**' ;

A.5.28 Names

name = identifier, ['`', identifier] ;

name list = name, { ',', name } ;

old name = identifier, '~' ;

A.6 State Designators

state designator = name
               | field reference
               | map or sequence reference ;

field reference = state designator, '**.**', identifier ;

map or sequence reference = state designator, '(', expression, ')' ;

A.7   Statements

```
statement  =  let statement
           |   let be statement
           |   def statement
           |   block statement
           |   general assign statement
           |   if statement
           |   cases statement
           |   sequence for loop
           |   set for loop
           |   index for loop
           |   while loop
           |   nondeterministic statement
           |   call statement
           |   specification statement
           |   start statement
           |   start list statement
           |   stop statement
           |   stop list statement
           |   duration statement
           |   cycles statement
           |   return statement
           |   always statement
           |   trap statement
           |   recursive trap statement
           |   exit statement
           |   error statement
           |   identity statement  ;
```

A.7.1   Local Binding Statements

let statement = '**let**', local definition, { ' **,** ', local definition },
                 '**in**', statement ;

local definition = value definition
                 | function definition ;

let be statement = '**let**', multiple bind, ['**be**', '**st**', expression], '**in**',
                     statement ;

def statement = '**def**', equals definition,
                 { '**;**', equals definition }, ['**;**'],
                 '**in**', statement ;

equals definition = pattern bind, '=', expression ;

A.7.2 Block and Assignment Statements

block statement = '(', { dcl statement },
                 statement, { ';', statement }, [';'], ')' ;

dcl statement = '**dcl**', assignment definition,
              { ',', assignment definition }, ';' ;

assignment definition = identifier, ':', type, [':=', expression] ;

general assign statement = assign statement
                      | multiple assign statement ;

assign statement = state designator, ':=', expression ;

multiple assign statement = '**atomic**', '(' assign statement, ';',
                   assign statement,
                   [{ ';', assign statement }], ')' ;

A.7.3 Conditional Statements

if statement = '**if**', expression, '**then**', statement,
           { elseif statement },
           ['**else**', statement] ;

elseif statement = '**elseif**', expression, '**then**', statement ;

cases statement = '**cases**', expression, ':',
           cases statement alternatives,
           [',', others statement], '**end**' ;

cases statement alternatives = cases statement alternative,
                       { ',', cases statement alternative } ;

cases statement alternative = pattern list, '->', statement ;

others statement = '**others**', '->', statement ;

A.7.4   Loop Statements

sequence for loop = '**for**', pattern bind, '**in**',
                    expression, '**do**', statement ;

set for loop = '**for**', '**all**', pattern, '**in set**', expression,
                '**do**', statement ;

index for loop = '**for**', identifier, '=', expression, '**to**', expression,
                    ['**by**', expression],
                    '**do**', statement ;

while loop = '**while**', expression, '**do**', statement ;

A.7.5   The Nondeterministic Statement

nondeterministic statement = '||', '(', statement,
                                { ',', statement }, ')' ;

A.7.6   Call and Return Statements

In VDM-SL a call statement looks like:

call statement = name, '(',
                    [expression list], ')' ;

In VDM++ and VDM-RT a call statement looks like:

call statement = [object designator, '.'],
                    name, '(', [expression list], ')', ;

object designator = name
                    | self expression
                    | new expression
                    | object field reference
                    | object apply ;

object field reference = object designator, '.', identifier ;

object apply = object designator, '(', [expression list], ')' ;

return statement = '**return**', [expression] ;

A.7.7   The Specification Statement

specification statement = '[', implicit operation body, ']' ;

A.7.8   Start and Start List Statements (VDM++ and VDM-RT)

start statement = '**start**', '(', expression, ')' ;

start list statement = '**startlist**', '(', expression, ')' ;

A.7.9   Stop and Stop List Statements (VDM++ and VDM-RT)

stop statement = '**stop**', '(', expression, ')' ;

stop list statement = '**stoplist**', '(', expression, ')' ;

A.7.10   The Duration and Cycles Statements (VDM-RT)

duration statement = '**duration**', '(', expression, ')',
                        statement ;

cycles statement = '**cycles**', '(', expression, ')',
                        statement ;

A.7.11   Exception Handling Statements

always statement = '**always**', statement, '**in**', statement ;

trap statement = '**trap**', pattern bind, '**with**', statement,
                    '**in**', statement ;

recursive trap statement = '**tixe**', traps, '**in**', statement ;

traps = '{', pattern bind, '|->', statement,
            { ',', pattern bind, '|->', statement }, '}' ;

exit statement = '**exit**', [expression] ;

A.7.12   The Error Statement

error statement = '**error**' ;

A.7.13   The Identity Statement

identity statement = '**skip**' ;

A.8   Patterns and Bindings

A.8.1   Patterns

pattern = pattern identifier
    | match value
    | set enum pattern
    | set union pattern
    | seq enum pattern
    | seq conc pattern
    | map enumeration pattern
    | map muinon pattern
    | tuple pattern
    | object pattern
    | record pattern ;

pattern identifier = identifier | '−' ;

match value = '(', expression, ')'
     | symbolic literal ;

set enum pattern = '{', [pattern list], '}' ;

set union pattern = pattern, 'union', pattern ;

seq enum pattern = '[', [pattern list], ']' ;

seq conc pattern = pattern, 'ˆ', pattern ;

map enumeration pattern = '{', maplet pattern list, '}'
       | '{', '|->', '}' ;

maplet pattern list = maplet pattern, { ',', maplet pattern } ;

maplet pattern = pattern, '|->', pattern ;

map muinon pattern = pattern, '**munion**', pattern ;

tuple pattern = 'mk_', '(', pattern, ',', pattern list, ')' ;

record pattern = 'mk_',[3] name, '(', [pattern list], ')' ;

object pattern = '**obj_**', identifier, '(', [field pattern list], ')'[4];

field pattern list = field pattern, { ',', field pattern } ;

field pattern = identifier, '|->', pattern ;

pattern list = pattern, { ',', pattern } ;

---

[3]**Note:** no delimiter is allowed
[4]**Note:** object pattern is only be used in VDM++ and VDM-RT.

A.8.2   Bindings

pattern bind =   pattern | bind ;

bind =   set bind | type bind ;

set bind =   pattern, '**in set**', expression ;

type bind =   pattern, ':', type ;

bind list =   multiple bind, { ',', multiple bind } ;

multiple bind =   multiple set bind
              | multiple type bind ;

multiple set bind =   pattern list, '**in set**', expression ;

multiple type bind =   pattern list, ':', type ;

type bind list =   type bind, { ',', type bind } ;

Appendix B

Lexical Specification

B.1 Characters

The characters that comprise a valid VDM specification are defined in terms of Unicode code-points. The actual character encoding of a VDM source file (for example UTF-8, ISO-Latin-1 or Shift-JIS) is not defined, and the tool support is responsible for converting whatever encoding is used into Unicode during the parse of the file.

All VDM keywords and delimiter tokens are composed of characters from the Basic Latin block ("ASCII" codepoints less than U+0080). On the other hand, user identifiers (variable names, function names and so on) can be composed of a rich variety of Unicode codepoints, reflecting the need for fully internationalized specifications.

All Unicode codepoints have a "category". Certain categories are entirely excluded from the set of codepoints that are permitted in identifiers. This prevents, say, puntuation characters from being used. On the other hand, to provide a degree of compatibility with the original VDM ISO standard, and for backward compatibility, there are different rules for the formation of user identifiers that only use ASCII characters. For example, the underscore is permitted in identifiers (U+005F), even though this is in the connecting punctuation category, which would not normally be allowed.

See `http://www.fileformat.info/info/unicode/category/index.htm` for more information about categories.

initial letter:	
if	codepoint < U+0100
then	Any character in categories Ll, Lm, Lo, Lt, Lu or U+0024 (a dollar sign)
else	Any character except categories Cc, Zl, Zp, Zs, Cs, Cn, Nd, Pc

following letter:	
if	codepoint < U+0100
then	Any character in categories Ll, Lm, Lo, Lt, Lu, Nd or U+0024 (a dollar sign) or U+005F (underscore) or U+0027 (apostrophe)
else	Any character except categories Cc, Zl, Zp, Zs, Cs, Cn

digit:

```
0   1   2   3   4   5   6   7   8   9
```

hexadecimal digit:

```
0   1   2   3   4   5   6   7   8   9
A   B   C   D   E   F
a   b   c   d   e   f
```

octal digit:

```
0   1   2   3   4   5   6   7
```

Table B.1: Character set

B.2 Symbols

The following kinds of symbols exist: keywords, delimiters, symbolic literals, and comments. The transformation from characters to symbols is given by the following rules; these use the same notation as the syntax definition but differ in meaning in that no separators may appear between adjacent terminals. Where ambiguity is possible otherwise, two consecutive symbols must be separated by a separator.

```
keyword = ‘#act’ | ‘#active’ | ‘#fin’ | ‘#req’ | ‘#waiting’ | ‘abs’
        | ‘all’ | ‘always’ | ‘and’ | ‘as’ | ‘async’ | ‘atomic’ | ‘be’
        | ‘bool’ | ‘by’ | ‘card’ | ‘cases’ | ‘char’ | ‘class’
        | ‘comp’ | ‘compose’ | ‘conc’ | ‘cycles’ | ‘dcl’ | ‘def’
        | ‘definitions’ | ‘dinter’ | ‘div’ | ‘dlmodule’ | ‘do’
        | ‘dom’ | ‘dunion’ | ‘duration’ | ‘elems’ | ‘else’ | ‘elseif’
        | ‘end’ | ‘error’ | ‘errs’ | ‘exists’ | ‘exists1’ | ‘exit’
        | ‘exports’ | ‘ext’ | ‘false’ | ‘floor’
        | ‘for’ | ‘forall’ | ‘from’ | ‘functions’ | ‘hd’ | ‘if’ | ‘in’
        | ‘inds’ | ‘inmap’ | ‘instance’ | ‘int’ | ‘inter’
        | ‘imports’ | ‘init’ | ‘inv’ | ‘inverse’ | ‘iota’ | ‘is’
        | ‘isofbaseclass’ | ‘isofclass’ | ‘lambda’ | ‘len’ | ‘let’
        | ‘map’ | ‘measure’ | ‘merge’ | ‘mod’ | ‘module’ | ‘mu’
        | ‘munion’ | ‘mutex’ | ‘nat’ | ‘nat1’ | ‘new’ | ‘nil’ | ‘not’ | ‘of’
        | ‘operations’ | ‘or’ | ‘others’ | ‘per’ | ‘periodic’ | ‘post’
        | ‘power’ | ‘pre’ | ‘private’ | ‘protected’ | ‘psubset’
        | ‘public’ | ‘rat’ | ‘rd’ | ‘real’ | ‘rem’ | ‘renamed’
        | ‘responsibility’ | ‘return’ | ‘reverse’ | ‘rng’
        | ‘samebaseclass’ | ‘sameclass’ | ‘self’ | ‘seq’ | ‘seq1’
        | ‘set’ | ‘skip’ | ‘specified’ | ‘sporadic’ | ‘st’ | ‘start’
        | ‘startlist’ | ‘state’ | ‘stop’ | ‘stoplist’
        | ‘struct’ | ‘subclass’ | ‘subset’ | ‘sync’
        | ‘system’ | ‘then’ | ‘thread’ | ‘threadid’ | ‘time’ | ‘tixe’
        | ‘tl’ | ‘to’ | ‘token’ | ‘traces’ | ‘trap’ | ‘true’ | ‘types’
        | ‘undefined’ | ‘union’ | ‘uselib’ | ‘values’
        | ‘variables’ | ‘while’ | ‘with’ | ‘wr’ | ‘yet’ | ‘RESULT’ ;

identifier = initial letter, { following letter } ;
```

Note that in VDM-RT the CPU and BUS classes are reserved and cannot be redefined by the user. These two predefined classes contain the functionality described in Section 13.3 above.

All identifiers beginning with one of the reserved prefixes are reserved: **init_**, **inv_**, **is_**, **mk_**, **post_** and **pre_**.

```
type variable identifier = ‘@’, identifier ;
```

is basic type = '**is_**', ('**bool**' | '**nat**' | '**nat1**' | '**int**' | '**rat**'
            | '**real**' | '**char**' | '**token**') ;

symbolic literal = numeric literal | boolean literal
                  | nil literal | character literal | text literal
                  | quote literal ;

numeral = digit, { digit } ;

numeric literal = decimal literal | hexadecimal literal ;

exponent = ('E' | 'e'), ['+' | '−'], numeral ;

decimal literal = numeral, ['.', digit, { digit }], [exponent] ;

hexadecimal literal = ('0x' | '0X'), hexadecimal digit, { hexadecimal digit } ;

boolean literal = '**true**' | '**false**' ;

nil literal = '**nil**' ;

character literal = ' ' ', character | escape sequence
                | ' ' ' ;

escape sequence = '\\' | '\r' | '\n' | '\t' | '\f' | '\e' | '\a'
               | '\x' hexadecimal digit,hexadecimal digit
               | '\u' hexadecimal digit,hexadecimal digit,
                  hexadecimal digit,hexadecimal digit
               | '\c' character
               | '\' octal digit, octal digit, octal digit
               | '\"' | '\'' | ;

text literal = ' " ', { ' \" ' | character | escape sequence }, ' " ' ;

quote literal = '<', identifier, '>' ;

Single-line comment = '−−', { character − newline }, newline ;

Multiple-line comment = '/*', { character }, '*/' ;

The escape sequences given above are to be interpreted as follows:

Sequence	Interpretation
'\\'	U+005C (backslash character)
'\r'	U+000D (return character)
'\n'	U+000A (newline character)
'\t'	U+0009 (tab character)
'\f'	U+000C (formfeed character)
'\e'	U+001B (escape character)
'\a'	U+0007 (alarm (bell))
'\x' hexadecimal digit, hexadecimal digit	U+00xy (hex representation of character (e.g. \x41 is 'A'))
'\u' hexadecimal digit, hexadecimal digit, hexadecimal digit, hexadecimal digit	U+abcd (hex representation of character (e.g. \u0041 is 'A'))
'\c' character	U+00nn (control character) (e.g. \cA ≡ \x01)
'\' octal digit, octal digit, octal digit	U+00nn (octal representation of character)
'\"'	U+0022 (double quote)
'\''	U+0027 (apostrophe)

Table B.2: Escape sequences

Appendix C

Operator Precedence

The precedence ordering for operators in the concrete syntax is defined using a two-level approach: operators are divided into families, and an upper-level precedence ordering, $>$, is given for the families, such that if families F_1 and F_2 satisfy

$$F_1 > F_2$$

then every operator in the family F_1 is of a higher precedence than every operator in the family F_2.

The relative precedences of the operators within families is determined by considering type information, and this is used to resolve ambiguity. The type constructors are treated separately, and are not placed in a precedence ordering with the other operators.

There are six families of operators, namely Combinators, Applicators, Evaluators, Relations, Connectives and Constructors:

Combinators: Operations that allow function and mapping values to be combined, and function, mapping and numeric values to be iterated.

Applicators: Function application, field selection, sequence indexing, etc.

Evaluators: Operators that are non-predicates.

Relations: Operators that are relations.

Connectives: The logical connectives.

Constructors: Operators that are used, implicitly or explicitly, in the construction of expressions; e.g. **if-then-elseif-else**, '`|->`', '`...`', etc.

The precedence ordering on the families is:

combinators $>$ applicators $>$ evaluators $>$ relations $>$ connectives $>$ constructors

C.1 The Family of Combinators

These combinators have the highest family priority.

combinator = iterate | composition ;

iterate = '**' ;

composition = '**comp**' ;

precedence level	combinator
1	**comp**
2	**iterate**

C.2 The Family of Applicators

All applicators have equal precedence.

applicator = subsequence
              | apply
              | function type instantiation
              | field select ;

subsequence = expression, '(', expression, ',', '...', ',',
                expression, ')' ;

apply = expression, '(', [expression list], ')' ;

function type instantiation = expression, '[', type, { ',', type }, ']' ;

field select = expression, '.', identifier ;

C.3 The Family of Evaluators

The family of evaluators is divided into nine groups, according to the type of expression they are used in.

```
evaluator  =  arithmetic prefix operator
            |   set prefix operator
            |   sequence prefix operator
            |   map prefix operator
            |   map inverse
            |   arithmetic infix operator
            |   set infix operator
            |   sequence infix operator
            |   map infix operator  ;
```

arithmetic prefix operator = '+' | '−' | '**abs**' | '**floor**' ;

set prefix operator = '**card**' | '**power**' | '**dunion**' | '**dinter**' ;

sequence prefix operator = '**hd**' | '**tl**' | '**len**'
                          | '**inds**' | '**elems**' | '**conc**' ;

map prefix operator = '**dom**' | '**rng**' | '**merge**' | '**inverse**' ;

arithmetic infix operator = '+' | '−' | '*' | '/' | '**rem**' | '**mod**' | '**div**' ;

set infix operator = '**union**' | '**inter**' | '\' ;

sequence infix operator = '^' ;

map infix operator = '**munion**' | '++' | '<:' | '<−:' | ':>' | ':−>' ;

The precedence ordering follows a pattern of analogous operators. The family is defined in the following table.

C.4   The Family of Relations

This family includes all the relational operators whose results are of type **bool**.

relation = relational infix operator | set relational operator ;

relational infix operator = '=' | '<>' | '<' | '<=' | '>' | '>=' ;

set relational operator = '**subset**' | '**psubset**' | '**in set**' | '**not in set**' ;

thb

precedence level	arithmetic	set	map	sequence
1	+ −	**union** \	**munion** ++	^
2	* / **rem** **mod** **div**	**inter**		
3			**inverse**	
4			<: <−:	
5			:> :−>	
6	(unary) + (unary) − **abs** **floor**	**card** **power** **dinter** **dunion**	**dom** **rng** **merge**	**len** **elems** **hd tl** **conc** **inds**

Table C.1: Operator precedence

precedence level	relation	
1	<=	<
	>=	>
	=	<>
	subset	**psubset**
	in set	**not in set**

All operators in the Relations family have equal precedence. Typing dictates that there is no meaningful way of using them adjacently.

C.5   The Family of Connectives

This family includes all the logical operators whose result is of type **bool**.

connective = logical prefix operator | logical infix operator ;

logical prefix operator = '**not**' ;

logical infix operator = '**and**' | '**or**' | '=>' | '<=>' ;

precedence level	connective
1	<=>
2	=>
3	**or**
4	**and**
5	**not**

C.6   The Family of Constructors

This family includes all the operators used to construct a value. Their priority is given either by brackets, which are an implicit part of the operator, or by the syntax.

C.7   Grouping

The grouping of operands of the binary operators are as follows:

Combinators: Right grouping.

Applicators: Left grouping.

Connectives: The '=>' operator has right grouping. The other operators are associative and therefore right and left grouping are equivalent.

Evaluators: Left grouping[1].

Relations: No grouping, as it has no meaning.

Constructors: No grouping, as it has no meaning.

C.8   The Type Operators

Type operators have their own separate precedence ordering, as follows:

1. Function types: `->`, `+>` (right grouping).

2. Union type: `|` (left grouping).

3. Other binary type operators: `*` (no grouping).

4. Map types: **map** ... **to** ... and **inmap** ... **to** ... (right grouping).

5. Unary type operators: **seq of**, **seq1 of**, **set of**.

---

[1]Except the "map domain restrict to" and the "map domain restrict by" operators which have a right grouping. This is not standard.

Appendix D

Differences between the Concrete Syntaxes

Whe VDM was originally developed a mathematical syntax was used and this have also been retained in the ISO/VDM-SL standard. However, most VDM tools today mainly use an ASCII syntax. Below is a list of the symbols which are different in the mathematical syntax and the ASCII syntax:

Mathematical syntax	ASCII syntax
·	`&`
$\times$	`*`
$\leq$	`<=`
$\geq$	`>=`
$\neq$	`<>`
$\xrightarrow{o}$	`==>`
$\rightarrow$	`->`
$\Rightarrow$	`=>`
$\Leftrightarrow$	`<=>`
$\mapsto$	`|->`
$\triangleq$	`==`
$\uparrow$	`**`
$\dagger$	`++`
$\underline{m}$	**`munion`**
$\triangleleft$	`<:`
$\triangleright$	`:>`
$\triangleleft$	`<-:`
$\triangleright$	`:->`
$\subset$	**`psubset`**
$\subseteq$	**`subset`**
$\curvearrowright$	`^`
$\cap$	**`dinter`**
$\cup$	**`dunion`**

215

Mathematical syntax	ASCII syntax
$\mathcal{F}$	`power`
$\ldots$-set	`set of ...`
$\ldots^{*}$	`seq of ...`
$\ldots^{+}$	`seq1 of ...`
$\ldots\xrightarrow{m}\ldots$	`map ...  to ...`
$\ldots\xleftrightarrow{m}\ldots$	`inmap ...  to ...`
μ	`mu`
$\mathbb{B}$	`bool`
$\mathbb{N}$	`nat`
$\mathbb{Z}$	`int`
$\mathbb{R}$	`real`
$\neg$	`not`
$\cap$	`inter`
$\cup$	`union`
$\in$	`in set`
$\notin$	`not in set`
$\wedge$	`and`
$\vee$	`or`
$\forall$	`forall`
$\exists$	`exists`
$\exists!$	`exists1`
λ	`lambda`
ι	`iota`
$\ldots^{-1}$	`inverse ...`

Index